

# **CSE 431**

# **Computer Architecture**

# **Fall 2022**

## **VLIW Datapaths**

Kiwan Maeng

[Adapted from *Computer Organization and Design, 5<sup>th</sup> Edition*,

Patterson & Hennessy, © 2014, MK

With additional thanks/credits to Amir Roth, Milo Martin]

# How Was Your Exam?


- ❑ Let's keep up the good work!
- ❑ Exam results will be published as soon as possible.
- ❑ Trying to give as much partial credits as possible!
- ❑ So far, it seems like slightly over 20% did a great job.

# Reviewing Q2-1...

Q2. Answer the following questions (15pt total).

- Q2-1. Write a MIPS assembly program that is equivalent to the following C program. Assume that `tmp` is stored in register `$t0`, and the base address of `Array` is stored in register `$s0`. (5pt)

```
int32_t Array[100]; int32_t tmp;  
...  
Array[3] = tmp;
```



## Quiz 3 is Out!

- ❑ Due 10/3 11:59PM
- ❑ No extension, as always

# Project 1 (due 11/10) Will Be Out, But...

- ❑ Implementing a simple dynamic superscalar processor over a skeleton code
- ❑ You cannot do it yet because we slowed down the course and did not learn dynamic superscalar...
- ❑ Still, you will have enough time

# Review: Pipeline Hazards

## ❑ Structural hazards

- ❑ Design pipeline to eliminate structural hazards

## ❑ Data hazards – read after write (RAW)

- ❑ Use data forwarding inside the pipeline
- ❑ For those cases that forwarding won't solve (e.g., load-use) include hazard hardware to insert stalls in the instruction stream

## ❑ Control hazards – `beq, bne, j, jr, jal, jalr`

- ❑ Move decision point as early in the pipeline as possible (e.g., Decode) – reduces number of stalls at the cost of additional hardware
- ❑ Predict – with dynamic branch prediction, can reduce the impact of control hazard flushes even further if the branch prediction (contained in a BHT in the Fetch stage) is correct (shoot for 97%+ accuracy) and if the branched-to instruction is cached (in a BTB also in the Fetch stage)

# Extracting Yet *More* Performance

- ❑ Increase the depth of the pipeline to increase the clock rate (CPI still 1, IC unchanged) – **superpipelining**
  - ❑ Higher parallelism and better performance
  - ❑ The more stages in the pipeline, the more forwarding/hazard hardware needed and the more pipeline latch overhead (i.e., the pipeline latch accounts for a larger and larger percentage of the clock cycle time)
- ❑ Fetch (and execute) more than one instructions at one time (expand every pipeline stage to accommodate multiple instructions) – **multiple-issue**
  - ❑ The instruction execution rate, CPI, will be less than 1, so instead we use **IPC**: instructions per clock cycle
    - E.g., a 6 GHz, four-way multiple-issue core can execute at a peak rate of 24 billion instructions per second with a best-case CPI of 0.25 or a best-case IPC of 4
    - How many instructions are in flight together?
  - ❑ More susceptible to resource conflicts

# Types of Code Parallelism

- ❑ **Instruction-level parallelism (ILP)** of a program – a measure of the average number of instructions in a program that a core *might* be able to execute at the same time
  - ❑ The IPC of an optimal processor
  - ❑ Mostly determined by the number of true (data) dependencies and procedural (control) dependencies in relation to the number of other instructions
  - ❑ E.g.,  
add \$1, \$1, \$2  
sub \$3, \$2, \$2  
add \$4, \$2, \$2  
sub \$5, \$6, \$7  
add \$8, \$9, \$10
  - ❑ E.g.,  
add \$1, \$1, \$2  
sub \$3, \$1, \$2  
add \$4, \$3, \$2  
sub \$5, \$6, \$7  
add \$8, \$9, \$10



# Types of Code Parallelism

## □ Data-level parallelism (DLP)

```
DO I = 1 TO 100  
    A[I] = A[I] + 1  
CONTINUE
```

- The example shown has lots of data parallelism but almost no instruction parallelism if compiled in the obvious way.
- Loop unrolling is a common compiler optimization. If completely unrolled, and if we had 100 arithmetic/addressing units and 100 memory ports, we could achieve a speedup of 100 over a scalar core.
  - It in a sense converts data parallelism to instruction parallelism

# Datapath (Core) Parallelism

- ❑ Machine-level parallelism of a datapath – a measure of the ability of the datapath to take advantage of the ILP of the program
  - ❑ Determined by the number of instructions that can be fetched and executed at the same time
- ❑ To achieve high performance, need both instruction-level parallelism and machine-level parallelism
- ❑ Some additional examples
  - ❑ SIMD instructions, short-vector packed data
  - ❑ Multithreading (MT)
  - ❑ Simultaneous Multithreading (SMT) (aka Hyperthreading)
  - ❑ Multicore, homogeneous and heterogeneous (GPGPUs)
  - ❑ Multiprocessor
  - ❑ others?

# Static vs Dynamic Processors

- ❑ There are two major ways of implementing a multiple-issue processor, with the main difference being the division of work between the **compiler** (static) and the **hardware** (dynamic)
- ❑ Although we describe these as distinct approaches, in reality one approach often borrows techniques from the other, and neither of them can claim to be perfectly *pure*

# Issue Structure vs Scheduling

- ❑ There are two primary and distinct responsibilities that must be dealt with in a multiple-issue pipeline:
- ❑ Packaging instructions into issue slots
  - ❑ Static-issue design vs Dynamic-issue design
- ❑ Scheduling
  - ❑ Dealing with data and control hazards

# Taxonomy of Multiple-Issue Machines

Common name	Issue structure	Hazard detection	Scheduling	Distinguishing characteristic	Examples
Superscalar (static)	Dynamic	Hardware	Static	In-order execution	Mostly in the embedded space: MIPS and ARM, including the ARM Cortex A8
Superscalar (dynamic)	Dynamic	Hardware	Dynamic	Some out-of-order execution, but no speculation	None at the present
Superscalar (speculative)	Dynamic	Hardware	Dynamic with speculation	Out-of-order execution with speculation	Intel Core i3, i5, i7; AMD Phenom; IBM Power 7
VLIW/LIW	Static	Primarily software	Static	All hazards determined and indicated by compiler (often implicitly)	Most examples are in signal processing, such as the TI C6x
EPIC	Primarily static	Primarily software	Mostly static	All hazards determined and indicated explicitly by the compiler	Itanium

# Multiple Instruction Issue Possibilities

- ❑ Fetch and issue **more than one** instruction in a cycle

## 1. **Statically-scheduled (in-order)**

- ❑ **Very Long Instruction Word (VLIW)** e.g., TransMeta (4-wide)
  - Compiler figures out what can be done in parallel, so the hardware can be dumb and low power (instruction issue is static)
  - Compiler must group parallel instr's, requires new binaries
- ❑ **SuperScalar** e.g., Pentium (2-wide), ARM Cortex-A8 (2-wide)
  - Hardware figures out what can be done in parallel
  - Executes unmodified sequential programs (but issue is dynamic)
- ❑ **Explicitly Parallel Instruction Computing (EPIC)** e.g., Intel Itanium (6-wide)
  - A compromise: compiler does some, hardware does the rest

## 2. **Dynamically-scheduled (out-of-order) SuperScalar**

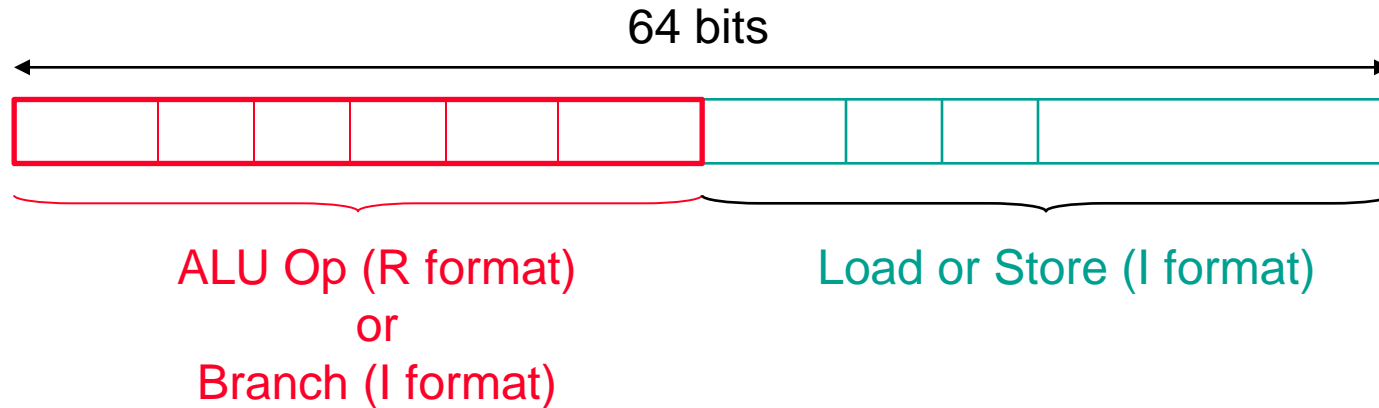
- ❑ Hardware dynamically determines what can be done in parallel (can extract much more ILP with OOO processing)
- ❑ E.g., Intel Core i7 (4-wide, 8-way SMT, 4 cores/chip), IBM Power8 (8-wide, 8-way SMT, 12 cores/chip)

# VLIW Multiple Issue Datapaths

- ❑ VLIW multiple-issue datapath has the compiler to **statically** decide which instructions to issue and execute simultaneously
  - ❑ **Issue packet** – the set of instructions that are *bundled* together and issued in one clock cycle – think of it as one **large** instruction with multiple operations
    - **Compiler guarantees that the instr's within a packet are independent** (this usually means some of the instr's in the packet are `nops`)
  - ❑ The mix of instructions in the **packet** (bundle) is usually restricted – a single “instruction” with several predefined, “slotted” fields
  - ❑ The compiler does static branch prediction and code scheduling (with renaming) to reduce control hazards and eliminate WAW & WAR data hazards
- ❑ VLIWs have
  - ❑ Multiple functional units
  - ❑ Multi-ported register files
  - ❑ Wide program bus

# An Example: A VLIW MIPS

- ❑ Consider a 2-wide issue MIPS with a 2-instr packet that is fetched, decoded and issued for execution as a pair





# Working Around Structural Hazards

- ❑ FETCH – 2-wide issue packets so fetch a 64-bit “instruction” packet (cache blocks of at least 8B)

Address	Instruction type	Pipeline Stages						
n	ALU/branch	IF	ID	EX	MEM	WB		
n + 4	Load/store	IF	ID	EX	MEM	WB		
n + 8	ALU/branch		IF	ID	EX	MEM	WB	
n + 12	Load/store		IF	ID	EX	MEM	WB	
n + 16	ALU/branch			IF	ID	EX	MEM	WB
n + 20	Load/store			IF	ID	EX	MEM	WB

- ❑ ID (and WB) – Need a 4 read port and 2 write port RF and two decoders
- ❑ EX – Need a separate memory address adder
- ❑ MEM – Only one of the pair touches data memory (so no structural hazard there)

# Handling Data Hazard

## □ Data hazards ?

- Load-use (RAW) hazards have to be split by the compiler into two packets, while name dependence (WAR, WAW) can be avoided by renaming

lw	\$2,	400 (\$3)
add	\$1,	\$2, \$4

lw	\$2,	400 (\$3)
add	\$2,	\$3, \$4

# Handling Data Hazard

- ❑ To what extent can the compiler (or the datapath) reorder instructions? Are there execution-order constraints?

Original	possible?	possible?
instr 1 instr 2 consecutive	instr 2 instr 1 consecutive	instr 1 and instr 2 simultaneous

- ❑ **Instruction dependencies** imply that reordering instructions is not possible

- ❑ true dependence (or, data dep., flow dep.) (can~~not~~ reorder)

a = .

. = a      **RAW, read after write**

- ❑ anti-dependence (renaming allows reordering)

. = a

a = .      **WAR, write after read**

- ❑ output dependence (renaming allows reordering)

a = .

a = .      **WAW, write after write**

- ❑ Anti and output dependencies are also called **Name Dependencies**

# A MIPS VLIW (2-issue) Pipelined Datapath

- ❑ Clearly, this *two-issue* processor can improve performance by up to a factor of *two*
- ❑ Doing so, however, requires that *twice* as many instructions be overlapped in execution
- ❑ This additional overlap increases the relative performance loss from *data* and *control hazards*
  - ❑ Example: load-use hazard, ALU instructions
  - ❑ In a two-issue, five-stage pipeline, the result of a load instruction cannot be used on the next clock cycle, meaning that next two instructions cannot use the load result without stalling
- ❑ To effectively exploit the parallelism available in a multiple-issue processor, more ambitious compiler or hardware scheduling techniques are needed, and static multiple issue (like VLIW) requires that the compiler take on this role

# Code Scheduling Example

## ❑ Consider the following loop code

```
lp:   lw      $t0, 0($s1)    # $t0=array element
      addu    $t0, $t0, $s2  # add scalar in $s2
      sw      $t0, 0($s1)    # store result
      addi    $s1, $s1, -4   # decrement pointer
      bne     $s1, $0, lp    # branch if $s1 != 0
```

## ❑ Compiler “schedules” the instructions to avoid pipeline stalls

- ❑ Instructions in one bundle *must* be independent
- ❑ Must separate load-use instructions from their loads by one cycle (we are assuming a use latency of one cycle)
- ❑ Notice that the first two instructions have a load-use dependency (in red), the second and third as well as last two have true (RAW) data dependencies (in blue and green)
- ❑ Assume branches are perfectly predicted by the hardware

# The Scheduled Code (Not Unrolled)

	ALU or branch	Data transfer	CC
lp:		lw     \$t0, 0(\$s1)	1
	addi    \$s1, \$s1, -4		2
	addu    \$t0, \$t0, \$s2		3
	bne     \$s1, \$0, lp	sw     \$t0, 4(\$s1)	4
			5

- ❑ Note that displacement value for the sw has to be adjusted because the addi has been scheduled *before* it rather than *after* it as in the original code
- ❑ 4 clock cycles to execute 5 instructions for a
  - ❑ CPI of 0.8 (versus the best case of 0.5)
  - ❑ IPC of 1.25 (versus the best case of 2.0)
  - ❑ Three `nops` (that don't count towards performance !!)

# Multiple-Issue Datapath Responsibilities

- ❑ Must handle, with a combination of hardware and software fixes, the fundamental limitations of
  - ❑ How many instructions to **issue** in one clock cycle (issue packet)
    - Instruction issue is static
  - ❑ Storage (data) dependencies → data hazards
    - Limitation more severe in an in-order SuperScalar/VLIW core
  - ❑ Procedural dependencies → control hazards
    - Ditto, but even more severe
    - Use dynamic branch prediction to help resolve the ILP issue
    - Use loop unrolling (in the compiler) to increase ILP and reduce the occurrence of branches
  - ❑ Resource conflicts → structural hazards
    - A multiple-issue datapath has a much larger number of potential resource conflicts
    - Functional units may have to arbitrate for result buses and RF write ports
    - Resource conflicts can often be eliminated by duplicating the resource or by pipelining the resource

# Need for Loop Unrolling

- ❑ Compiler needs more instructions to fill empty slots
  - ❑ One loop iteration may *not* provide enough operations
  - ❑ We need more (useful) operations relative to the number of overhead instructions
  - ❑ This is where “loop unrolling” comes into the picture

```
for (i=0; i<100; i++) {  
    A[i] += B[i];  
}
```

```
for (i=0; i<100; i+=4) {  
    A[i] += B[i];  
    A[i+1] += B[i+1];  
    A[i+2] += B[i+2];  
    A[i+3] += B[i+3];  
}
```



# Loop Unrolling

- ❑ Compiler loop unrolling – multiple copies of the loop body are made and instructions from different iterations are scheduled together as a way to increase ILP
- ❑ Apply loop unrolling (4 times for our assembly code example) and then **schedule** the resulting code
  - ❑ Reduces loop-control overhead (fewer checks)
  - ❑ Schedule instr's so as to avoid load-use and other RAW hazards (more scope for reordering)
  - ❑ Schedule instr's so as to obey **loop-carried** dependencies (RAW), e.g., store in one loop followed by a load of the same register in the next loop
- ❑ During unrolling the compiler applies **register renaming** to eliminate all data dependencies that are *not* true data dependencies, i.e., WAW and WAR data hazards
  - ❑ This means loop unrolling leads to more register usage

# Unrolled Code Example

```
lp:   lw      $t0, 0($s1)      # $t0=array element
      lw      $t1, -4($s1)     # $t1=array element
      lw      $t2, -8($s1)     # $t2=array element
      lw      $t3, -12($s1)    # $t3=array element
      addu    $t0, $t0, $s2     # add scalar in $s2
      addu    $t1, $t1, $s2     # add scalar in $s2
      addu    $t2, $t2, $s2     # add scalar in $s2
      addu    $t3, $t3, $s2     # add scalar in $s2
      sw      $t0, 0($s1)      # store result
      sw      $t1, -4($s1)     # store result
      sw      $t2, -8($s1)     # store result
      sw      $t3, -12($s1)    # store result
      addi    $s1, $s1, -16     # decrement pointer
      bne     $s1, $0, lp      # branch if $s1 != 0
```

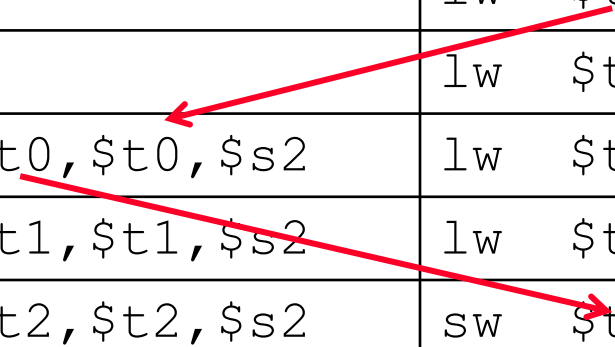
## The Scheduled Code (Unrolled) ... almost

	ALU or branch	Data transfer	CC
lp:		lw \$t0, 0(\$s1)	1
		lw \$t1, -4(\$s1)	2
		lw \$t2, -8(\$s1)	3
		lw \$t3, -12(\$s1)	4
		sw \$t0, 0(\$s1)	5
		sw \$t1, -4(\$s1)	6
		sw \$t2, -8(\$s1)	7
		sw \$t3, -12(\$s1)	8

- ❑ First schedule the data transfers (which can't be done in less than 8 cycles)
  - ❑ Notice the abundant use of registers

## The Scheduled Code (Unrolled) ... almost

	ALU or branch	Data transfer	CC
lp:		lw \$t0, 0(\$s1)	1
		lw \$t1, -4(\$s1)	2
	addu \$t0, \$t0, \$s2	lw \$t2, -8(\$s1)	3
	addu \$t1, \$t1, \$s2	lw \$t3, -12(\$s1)	4
	addu \$t2, \$t2, \$s2	sw \$t0, 0(\$s1)	5
	addu \$t3, \$t3, \$s2	sw \$t1, -4(\$s1)	6
		sw \$t2, -8(\$s1)	7
	bne \$s1, \$0, lp	sw \$t3, -12(\$s1)	8



- ❑ Next schedule the data use instr's in 4 cycles, being sure to leave (at least) one cycle between the load and its data use
- ❑ And schedule the branch instr in the last slot

## The Scheduled Code (Unrolled) ... almost

	ALU or branch	Data transfer	CC
lp:		lw \$t0, 0(\$s1)	1
		lw \$t1, -4(\$s1)	2
	addu \$t0, \$t0, \$s2	lw \$t2, -8(\$s1)	3
	addu \$t1, \$t1, \$s2	lw \$t3, -12(\$s1)	4
	addu \$t2, \$t2, \$s2	sw \$t0, 0(\$s1)	5
	addu \$t3, \$t3, \$s2	sw \$t1, -4(\$s1)	6
		sw \$t2, -8(\$s1)	7
	bne \$s1, \$0, lp	sw \$t3, -12(\$s1)	8

❑ One final instruction to schedule, where do we put

```
addi    $s1, $s1, -16
```

# The Scheduled Code (Unrolled)

	ALU or branch	Data transfer	CC
lp:	addi \$s1,\$s1,-16	lw \$t0,0(\$s1)	1
		lw \$t1,12(\$s1)	2
	addu \$t0,\$t0,\$s2	lw \$t2,8(\$s1)	3
	addu \$t1,\$t1,\$s2	lw \$t3,4(\$s1)	4
	addu \$t2,\$t2,\$s2	sw \$t0,16(\$s1)	5
	addu \$t3,\$t3,\$s2	sw \$t1,12(\$s1)	6
		sw \$t2,8(\$s1)	7
	bne \$s1,\$0,lp	sw \$t3,4(\$s1)	8


- ❑ Notice the adjustment in the memory address offsets
- ❑ Eight clock cycles to execute 14 instructions for a
  - ❑ CPI of 0.57 (versus the best case of 0.5)
  - ❑ IPC of 1.8 (versus the best case of 2.0), but at the cost of code size and more register use

# Steps for Loop Unrolling and Scheduling

- ❑ Determine that unrolling the loop would be useful (How?)
- ❑ Use different registers to avoid unnecessary constraints (What if not done?)
- ❑ Eliminate extra test and branch instructions
- ❑ Schedule the code

Loop unrolling and scheduling are complementary!

```
do i=1,N
  Dy(i) = Dy(i) + Da*Dx(i)
end do
```



```
do i=1,N,4
  Dy(i) = Dy(i) + Da*Dx(i)
  Dy(i+1) = Dy(i+1) + Da*Dx(i+1)
  Dy(i+2) = Dy(i+2) + Da*Dx(i+2)
  Dy(i+3) = Dy(i+3) + Da*Dx(i+3)
end do
```

# Speculation in VLIW Cores

- ❑ Speculation is used to allow execution of future instr's that (may) depend on the speculated instruction
  - ❑ Speculate on the outcome of a conditional branch (**branch prediction**)
  - ❑ Speculate that a store (for which we don't yet know the address) that precedes a load does not refer to the same address, allowing the load to be scheduled before the store (**load speculation**)
- ❑ Are these two the same or not?

```
sw $t2, 0($s3)
```

```
lw $t1, 0($s1)
```

```
lw $t1, 0($s1)
```

```
sw $t2, 0($s3)
```

- ❑ Depends on whether  $\$s3 == \$s1$ !



# Speculation in VLIW Cores

- ❑ Speculation is used to allow execution of future instr's that (may) depend on the speculated instruction
  - ❑ Speculate on the outcome of a conditional branch (**branch prediction**)
  - ❑ Speculate that a store (for which we don't yet know the address) that precedes a load does not refer to the same address, allowing the load to be scheduled before the store (**load speculation**)
- ❑ What if the speculation was wrong?
  - ❑ In a VLIW core the compiler inserts additional instructions that check the accuracy of the speculation and provides a fix-up routine to use when the speculation was incorrect
- ❑ Ignore and/or buffer **exceptions** created by speculatively executed instructions until it is clear that they should really occur (i.e., not allowed to change the machine state until you are sure)

# Compiler Support for VLIW Cores

- ❑ The compiler packs groups of **independent** instructions into the bundle
  - ❑ Done by code re-ordering
- ❑ The compiler uses loop unrolling to expose more ILP
  - ❑ Loop unrolling also reduces the number of conditional branches
- ❑ The compiler uses register renaming to solve name dependencies (WAR [anti] and WAW [output]) and ensures no load-use hazards occur by scheduling load-use instr's appropriately
- ❑ VLIWs primarily depend on the *compiler* for branch prediction
- ❑ The compiler predicts memory bank references to help minimize memory bank conflicts

# VLIW Advantages and Disadvantages

## ❑ Advantages

- ❑ Simpler hardware (potentially less power hungry)
  - Compiler does most of the work
- ❑ Potentially more scalable
  - Allow more instr's per VLIW bundle and add more FUs

## ❑ Disadvantages

- ❑ Programmer/compiler complexity and longer compilation times
  - Deep pipelines can be confusing as to what can be handled with forwarding and what needs to be stalled
- ❑ Lock-step operation, i.e., one slow can stall the entire packet
- ❑ Object (binary) code incompatibility – recompilation need
- ❑ Needs lots of program memory bandwidth
- ❑ Code bloat
  - `nops` are a waste of program memory space
  - Loop unrolling to expose more ILP uses more program memory space as well

# VLIW Philosophy

- ❑ Philosophy similar to RISC (simple instructions and hardware)
  - ❑ Except multiple instructions in parallel
  
- ❑ RISC (John Cocke, 1970s, IBM 801 minicomputer)
  - ❑ Compiler does the hard work to translate high-level language code to simple instructions
    - And, to reorder simple instructions for high performance
  - ❑ Hardware does little translation/decoding → very simple
  
- ❑ VLIW (Fisher, ISCA 1983)
  - ❑ Compiler does the hard work to find instruction level parallelism
  - ❑ Hardware stays as simple and streamlined as possible
    - Executes each instruction in a bundle in lock step
    - Simple → higher frequency, easier to design

# Track Record of VLIWs

- ❑ Started with “horizontal microcode”
- ❑ Academic projects
  - ❑ Yale ELI-512 [Fisher, ‘85]
  - ❑ Illinois IMPACT [Hwu, ‘91]
- ❑ Commercial attempts
  - ❑ Multiflow [Colwell+Fisher, ‘85] → failed
  - ❑ Cydrome [Rau, ‘85] → failed
  - ❑ Motorola, TI, ... embedded (DSP) cores → successful
    - TI TMS320C62xx and others
    - Lucent/Motorola StarCoreSC140
  - ❑ Intel Itanium [Fisher+Rau, ‘97] → ??
    - <http://en.wikipedia.org/wiki/Itanium>
  - ❑ Transmeta Crusoe [Ditzel, ‘99] → mostly failed
    - <http://en.wikipedia.org/wiki/Transmeta> (x86 to molecules)
  - ❑ VLIW features have also been added to SoC designs
    - Tensilica’s Xtensa LX2 processor incorporates a technology named Flexible Length Instruction eXtensions (FLIX) that allows multi-operation instructions