

# CMPEN 431

## Introduction to Computer Architecture

### Fall 2022

## Multiprocessor (1): Cache Coherence (1)

Kiwan Maeng

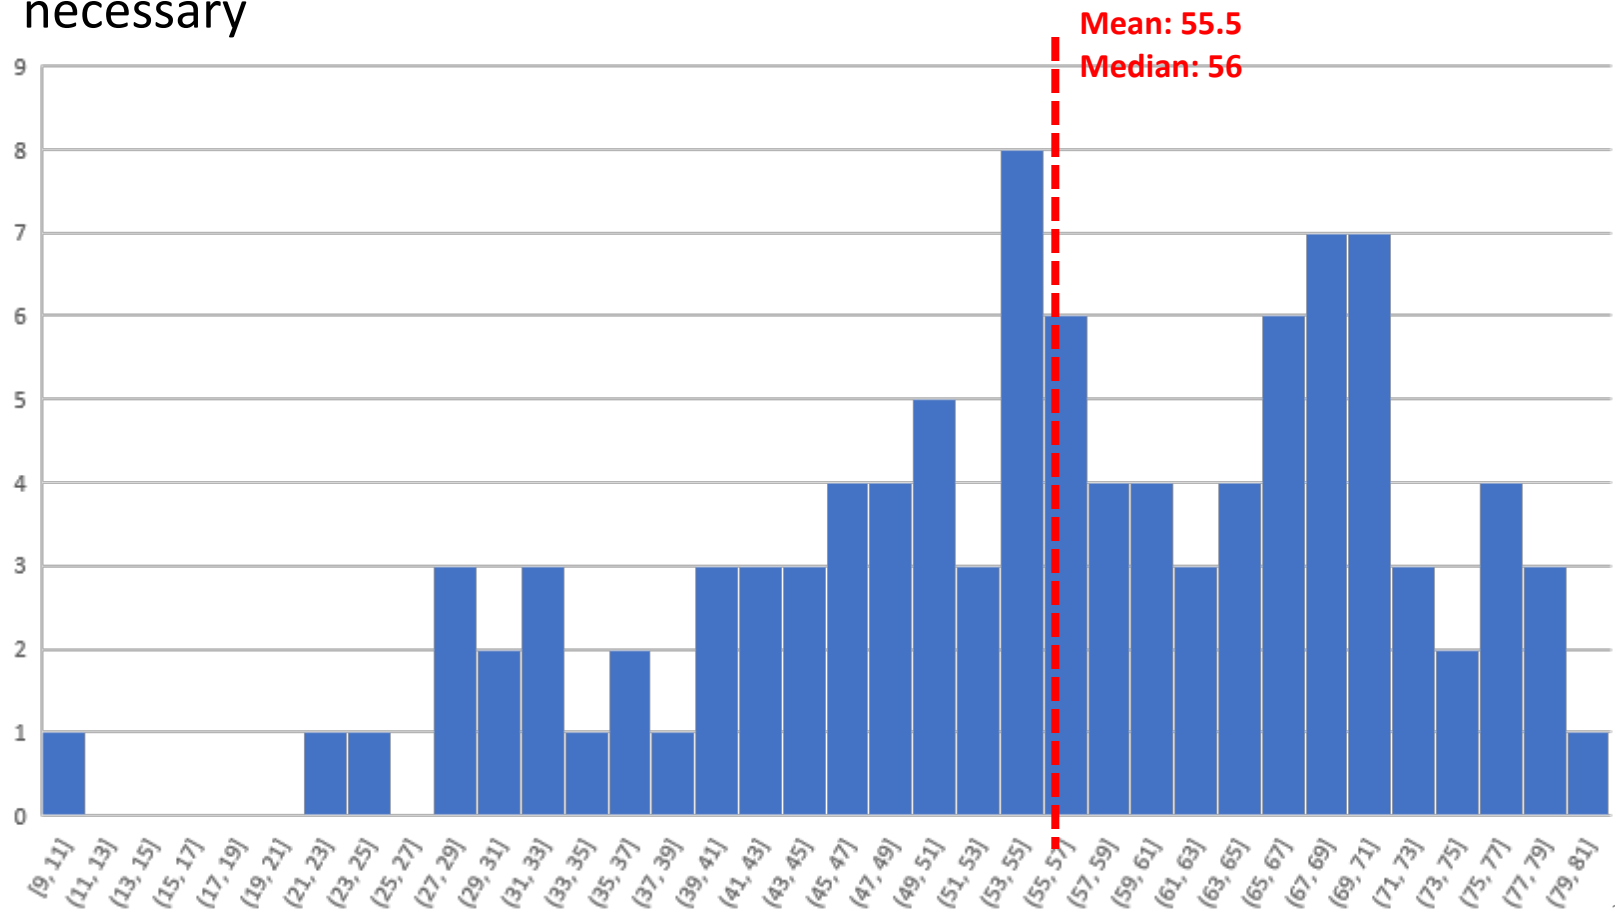
Slides adapted from ECE6100 @ GeorgiaTech  
by Hsien-Hsin Sean Lee &  
Computer Organization and Design, 5th Edition,  
Patterson & Hennessy, © 2014, MK

# Exam 2: Partial Answer Change


- Q2-2: What is the disadvantage of loop unrolling?
- The answer I taught during class (and the one I was expecting)
  - Binary size increase (use up more memory, increased I-cache misses, ...)
- The answer many students wrote
  - Increased use of register file
  - **I decided to give full credit to this answer as well.**
  - I never said this was a disadvantage
  - This is usually less of a problem because many other optimizations alleviate this problem
  - But there can be a case where this can be a problem... so I am giving full credit
  - Keep in mind that a much commonly discussed issue is the increase of binary size, which leads to increased instruction memory use + increased I-cache misses

# Exam 2 Results Are Out!


- Please check the solution & submit regrade request
- Submit regrade request by **11/10 (Thurs)** so that we can late-drop if necessary



# Exam 1 + 2 (each exam scaled to 100)

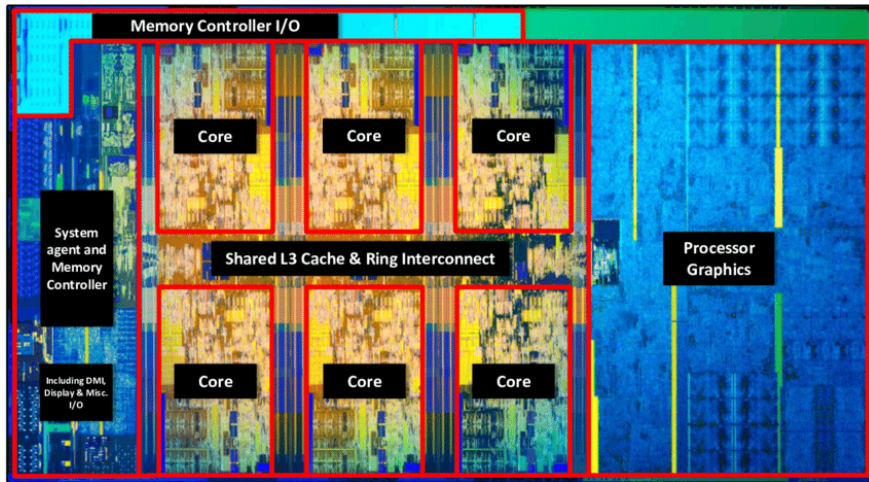


## Exam 1 + 2: Studying (or not studying) Actually Makes a Difference!



- Many students moved between the A-range and B-range
- Many students moved between the C-range and D-range


# Shared-memory Multiprocessors (SMP)



## Intel Coffelake Processor

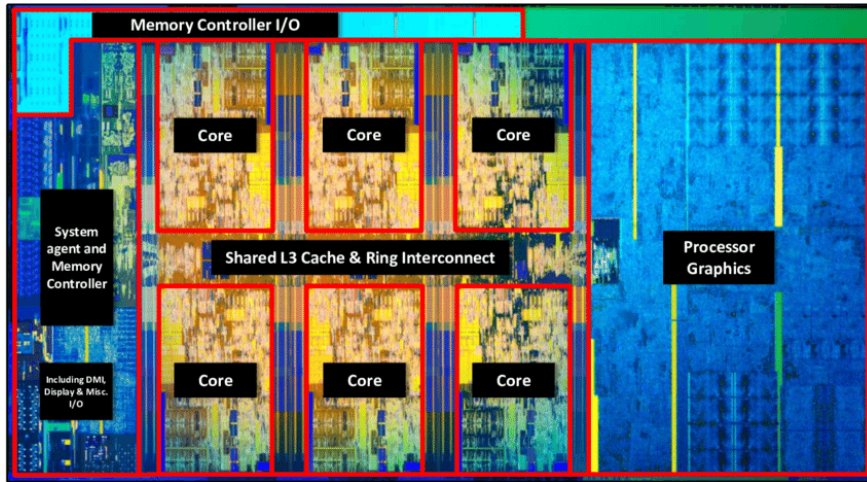
6 cores + graphics co-processor,  
Shared LLC

- Moore's Law: transistor counts kept increasing
- Pipelining / OoO was hitting the wall... what should we do?
  - Multicore/multiprocessors!
- Shared-memory multiprocessors (SMP): a single physical address space across all processors




## Conceptual diagram

# Shared-memory Multiprocessors (SMP)



## Intel Coffelake Processor

6 cores + graphics co-processor,  
Shared LLC




## Conceptual diagram

Complexities of an SMP

- **Cache coherence:** defines the ordering of writes to a single address location
- **Memory consistency:** defines the ordering of reads and writes to all memory location
- **Synchronization:** allowing only one processor to access data at a time

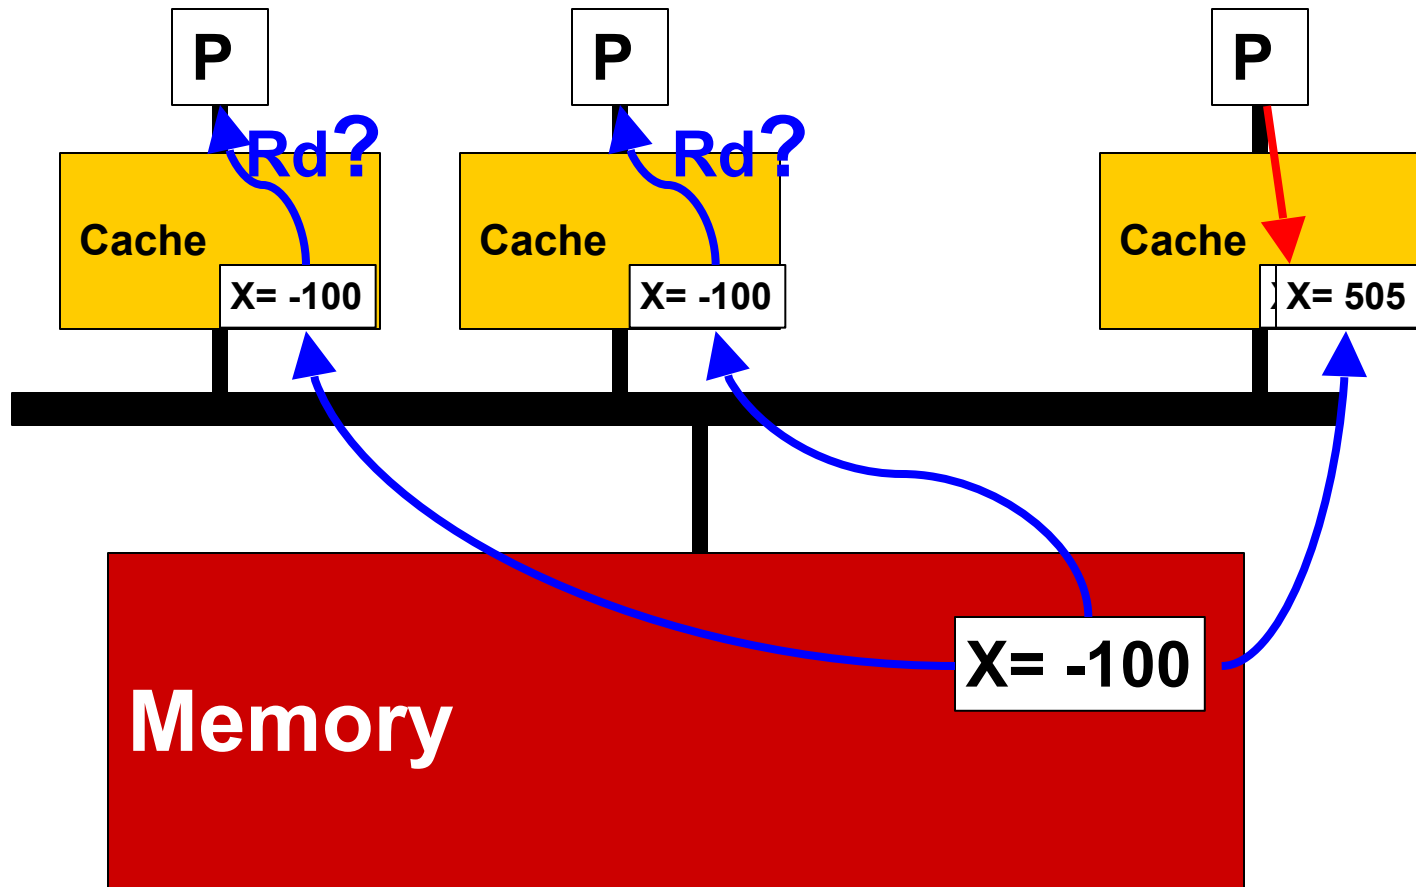
# Cache Coherence

- If P writes to a location X and reads X, if there were no other processors written to X in between, P should read the value it wrote.
- If P1 writes to a location X and P2 reads X, if there were no other writes to X in between and if the two were sufficiently separated in time, P2 should read the value P1 wrote.
- Two writes to the same location by any two processors are seen in the same order by all processors (writes are serializes).
- Seem obvious?

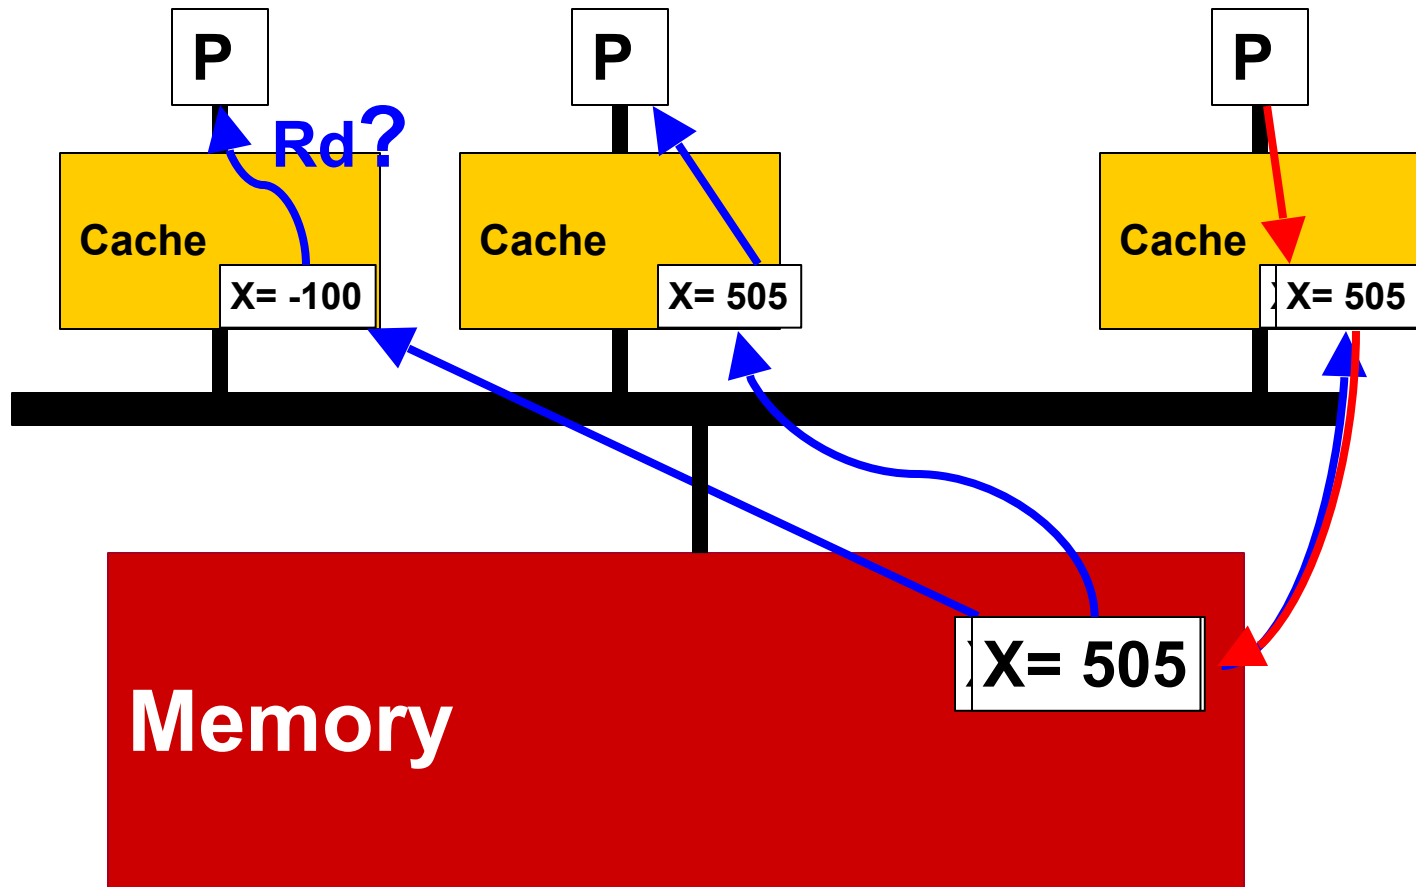




# Example (Writeback Cache)



# Example (Write-through Cache)



# Defining Coherence

- An MP is coherent if the results of any execution of a program can be reconstructed by a hypothetical *serial order*

## Implicit definition of coherence

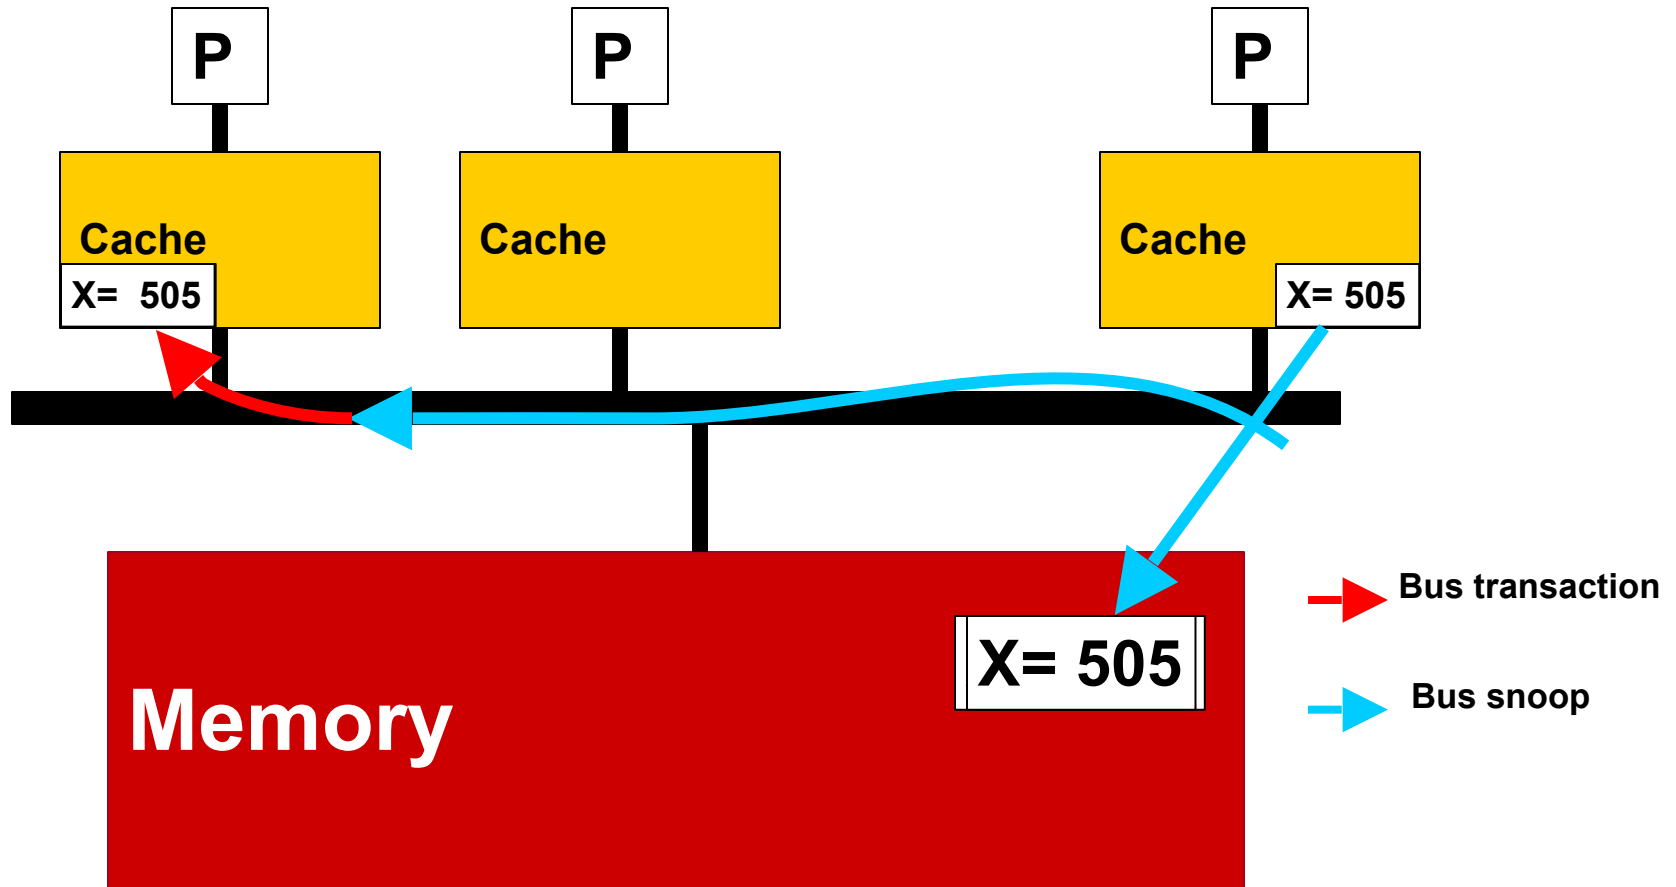
- Write propagation
  - Writes are visible to other processes
- Write serialization
  - All writes to the same location are seen in the same order by all processes (to “all” locations called write atomicity)
  - E.g.,  $w_1$  followed by  $w_2$  seen by a read from  $P_1$ , will be seen in the same order by all reads by other processors  $P_i$

# Bus Snooping based on Write-Through Cache

- All the writes will be shown as a transaction on the shared bus to memory
- Two protocols
  - Update-based Protocol
  - Invalidation-based Protocol

# Bus Snooping

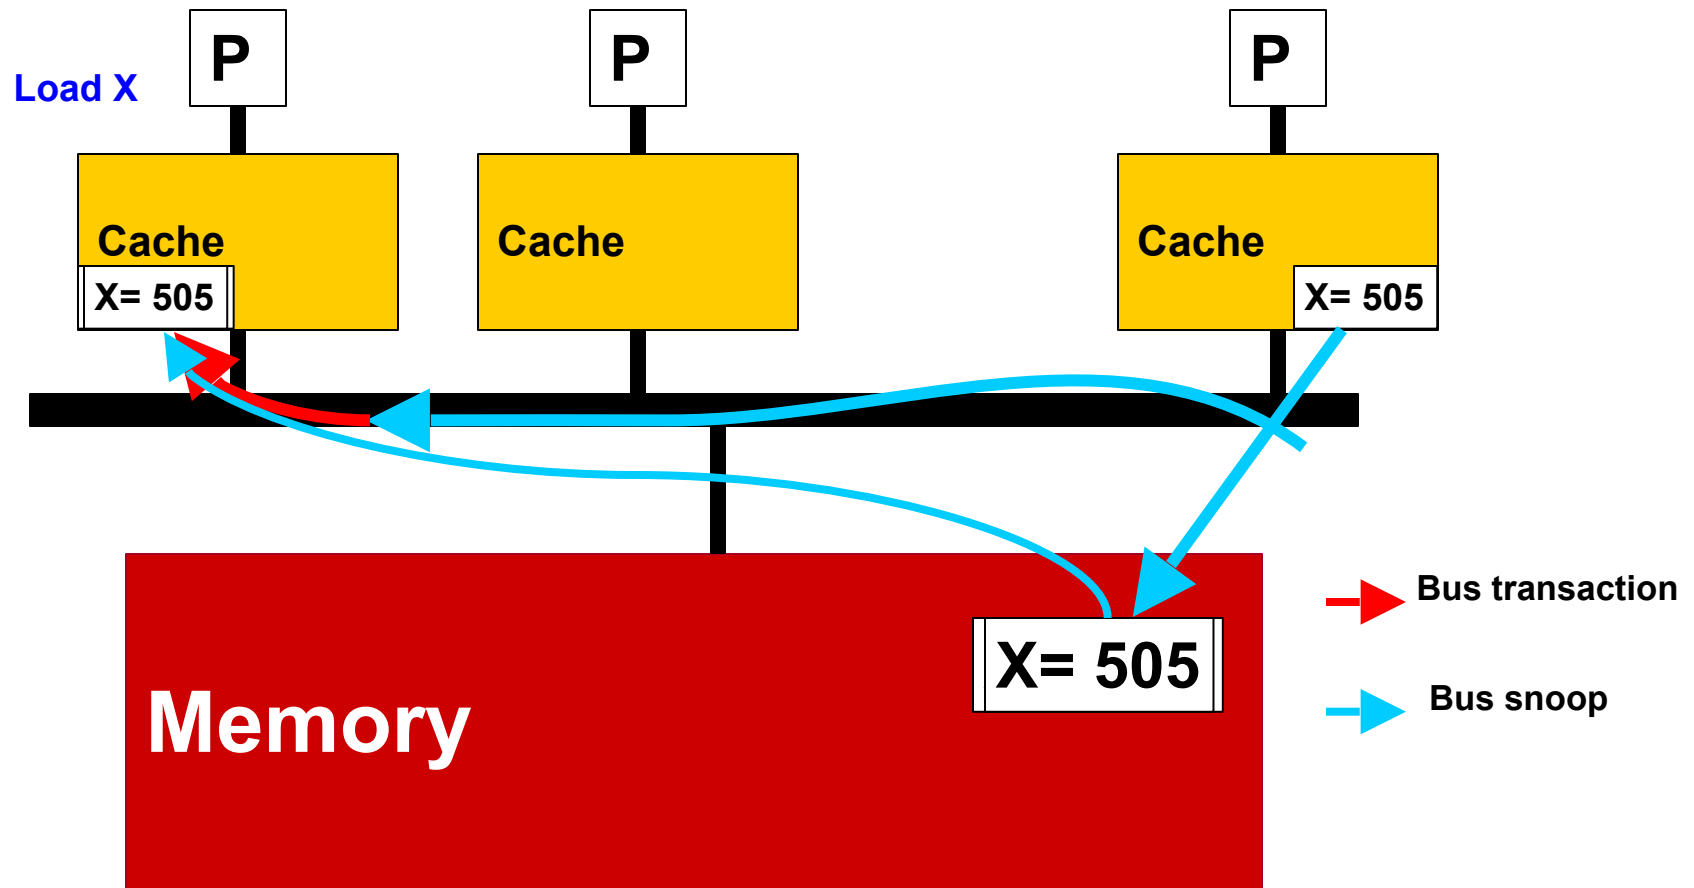
(Update-based Protocol on Write-Through cache)



- Each processor's cache controller constantly snoops on the bus
- Update local copies upon snoop hit


# Bus Snooping

(Invalidation-based Protocol on Write-Through cache)

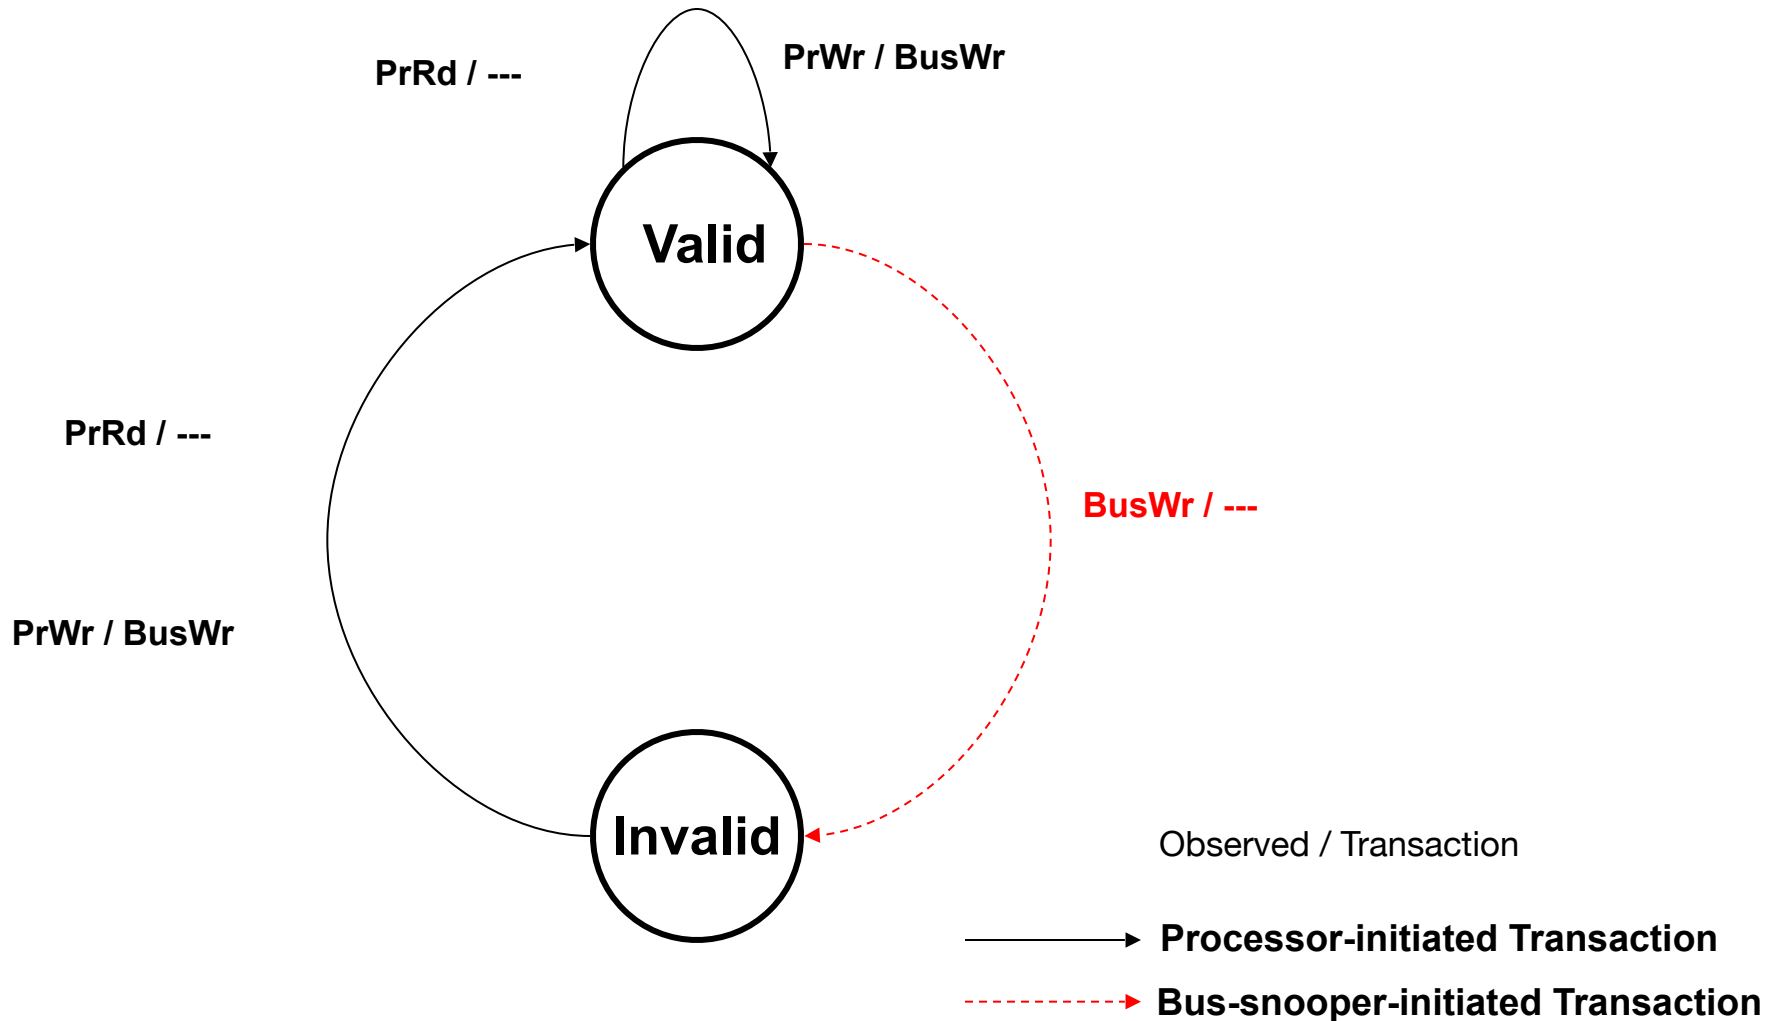


- Each processor's cache controller constantly snoops on the bus
- Invalidate local copies upon snoop hit

# A Simple Invalidation-based Coherence Protocol for a WT, No Write-Allocate Cache




# A Simple Invalidation-based Coherence Protocol for a WT, Write-Allocate Cache





# Bus Snooping

(Invalidation-based Protocol on Write-Through cache)



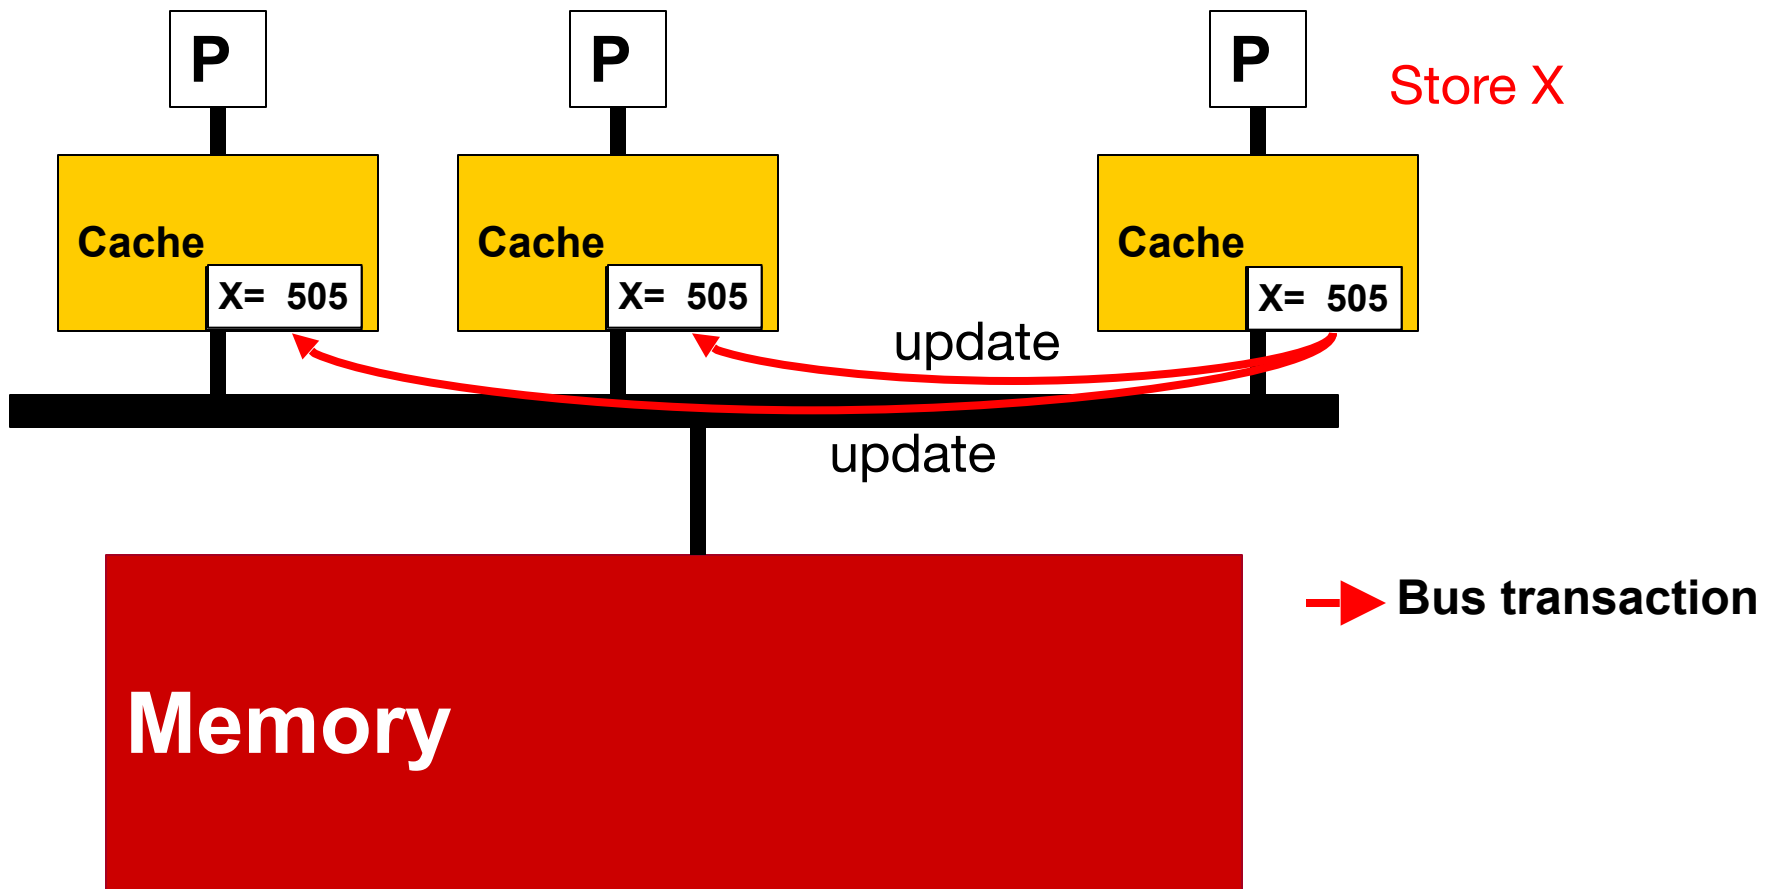
- Each processor's cache controller constantly snoops on the bus
- Invalidate local copies upon snoop hit

# How about Writeback Cache?

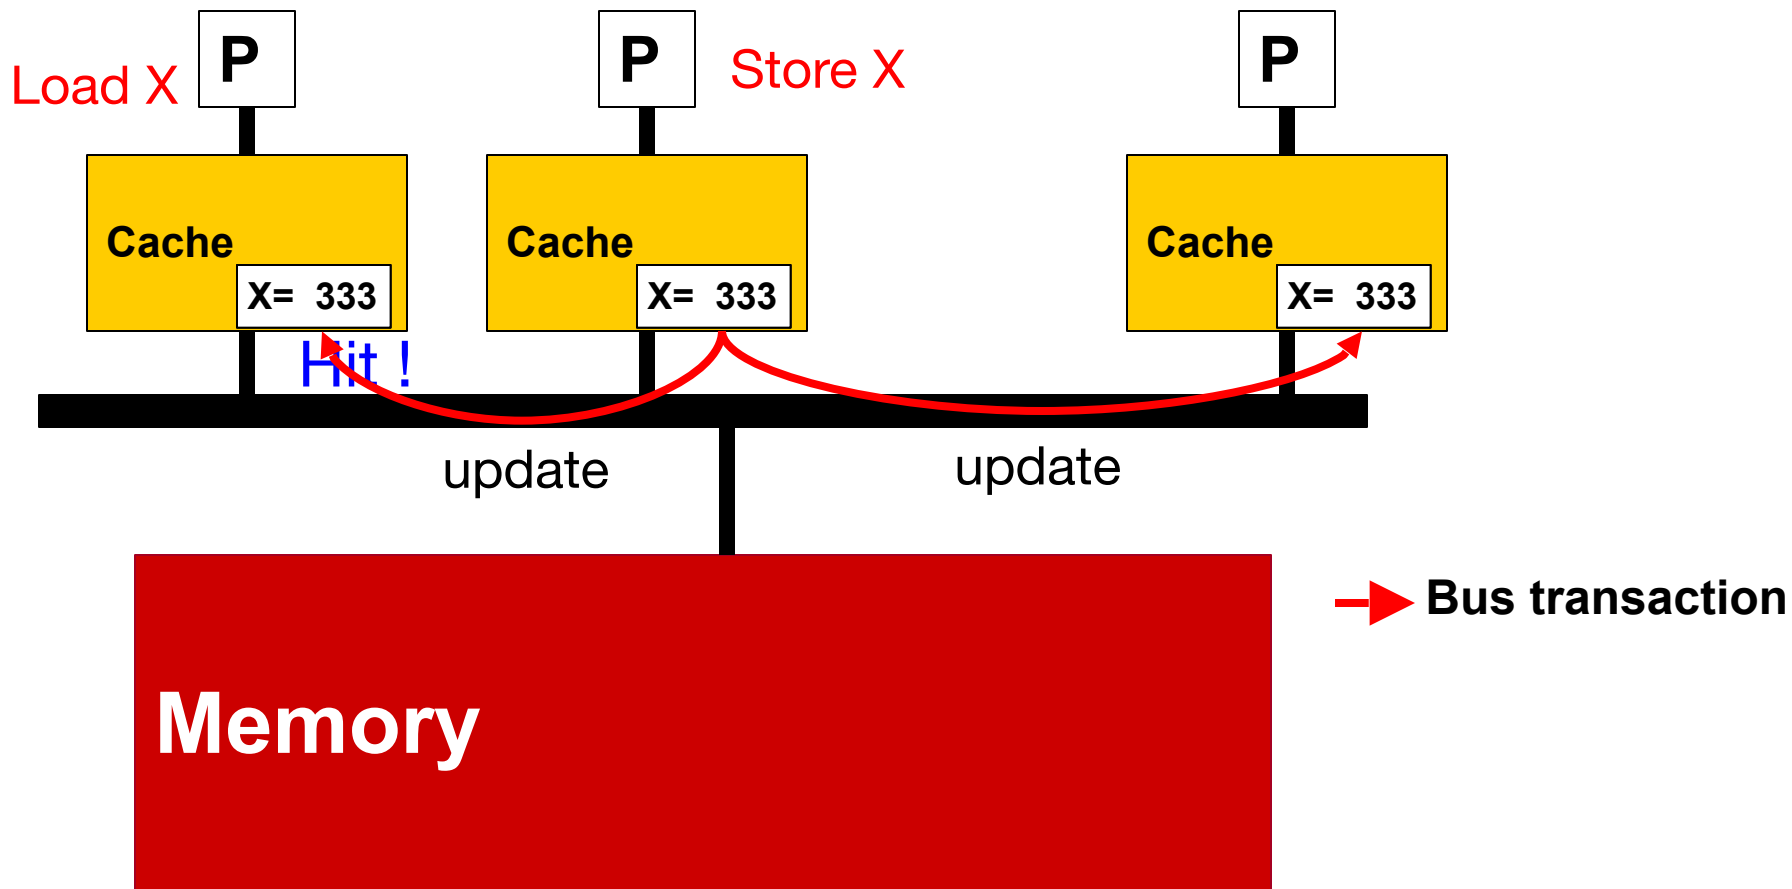
- Write-through wastes a lot of bus bandwidth (every update creates a traffic)
- WB cache to reduce bandwidth requirement
- The majority of local writes are hidden behind the processor nodes
- How to snoop?
- Write ordering

# Cache Coherence Protocols for WB caches


- A cache has an *exclusive* copy of a line if
  - It is the only cache having a valid copy
  - Memory may or may not have it
- Modified (dirty) cache line
  - The cache having the line is the *owner* of the line, because it must supply the block




- Update data for all processor nodes who share the same data
- For a processor node keeps updating the memory location, a lot of traffic will be incurred



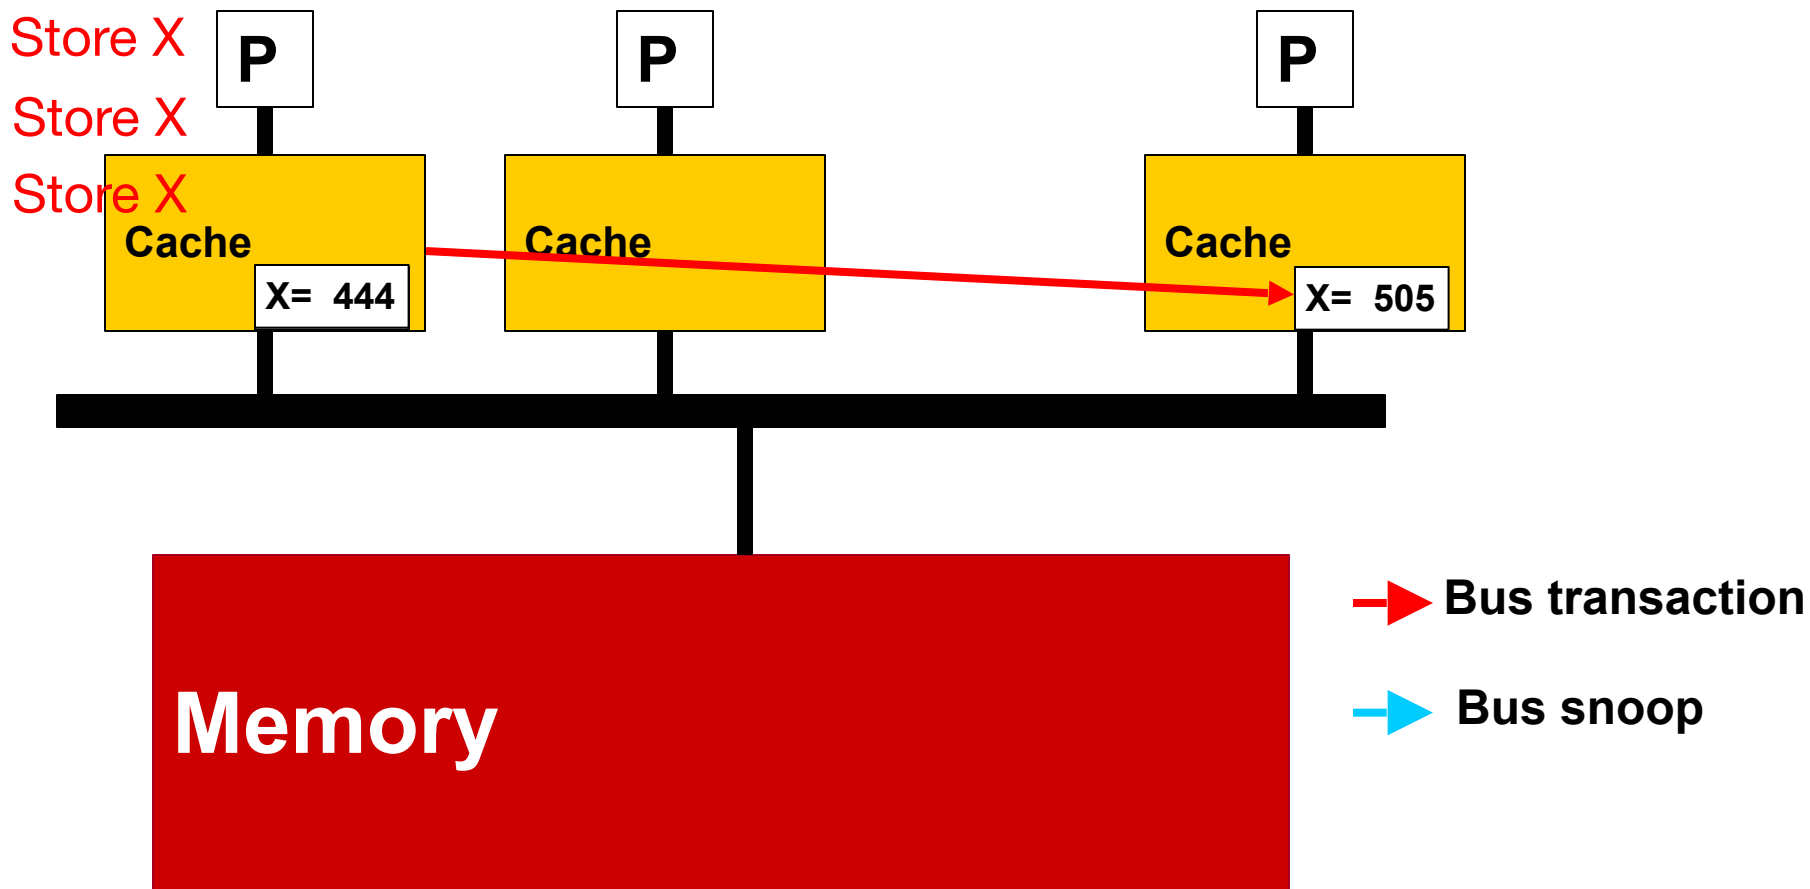
- Update data for all processor nodes who share the same data
- For a processor node keeps updating the memory location, **a lot of traffic** will be incurred



- Invalidate the data copies for the sharing processor nodes
- Reduced traffic when a processor node keeps updating the same memory location



- Invalidate the data copies for the sharing processor nodes
- Reduced traffic when a processor node keeps updating the same memory location



- Invalidate the data copies for the sharing processor nodes
- Reduced traffic when a processor node keeps updating the same memory location



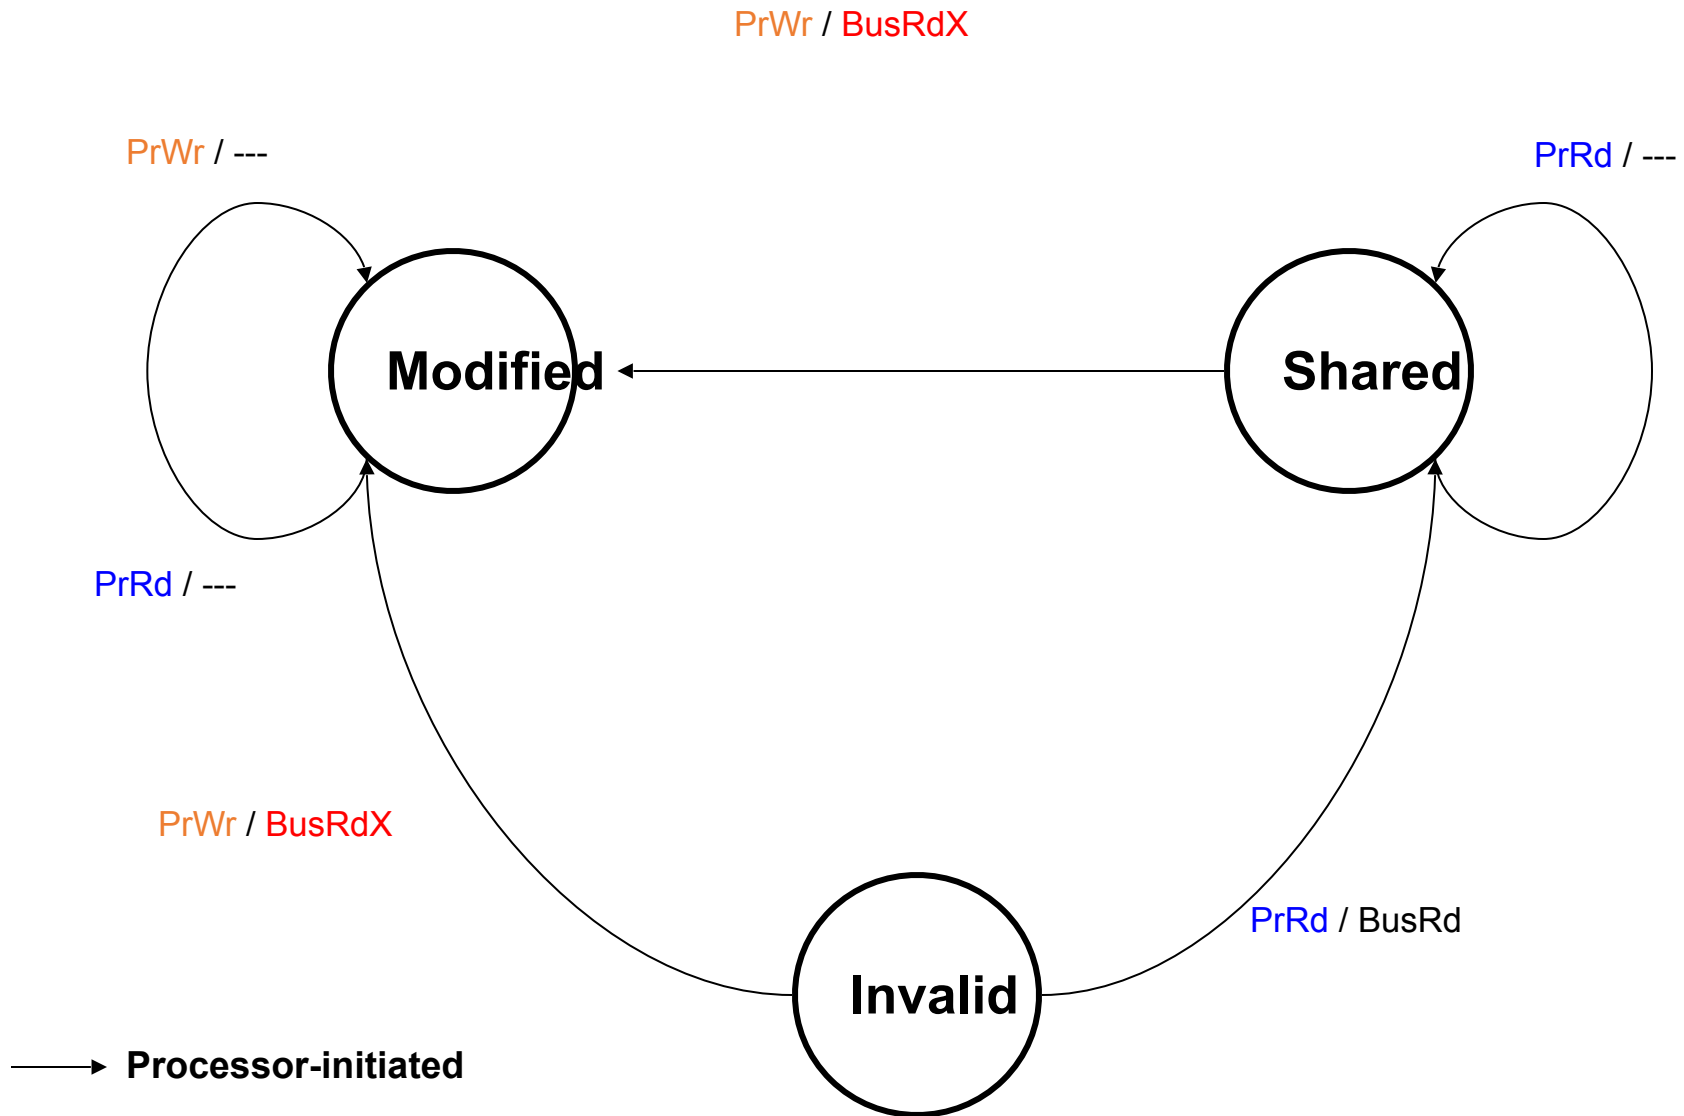
# MSI Writeback Invalidation Protocol

- **Modified**
    - Dirty
    - Only this cache has a valid copy
  - **Shared**
    - Memory is consistent
    - One or more caches have a valid copy
  - **Invalid**
- 
- Writeback protocol: A cache line can be written multiple times before the memory is updated.

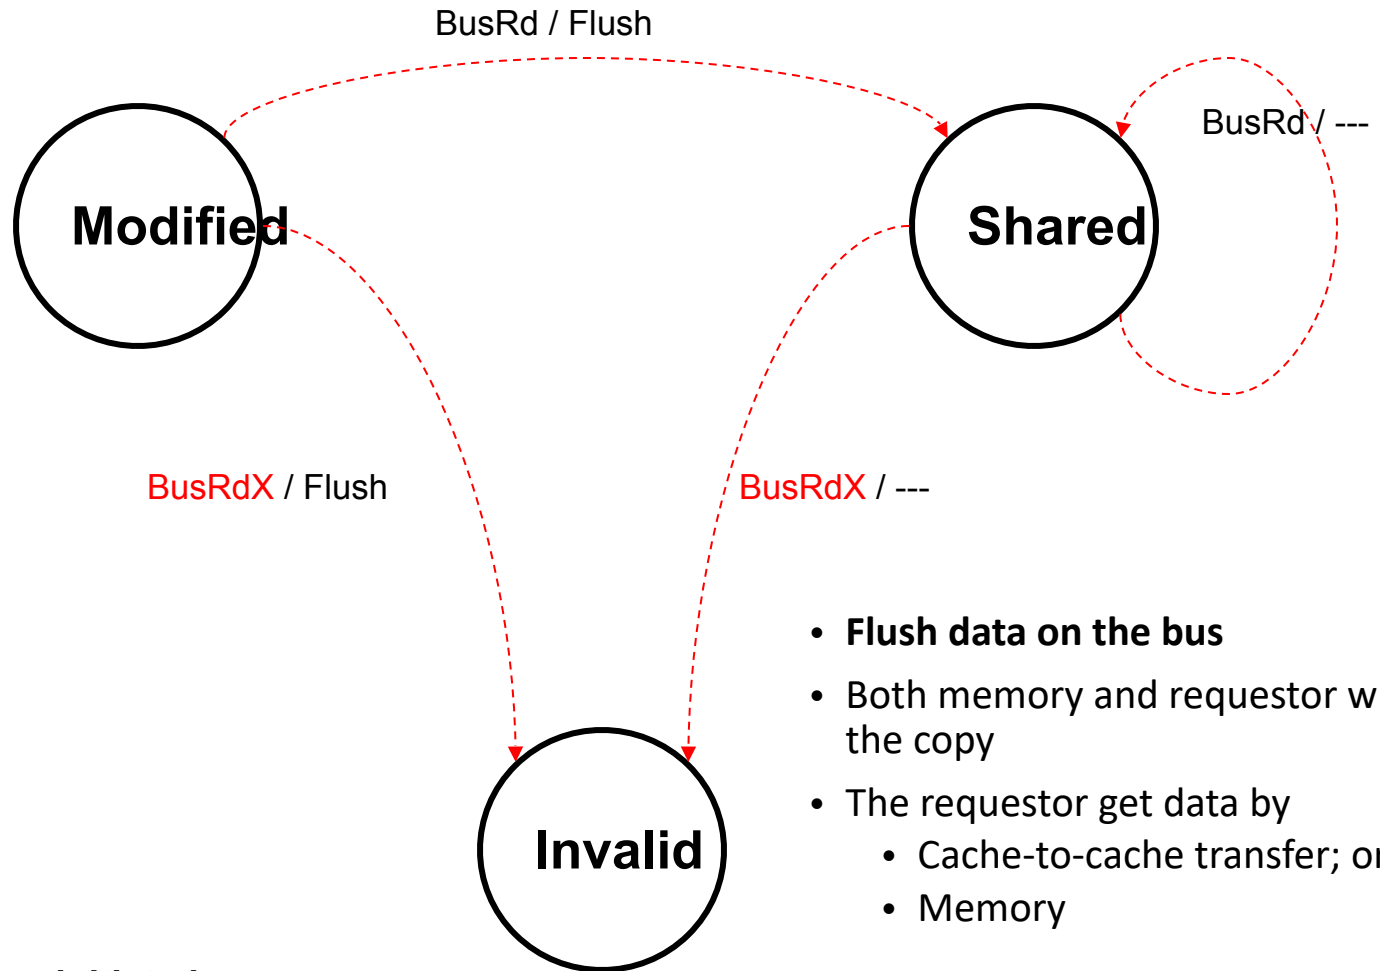
# MSI Writeback Invalidation Protocol

- Two types of request from the **processor**
  - PrRd
  - PrWr
- Three types of **bus transactions** post by cache controller
  - **BusRd**
    - PrRd misses the cache
    - Memory or another cache supplies the line
  - **BusRd eXclusive (Read-to-own)**
    - PrWr is issued to a line which is **not** in the **Modified** state
  - **BusWB**
    - Writeback due to replacement
    - Processor does not directly involve in initiating this operation

# MSI Writeback Invalidation Protocol (Processor Request)



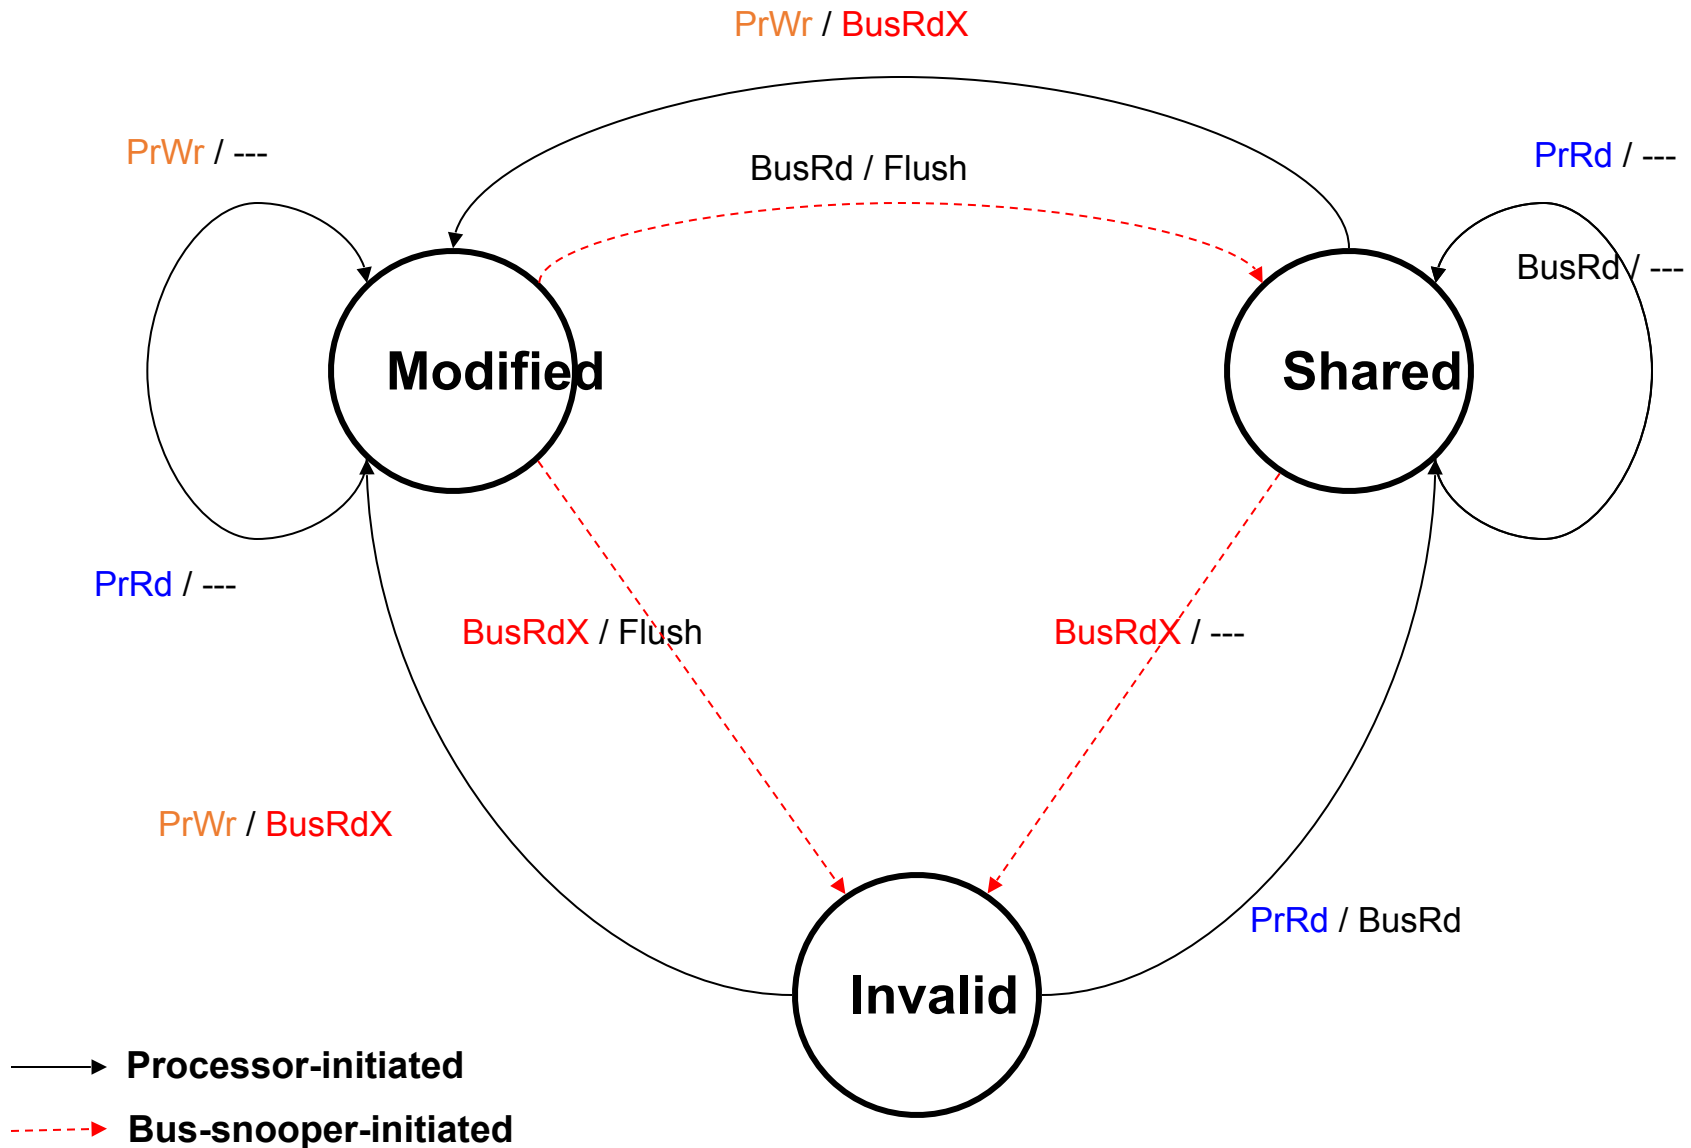
# MSI Writeback Invalidation Protocol (Bus Transaction)




-----> **Bus-snooper-initiated**

- **Flush data on the bus**
- Both memory and requestor will grab the copy
- The requestor get data by
  - Cache-to-cache transfer; or
  - Memory

# MSI Writeback Invalidation Protocol




# MSI Example



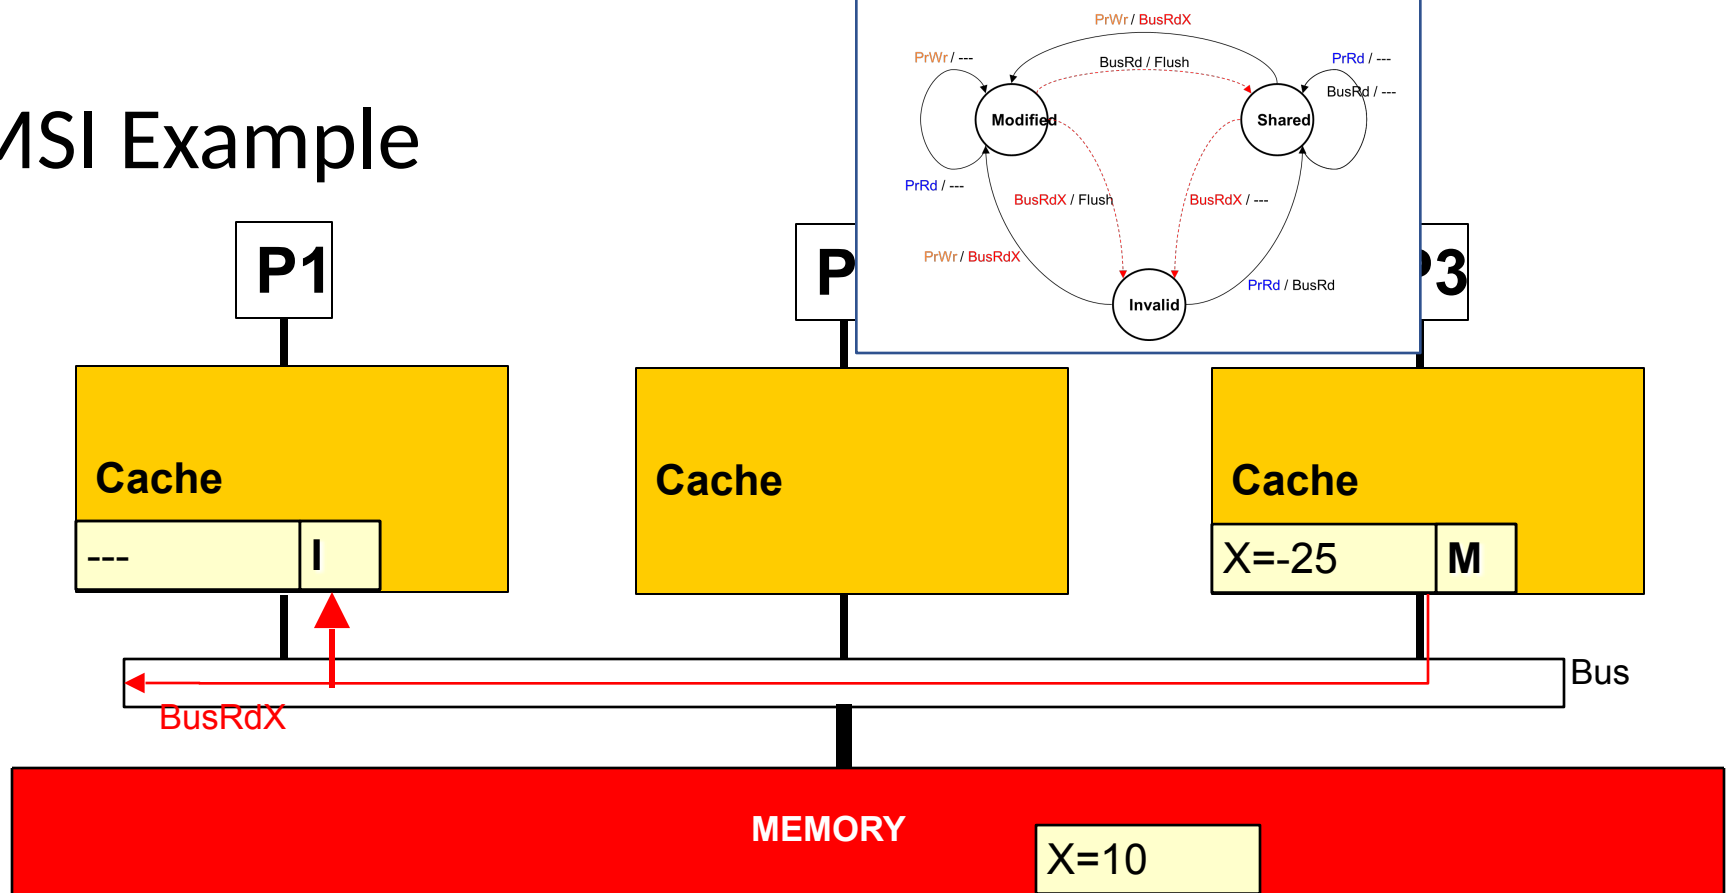
Processor Action	State in P1	State in P2	State in P3	Bus Transaction	Data Supplier
P1 reads X	S	---	---	BusRd	Memory

# MSI Example



Processor Action	State in P1	State in P2	State in P3	Bus Transaction	Data Supplier
P1 reads X	S	---	---	BusRd	Memory
P3 reads X	S	---	S	BusRd	Memory

# MSI Example

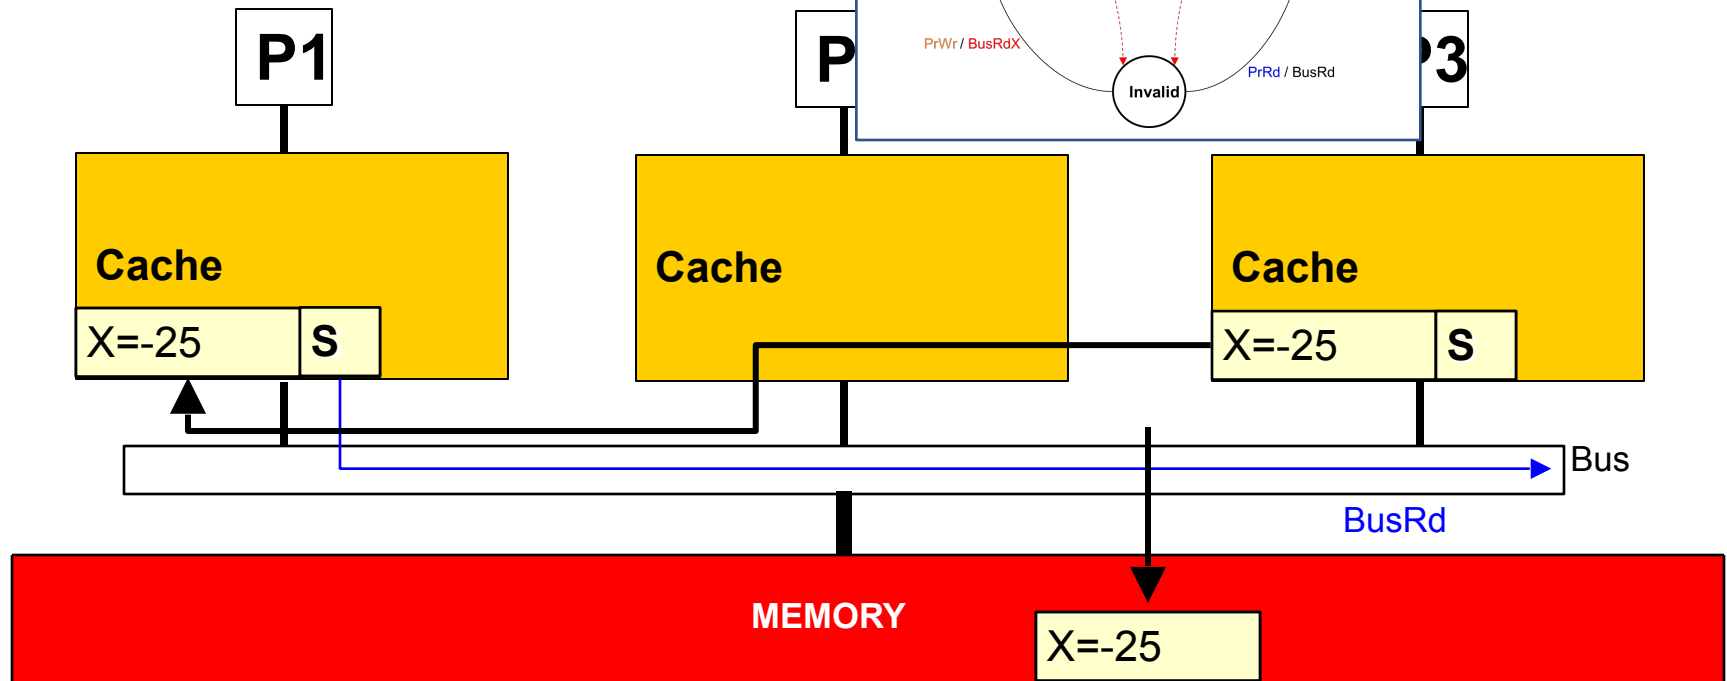


Processor Action	State in P1	State in P2	State in P3	Bus Transaction	Data Supplier
P1 reads X	S	---	---	BusRd	Memory
P3 reads X	S	---	S	BusRd	Memory
P3 writes X	I	---	M	BusRdX	P3 Cache

Does not come from memory if having “BusUpgrade”




# MSI Example



Processor Action	State in P1	State in P2	State in P3	Bus Transaction	Data Supplier
P1 reads X	S	---	---	BusRd	Memory
P3 reads X	S	---	S	BusRd	Memory
P3 writes X	I	---	M	BusRdX	P3 Cache
P1 reads X	S	---	S	BusRd	P3 Cache

# MSI Example



Processor Action	State in P1	State in P2	State in P3	Bus Transaction	Data Supplier
P1 reads X	S	---	---	BusRd	Memory
P3 reads X	S	---	S	BusRd	Memory
P3 writes X	I	---	M	BusRdX	P3 Cache
P1 reads X	S	---	S	BusRd	P3 Cache
P2 reads X	S	S	S	BusRd	Memory

# MSI Summary

- MSI ensures single-writer-multiple-reader
  - Often called the SWMR (pronounced as “swimmer”) invariant
- We can still do better
  - MESI, MOESI, ....
  - I will not go into further detail.