

CSE 431
Computer Architecture
Fall 2022
Introduction to Parallel Computers
Part II

**(based partially on the material in Introduction to Parallel Computing
by A. Grama, A. Gupta, G. Karypis, V. Kumar
Second Edition, Addison Wesley**

Kiwan Maeng

Slides adapted from Mahmut Taylan Kandemir (www.cse.psu.edu/~kandemir)

Project 2: Quick Run-Through

- ❑ `setSizeOffsetAndMaskFields(acache, size, assoc, blocksize)`
 - ❓ You must set the following variables
 - ❓ `acache->numways`
 - ❓ `acache->blocksize`
 - ❓ `acache->numsets`
 - ❓ `acache->numBitsForBlockOffset`
 - ❓ `acache->numBitsForIndex`
 - ❓ `acache->VAImask`
 - ❓ `acache->VATmask`
 - ❓ Try printing these out to see if you are getting reasonable results before doing anything!!!

Project 2: Quick Run-Through

- ❑ `getindex(acache, address)`
 - ❑ Get the index bits from the address
- ❑ `gettag(acache, address)`
 - ❑ Get the tag bits from the address

Project 2: Quick Run-Through

- ❑ writeback(acache, index, oldestway)
 - ❑ Write back the oldestway of the index to the next level (you only need to care about the data block)
 - ❑ 0. look at csim.h for the structure of acache
 - ❑ 1. select the block to write back
 - ❑ 2. write the block out to the next level
 - Calculate the address from the cache parameters
 - Get the value from the word in a block (you can only do it word by word)
 - Call Store(acache->nextcache, address, value)
 - Note that for multiword block cache, you have to do this multiple times!!!!

Project 2: Quick Run-Through

- ❑ `fill(acache, index, oldestway, address)`
 - ❓ Read the block from the next level and put it in the current block (you only need to care about the data block)
 - ❓ 0. look at `csim.h` for the structure of `acache`
 - ❓ 1. select the block to write into
 - ❓ 2. read the block from the next level and put it in the block
 - Calculate the address you want to read from ****this can be a bit tricky****
 - `Load(acache->nextcache, address)` to read a word from the address
 - Put it in the block
 - Note that for multiword block cache, you have to do this multiple times!!!!
 - ❓ If you use a similar logic with the `writeback()` to calculate the address, it also kinda works... (because the tag is already set)

Project 2: Quick Run-Through

- ❑ This is a 64-bit machine!! So things are different than what we learned.
 - ❑ Byte offset is 3
 - ❑ Word size is 8-byte
- ❑ THIS IS MUCH HARDER!!!!!! START NOW!!!!!!
- ❑ Use Canvas discussion as much as possible (although it is bad..)

Dichotomy of Parallel Computing Platforms

- ❑ An explicitly parallel program must specify concurrency and interaction between concurrent subtasks.
- ❑ The former is sometimes also referred to as the control structure and the latter as the communication model.

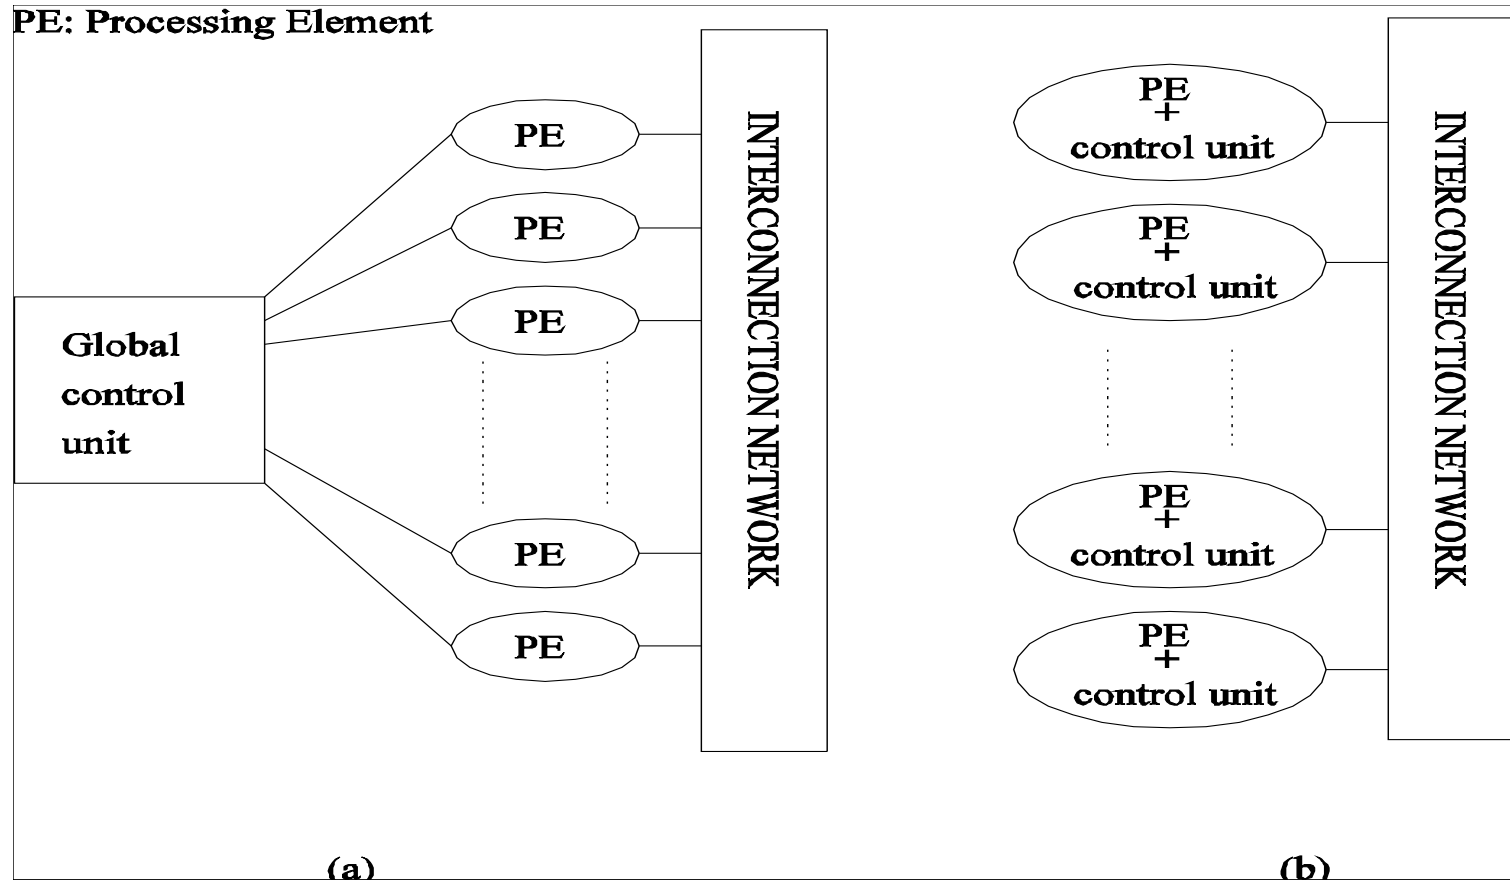
Control Structure of Parallel Programs

- ❑ Parallelism can be expressed at various levels of granularity - from instruction level to processes.
- ❑ Between these extremes exist a range of models, along with corresponding architectural support.

Control Structure of Parallel Programs

- ❑ Processing units in parallel computers either operate under the centralized control of a single control unit or work independently.
- ❑ If there is a single control unit that dispatches the same instruction to various processors (that work on different data), the model is referred to as **Single Instruction, Multiple Data** (SIMD).
- ❑ If each processor has its own control control unit, each processor can execute different instructions on different data items. This model is called **Multiple Instruction, Multiple Data** (MIMD).

SIMD and MIMD Processors

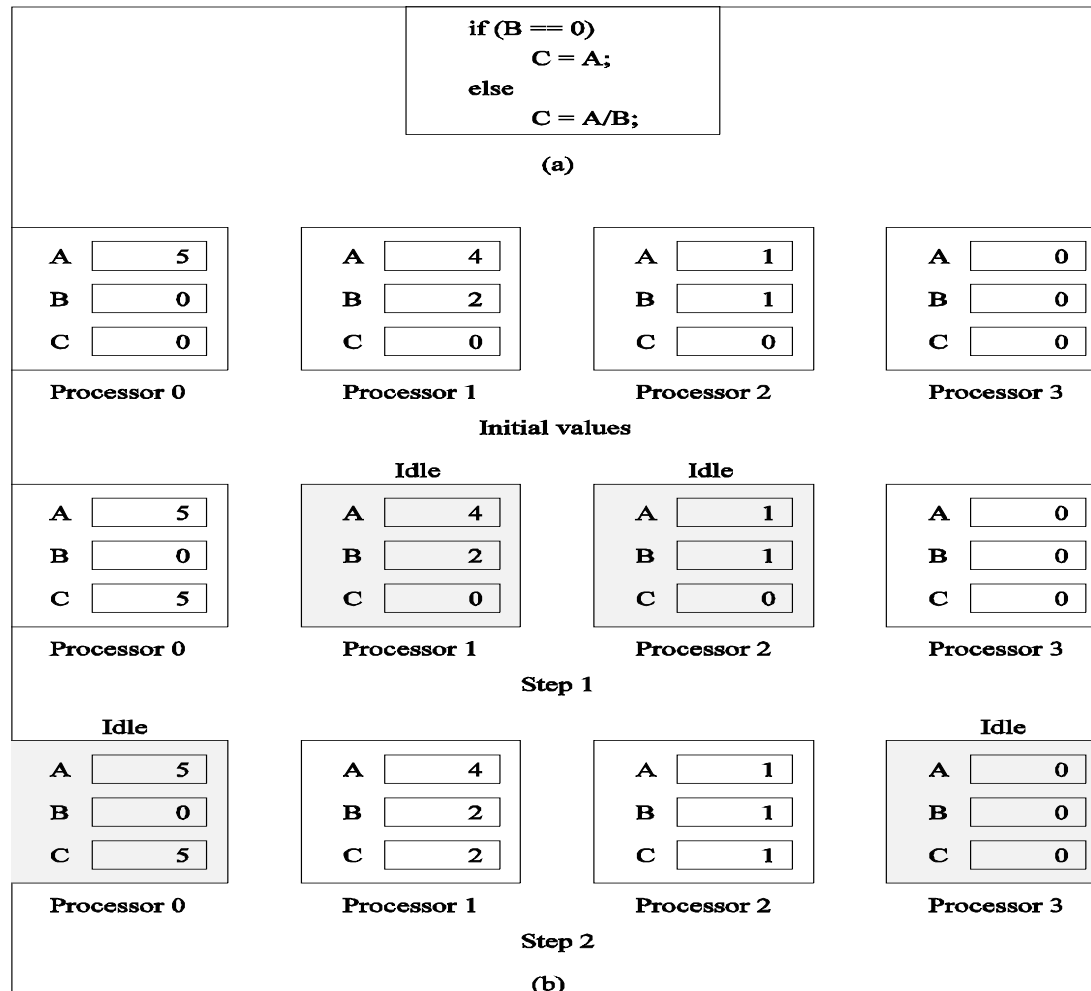


A typical SIMD architecture (a) and a typical MIMD architecture (b).

SIMD Processors

- ❑ Some of the earliest parallel computers such as the Illiac IV, MPP, DAP, CM-2, and MasPar MP-1 belonged to this class of machines.
- ❑ Variants of this concept have found use in co-processing units such as the **MMX** units in Intel processors and **DSP** chips such as the Sharc.
- ❑ SIMD relies on the regular structure of computations (such as those in image processing).
- ❑ It is often necessary to selectively turn off operations on certain data items. For this reason, most SIMD programming paradigms allow for an "activity mask", which determines if a processor should participate in a computation or not.

Conditional Execution in SIMD Processors



Executing a conditional statement on an SIMD computer with four processors: (a) the conditional statement; (b) the execution of the statement in two steps.

MIMD Processors

- ❑ In contrast to SIMD processors, MIMD processors can execute different programs on different processors.
- ❑ A variant of this, called single program multiple data streams (SPMD) executes the same program on different processors.
- ❑ It is easy to see that SPMD and MIMD are closely related in terms of programming flexibility and underlying architectural support.
- ❑ Examples of such platforms include current generation multiprocessor PCs and workstation clusters.

SIMD-MIMD Comparison

- ❑ SIMD computers require less hardware than MIMD computers (due to single control unit).
- ❑ However, since SIMD processors are specially designed, they tend to be expensive and have long design cycles.
- ❑ Not all applications are naturally suited to SIMD processors (image processing applications are).
- ❑ In contrast, platforms supporting the SPMD paradigm can be built from inexpensive off-the-shelf components with relatively little effort in a short amount of time.

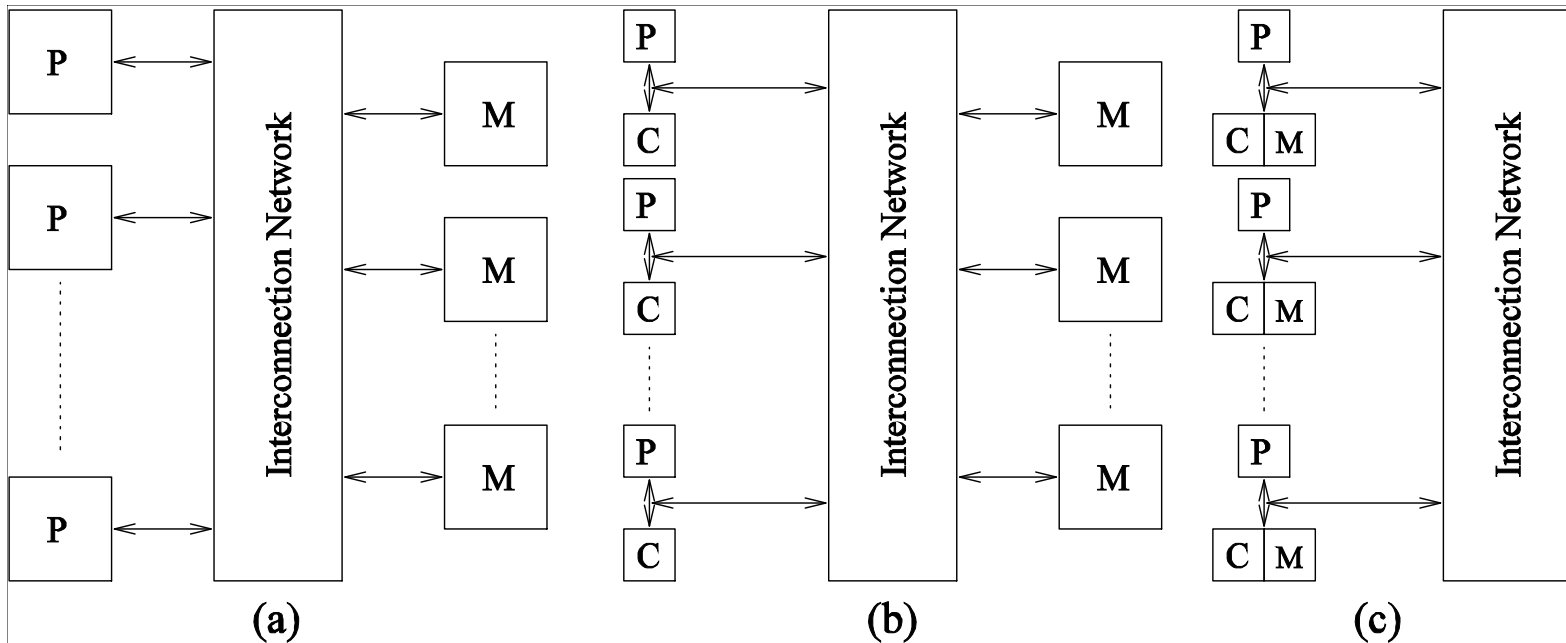
Communication Model of Parallel Platforms

- ❑ There are two primary forms of data exchange between parallel tasks - accessing a shared data space and exchanging messages.
- ❑ Platforms that provide a shared data space are called shared-address-space machines or multiprocessors.
- ❑ Platforms that support messaging are also called message-passing platforms or multi-computers.

Shared-Address-Space Platforms

- ❑ Part (or all) of the memory is accessible to all processors.
- ❑ Processors interact by modifying data objects stored in this shared-address-space.
- ❑ If the time taken by a processor to access any memory word in the system global or local is identical, the platform is classified as a uniform memory access (UMA), else, a non-uniform memory access (NUMA) machine.

NUMA and UMA Shared-Address-Space Platforms



Typical shared-address-space architectures: (a) Uniform-memory access shared-address-space computer; (b) Uniform-memory access shared-address-space computer with caches and memories; (c) Non-uniform-memory access shared-address-space computer with local memory only.

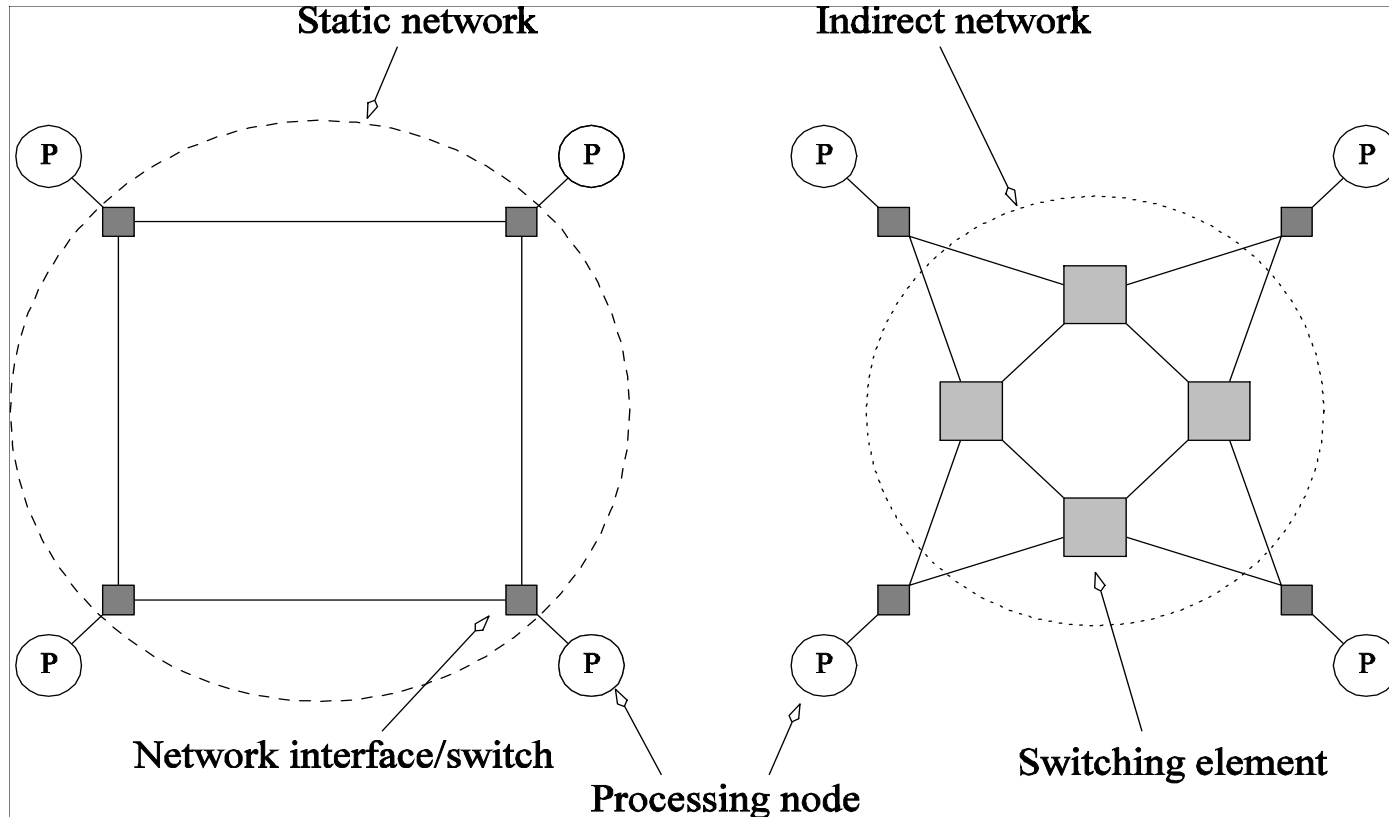
Message-Passing Platforms

- ❑ These platforms comprise of a set of processors and their own (exclusive) memory.
- ❑ Instances of such a view come naturally from clustered workstations and non-shared-address-space multicomputers.
- ❑ These platforms are programmed using (variants of) send and receive primitives.
- ❑ Libraries such as MPI provide such primitives.

Interconnection Networks for Parallel Computers

- ❑ Interconnection networks carry data between processors and to memory.
- ❑ Interconnects are made of switches and links (wires, fiber).
- ❑ Interconnects are classified as **Static** or **Dynamic**.
- ❑ Static networks consist of point-to-point communication links among processing nodes and are also referred to as **Direct** networks.
- ❑ Dynamic networks are built using switches and communication links. Dynamic networks are also referred to as **Indirect** networks.

Static and Dynamic Interconnection Networks



Classification of interconnection networks: (a) a static network; and (b) a dynamic network.

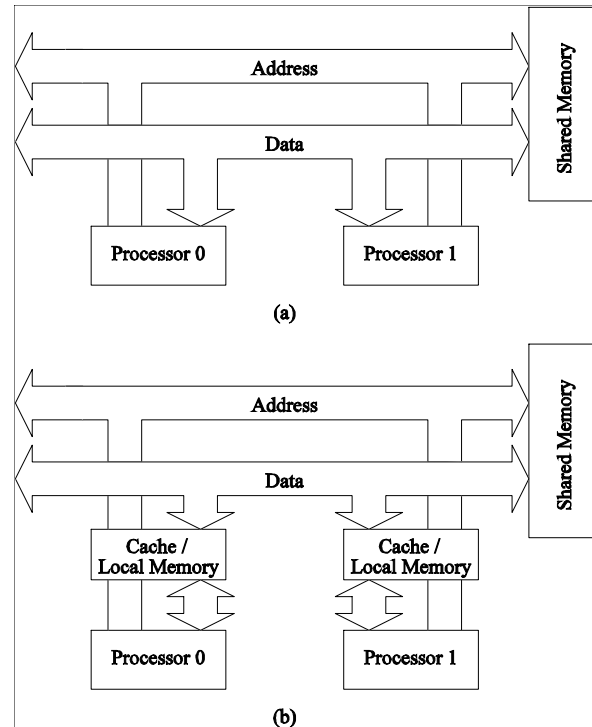
Network Topologies

- ❑ A variety of network topologies have been proposed and implemented.
- ❑ These topologies tradeoff performance for cost.
- ❑ Commercial machines often implement hybrids of multiple topologies for reasons of packaging, cost, and available components.

Network Topologies: Buses

- ❑ Some of the simplest and earliest parallel machines used buses.
- ❑ All processors access a common bus for exchanging data.
- ❑ The distance between any two nodes is $O(1)$ in a bus. The bus also provides a convenient broadcast media.
- ❑ However, the bandwidth of the shared bus is a major bottleneck.
- ❑ Typical bus based machines are limited to dozens of nodes. Sun Enterprise servers and Intel Pentium based shared-bus multiprocessors are examples of such architectures.

Network Topologies: Buses




Bus-based interconnects (a) with no local caches; (b) with local memory/caches.

Since much of the data accessed by processors is local to the processor, a local memory can improve the performance of bus-based machines.

Network Topologies: Crossbars

A crossbar network uses an $p \times m$ grid of switches to connect p inputs to m outputs in a non-blocking manner.

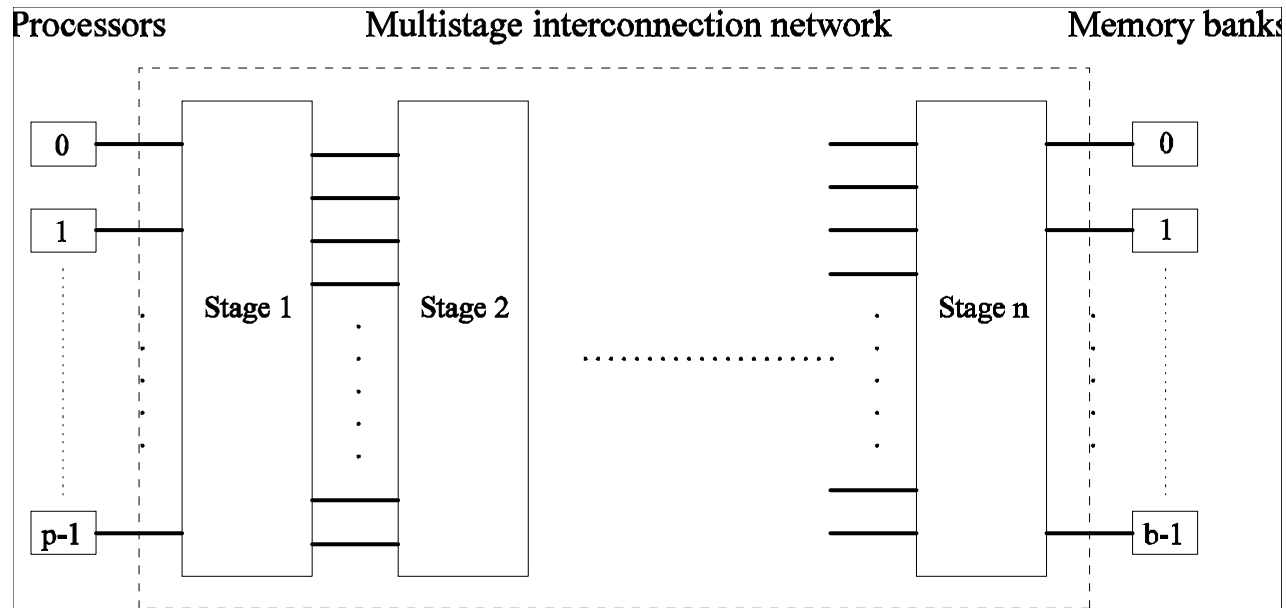


A completely non-blocking crossbar network connecting p processors to b memory banks.

Network Topologies: Multistage Networks

- ❑ Crossbars have excellent performance scalability but poor cost scalability.
- ❑ Buses have excellent cost scalability, but poor performance scalability.
- ❑ Multistage interconnects strike a *compromise* between these extremes.

Network Topologies: Multistage Networks




The schematic of a typical multistage interconnection network.

Network Topologies: Multistage Omega Network

- ❑ One of the most commonly used multistage interconnects is the Omega network.
- ❑ This network consists of $\log p$ stages, where p is the number of inputs/outputs.
- ❑ At each stage, input i is connected to output j if:


$$j = \begin{cases} 2i, & 0 \leq i \leq p/2 - 1 \\ 2i + 1 - p, & p/2 \leq i \leq p - 1 \end{cases}$$

Network Topologies: Completely Connected and Star Connected Networks



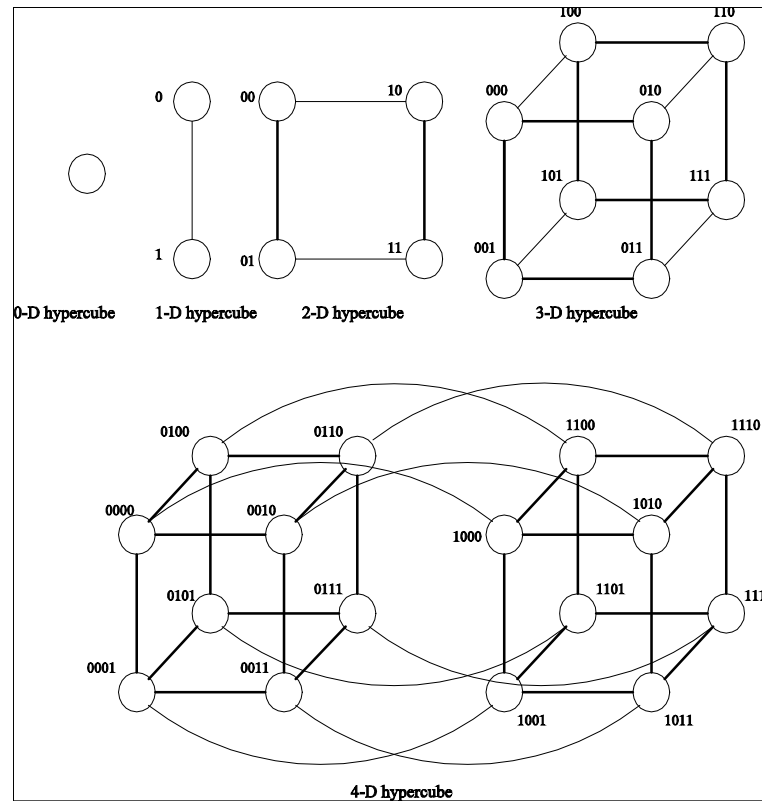
- (a) A completely-connected network of eight nodes;
(b) a star connected network of nine nodes.

Network Topologies: Meshes



Two and three dimensional meshes: (a) 2-D mesh with no wraparound; (b) 2-D mesh with wraparound link (2-D torus); and (c) a 3-D mesh with no wraparound.

Network Topologies: Hypercubes




Construction of hypercubes from hypercubes of lower dimension.

Network Topologies: Properties of Hypercubes


- ❑ The distance between any two nodes is at most $\log p$.
- ❑ Each node has $\log p$ neighbors.
- ❑ The distance between two nodes is given by the number of bit positions at which the two nodes differ.

Network Topologies: Tree-Based Networks



Complete binary tree networks: (a) a static tree network; and (b) a dynamic tree network.

Network Topologies: Fat Trees



A fat tree network of 16 processing nodes.

Manycores with On-Chip Networks

- An example 5x5 manycore system connected through 2D mesh on-chip network – also known as **Network-on-Chip** (NoC)
- A node contains a core, a private L1, private LLC/shared LLC bank, and network routing circuitry
- Data movements in NoC are caused by data accesses
- Directory accesses are not shown here

