




CSE 511: Operating Systems Design


Lectures 5,6

Page Tables
Kernel Booting

Traditional 48-bit paging (4 levels)



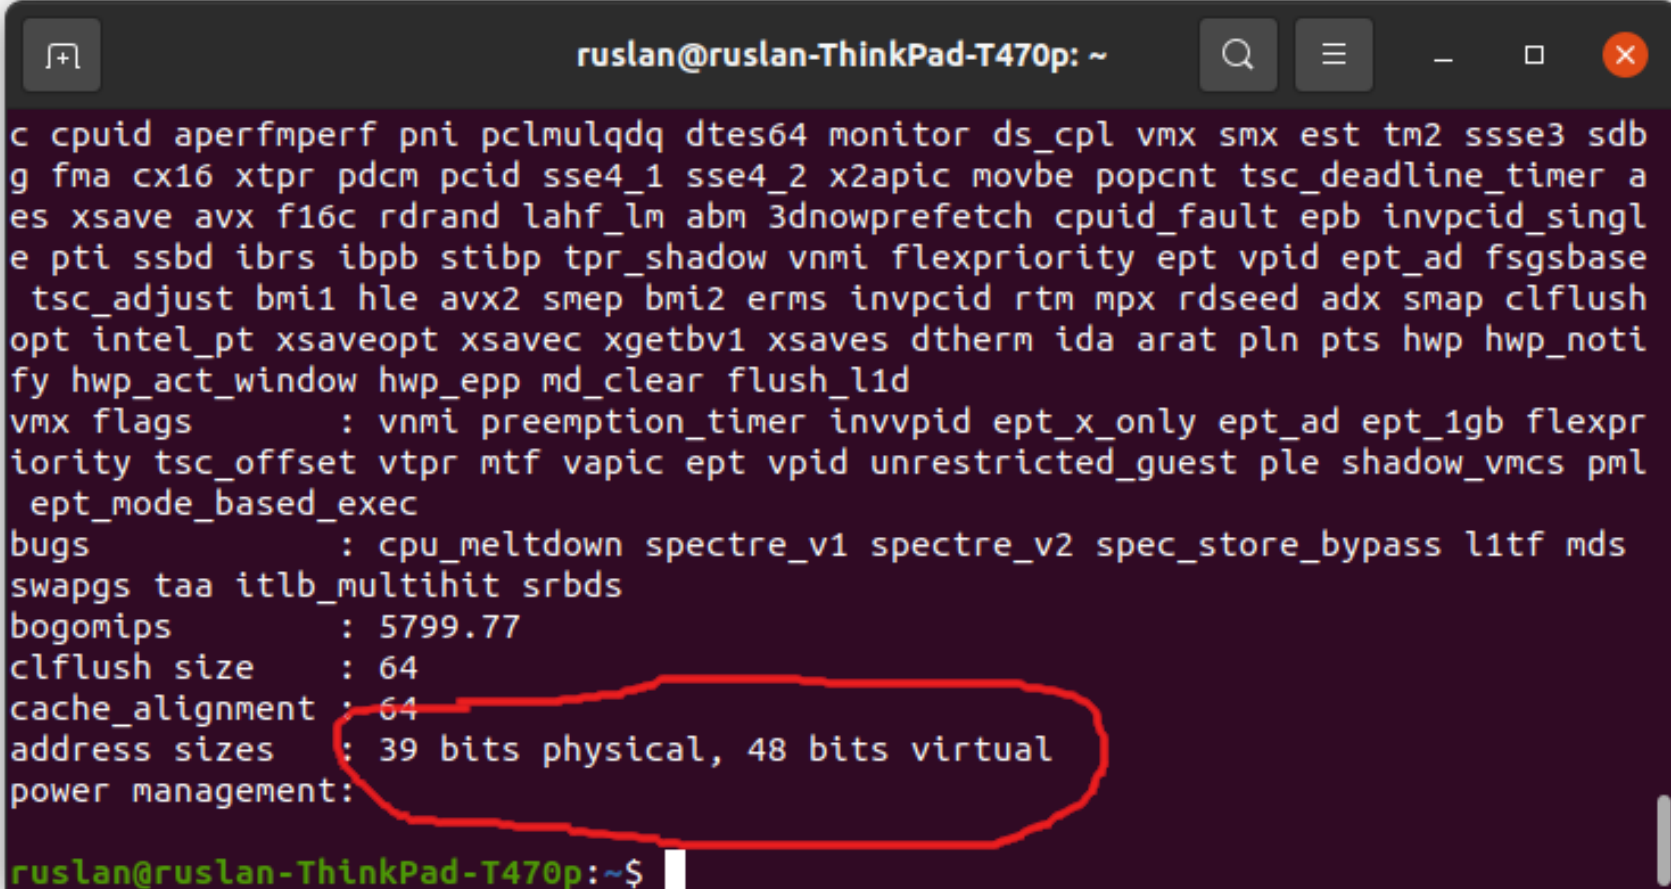
Extension: 57-bit paging (5 levels)



* The picture is taken from https://en.wikipedia.org/wiki/Intel_5-level_paging

Page Table

```
cat /proc/cpuinfo
```



```
ruslan@ruslan-ThinkPad-T470p: ~  
c cpuid aperfperf pni pclmulqdq dtes64 monitor ds_cpl vmx smx est tm2 ssse3 sdb  
g fma cx16 xtpr pdcm pcid sse4_1 sse4_2 x2apic movbe popcnt tsc_deadline_timer a  
es xsave avx f16c rdrand lahf_lm abm 3dnowprefetch cpuid_fault epb invpcid_singl  
e pti ssbd ibrs ibpb stibp tpr_shadow vnmi flexpriority ept vpid ept_ad fsgsbase  
tsc_adjust bmi1 hle avx2 smep bmi2 erms invpcid rtm mpx rdseed adx smap clflush  
opt intel_pt xsaveopt xsavec xgetbv1 xsaves dtherm ida arat pln pts hwp hwp_noti  
fy hwp_act_window hwp_epp md_clear flush_l1d  
vmx flags      : vmni preemption_timer invvpid ept_x_only ept_ad ept_1gb flexpr  
iority tsc_offset vtptr mtf vapic ept vpid unrestricted_guest ple shadow_vmcs pml  
ept_mode_based_exec  
bugs           : cpu_meltdown spectre_v1 spectre_v2 spec_store_bypass l1tf mds  
swapgs taa itlb_multihit srbds  
bogomips       : 5799.77  
clflush size   : 64  
cache_alignment : 64  
address sizes  : 39 bits physical, 48 bits virtual  
power management:  
ruslan@ruslan-ThinkPad-T470p:~$
```

100

[illegible]

Figure 5-18. 4-Kbyte PML4E—Long Mode

| | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
|-----------------------------|-----------|--|----|----|---|--|--|--|--|--|--|-----|--|-------------|--|-------------|----|-------------|-------------|-------------|-------------|----|---|---|---|---|---|---|
| 63 | 62 | | 52 | 51 | | | | | | | | | | | | | | | | | | 32 | | | | | | |
| N | Available | | | | Page-Directory Base Address (This is an architectural limit. A given implementation may support fewer bits.) | | | | | | | | | | | | | | | | | | | | | | | |
| X | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
| 31 | | | | | | | | | | | | | | | | 12 | 11 | | 9 | 8 | 7 | 6 | 5 | 4 | 3 | 2 | 1 | 0 |
| Page-Directory Base Address | | | | | | | | | | | | AVL | | I G N | | I G N | A | P C D | P W T | U / S | R / W | P | | | | | | |

Figure 5-19. 4-Kbyte PDPE—Long Mode

100

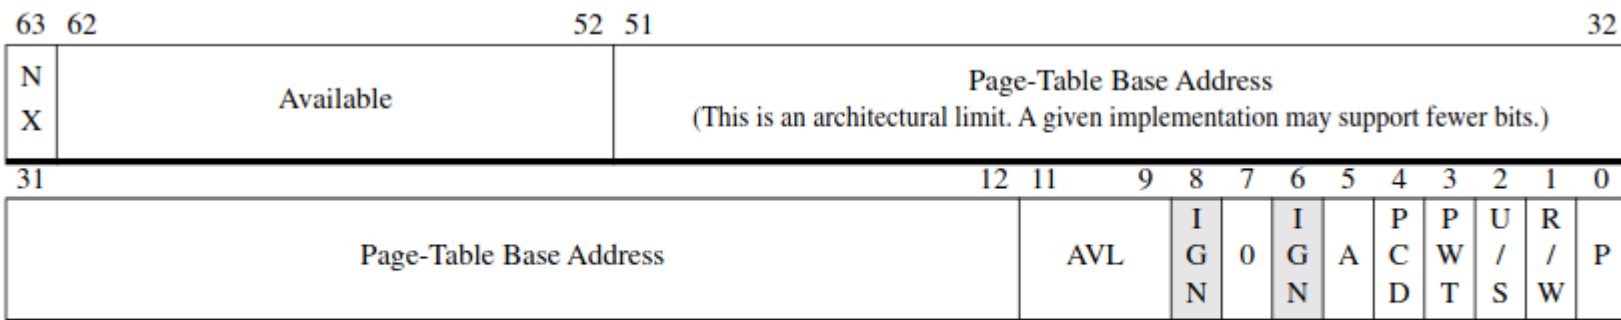


Figure 5-20. 4-Kbyte PDE—Long Mode




Figure 5-21. 4-Kbyte PTE—Long Mode

Example: PTE

```
typedef unsigned long long u64;

struct page_pte {
    u64 present:1;          // Bit P
    u64 writable:1;         // Bit R/W
    u64 user_mode:1;        // Bit U/S
    ...
    u64 page_address:40;    // 40+12 = 52-bit physical address (max)
    U64 avail:7;            // reserved, should be 0
    u64 pke:4;              // no MPK/PKE, should be 0
    u64 nonexecute:1;
};
```

Example: PTE

```
typedef unsigned long long u64;

struct page_pte {
    u64 present:1;          // Bit P
    u64 writable:1;         // Bit R/W
    u64 user_mode:1;        // Bit U/S
    ...
    u64 page_address:40;    // 40+12 = 52-bit physical address (max)
    u64 avail:7;           // reserved, should be 0
    u64 pke:4;             // no MPK/PKE, should be 0
    u64 nonexecute:1;
};

// _page_ address, i.e. memory_address / 4096 (or 2^12)
p->page_address = 0xZZZZZZZZZZULL; // ULL is unsigned long long
p->writable = 1;
p->present = 1;
p->user_mode = 0;
...
p->avail = 0;
p->pke = 0;
p->nonexecute = 0;
```


Example: Initializing 4 GB

For 4GB, we reference 1048576 physical pages

Using 2048 pages for PTEs, 4 pages for PDEs, 1 for PDPE, 1 for PMLE4E

Total: 2054 pages = 8413184 bytes, align at the 4096 boundary!

PTE:

```
struct page_pte *p; ... // Each entry is 8 bytes
```

```
for (int i = 0; i < 1048576; i++) { // 1048576/512 = 2048 pages
    p[i] = ... // physical pages 0,1,2,3,...,1048575 (absolute address)
}
```

PDE:

```
struct page_pde *pd = (struct page_pde *) (p + 1048576);
```

```
for (int j = 0; j < 2048; j++) { // 2048/512 = 4 pages
    struct page_pte *start_pte = p + 512 * j;
    page_addr = (u64) start_pte >> 12; // we record the page address
    ...
}
```

PDPE:

Reference 4 PDEs (1 page), everything else is empty

PMLE4E:

Just one reference to PDPE; everything else is empty



Example: Initializing 4 GB

- Do not use more than 1GB RAM! (qemu: 1024)
 - The Video RAM is after that (but before 4GB)

Aligning Pages

Allocate more space: e.g., $8413184 + 4095$

Align the allocated 'base':

```
(void *) (((unsigned long long) base + 4095) & (~4095ULL))
```

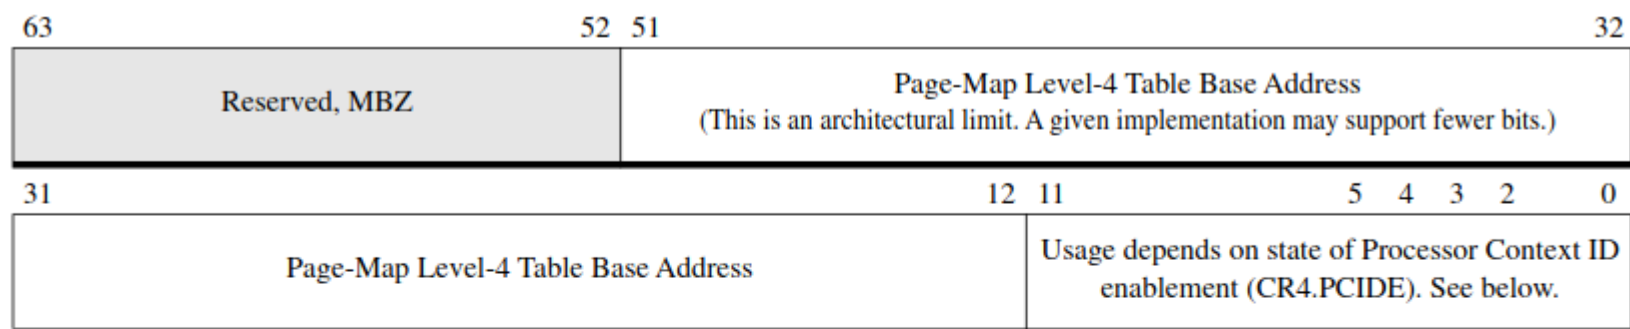
- Why long long?
 - The kernel uses 64-bit 'long' due to System V's ABI (aka the LP64 model)
 - The boot loader uses 32-bit 'long' due to EFI/Microsoft's ABI (aka the LLP64 model)
 - 'int' is 32 bit and 'long long' is 64 bit in either case

Loading Page Table

```
void write_cr3(unsigned long long cr3_value)
{
    asm volatile ("mov %0, %%cr3"
                  :
                  : "r" (cr3_value)
                  : "memory");
}
```

Loading Page Table

```
void write_cr3(unsigned long long cr3_value)
{
    asm volatile ("mov %0, %%cr3"
                  :
                  : "r" (cr3_value)
                  : "memory");
}
```



No PCID
(PCIDE=0)!

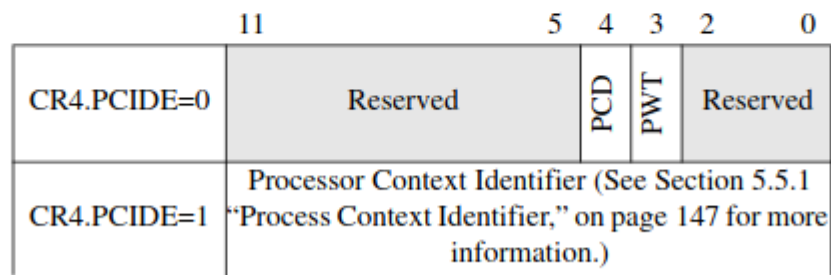


Figure 3-6. Control Register 3 (CR3)—Long Mode

* The picture is taken from <https://www.amd.com/system/files/TechDocs/24593.pdf>

Loading Page Table

Page-Level Writethrough (PWT) Bit. Bit 3. Page-level writethrough indicates whether the highest-level page-translation table has a writeback or writethrough caching policy. When PWT=0, the table has a writeback caching policy. When PWT=1, the table has a writethrough caching policy.

Page-Level Cache Disable (PCD) Bit. Bit 4. Page-level cache disable indicates whether the highest-level page-translation table is cacheable. When PCD=0, the table is cacheable. When PCD=1, the table is not cacheable.

PCD=0, PWT=0

OS Boot Process

- We have to load the main kernel image
 - An executable file (e.g., ELF) which is loaded at some **fixed** *virtual* address (e.g., 0x100000, right after the first 1MB of legacy BIOS/DOS memory)
 - Relocatable kernels are also possible
 - Kernel-address space layout randomization (KASLR) to enhance security
- The kernel image may load additional modules during the boot process
 - Typically relocatable executables

“Chicken-and-egg” Boot Problem

- We need to load file system and storage (or network) drivers to load kernel modules
 - But file system and storage (or network) drivers are often compiled as modules
 - We often cannot use firmware drivers once we left the OS boot loader
- What do we do?
 - The OS boot loader must preload critical (but not all!) modules

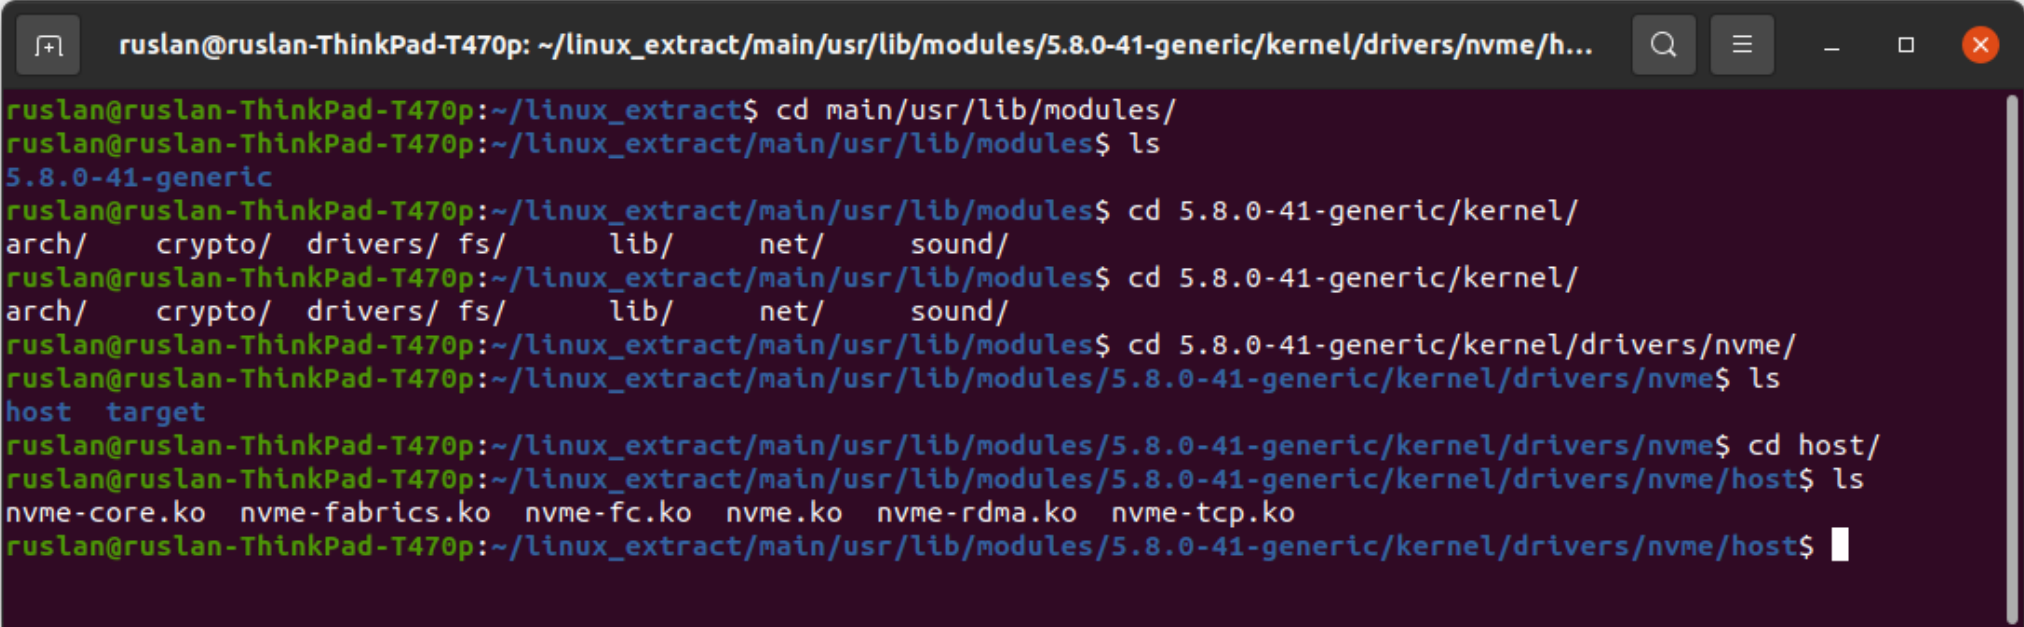
Example: Linux Protocol

- The Linux kernel executable image is compressed in the “/boot/vmlinuz” file
- All critical modules are placed in “initrd.img”
 - It also has other files (e.g., configuration, microcode, etc)
- For example, in GRUB (/boot/grub/grub.cfg):

```
linux /boot/vmlinuz-5.8.0-41-generic root=UUID=00000000-0000-0000-0000-000000000000 ro quiet splash $vt_handoff  
initrd /boot/initrd.img-5.8.0-41-generic
```

What is inside initrd.img?

```
$ mkdir linux_extract  
$ cd linux_extract/  
$ unmkinitramfs -v /boot/initrd.img-<version>-generic .
```



```
ruslan@ruslan-ThinkPad-T470p: ~/linux_extract/main/usr/lib/modules/5.8.0-41-generic/kernel/drivers/nvme/h...  
ruslan@ruslan-ThinkPad-T470p:~/linux_extract$ cd main/usr/lib/modules/  
ruslan@ruslan-ThinkPad-T470p:~/linux_extract/main/usr/lib/modules$ ls  
5.8.0-41-generic  
ruslan@ruslan-ThinkPad-T470p:~/linux_extract/main/usr/lib/modules$ cd 5.8.0-41-generic/kernel/  
arch/  crypto/  drivers/  fs/      lib/     net/     sound/  
ruslan@ruslan-ThinkPad-T470p:~/linux_extract/main/usr/lib/modules$ cd 5.8.0-41-generic/kernel/  
arch/  crypto/  drivers/  fs/      lib/     net/     sound/  
ruslan@ruslan-ThinkPad-T470p:~/linux_extract/main/usr/lib/modules$ cd 5.8.0-41-generic/kernel/drivers/nvme/  
ruslan@ruslan-ThinkPad-T470p:~/linux_extract/main/usr/lib/modules/5.8.0-41-generic/kernel/drivers/nvme$ ls  
host  target  
ruslan@ruslan-ThinkPad-T470p:~/linux_extract/main/usr/lib/modules/5.8.0-41-generic/kernel/drivers/nvme$ cd host/  
ruslan@ruslan-ThinkPad-T470p:~/linux_extract/main/usr/lib/modules/5.8.0-41-generic/kernel/drivers/nvme/host$ ls  
nvme-core.ko  nvme-fabrics.ko  nvme-fc.ko  nvme.ko  nvme-rdma.ko  nvme-tcp.ko  
ruslan@ruslan-ThinkPad-T470p:~/linux_extract/main/usr/lib/modules/5.8.0-41-generic/kernel/drivers/nvme/host$
```

Example: GRUB's "Multiboot" Protocol

- Can be used by many OS kernels if they implement this protocol (Examples: NetBSD, Xen, ...)
- For example for Xen (/boot/grub/grub.cfg)
 - We load the "kernel", called xen.gz (Xen)
 - We load modules
 - vmlinuz is a Xen "module", i.e., an OS kernel
 - initrd is another module (but for Linux)

```
multiboot2 /boot/xen.gz placeholder ${xen_rm_opts}
echo 'Loading Linux 5.8.0-41-generic ...'
module2 /boot/vmlinuz-5.8.0-41-generic placeholder root=UUID=00000000-
0000-0000-0000-000000000000 ro quiet splash
echo 'Loading initial ramdisk ...'
module2 --nounzip /boot/initrd.img-5.8.0-41-generic
```

State of the system

- Interrupts are disabled
 - Need to be set up
- Page tables can be very basic (“early boot”)
 - The kernel will set up its own page tables
- Other processors need to be brought up
 - Not clear if `EFI_MP_SERVICES_PROTOCOL` will change that in the future
- Little or no possibility to print anything right away
 - Some “early” `printk()` can still be available

Linux Boot Process

- For BIOS, there was no recourse
- For UEFI, we can specify CONFIG_EFI_STUB
 - The boot loader is embedded in the Linux kernel
 - ExitBootServices() is called by Linux
 - No support for Linux file systems, we have to place files on the EFI partition (even if the kernel has built-in drivers)
- EFI handover (new protocol)

```
linuxefi  /boot/efi/EFI/vmlinuz.efi  
initrdefi /boot/efi/EFI/initramfs.img
```

Linux Boot Process

- arch/x86/boot/header.S

```
ruslan@ruslan-ThinkPad-T470p: ~/linux-5.9.12/arch/x86
# offset 512, entry point

.globl _start
_start:
    # Explicitly enter this as bytes, or the assembler
    # tries to generate a 3-byte jump here, which causes
    # everything else to push off to the wrong offset.
    .byte 0xeb          # short (2-byte) jump
    .byte start_of_setup-1f

1:

    # Part 2 of the header, from the old setup.S

    .ascii "HdrS"       # header signature
    .word 0x020f         # header version number (>= 0x0105)
                        # or else old loadlin-1.5 will fail)
    .globl realmode_swch
realmode_swch: .word 0, 0 # default_switch, SETUPSEG
start_sys_seg: .word SYSSEG # obsolete and meaningless, but just
                        # in case something decided to "use" it
    .word kernel_version-512 # pointing to kernel version string
                        # above section of header is compatible

297,1-4 45%
```

Linux Boot Process

```
ruslan@ruslan-ThinkPad-T470p: ~/linux-5.9.12/arch/x86

.section ".entrytext", "ax"
start_of_setup:
# Force %es = %ds
    movw    %ds, %ax
    movw    %ax, %es
    cld

# Apparently some ancient versions of LILO invoked the kernel with %ss != %ds,
# which happened to work by accident for the old code. Recalculate the stack
# pointer if %ss is invalid. Otherwise leave it alone, LOADLIN sets up the
# stack behind its own code, so we can't blindly put it directly past the heap.

    movw    %ss, %dx
    cmpw    %ax, %dx    # %ds == %ss?
    movw    %sp, %dx
    je      2f          # -> assume %sp is reasonably set

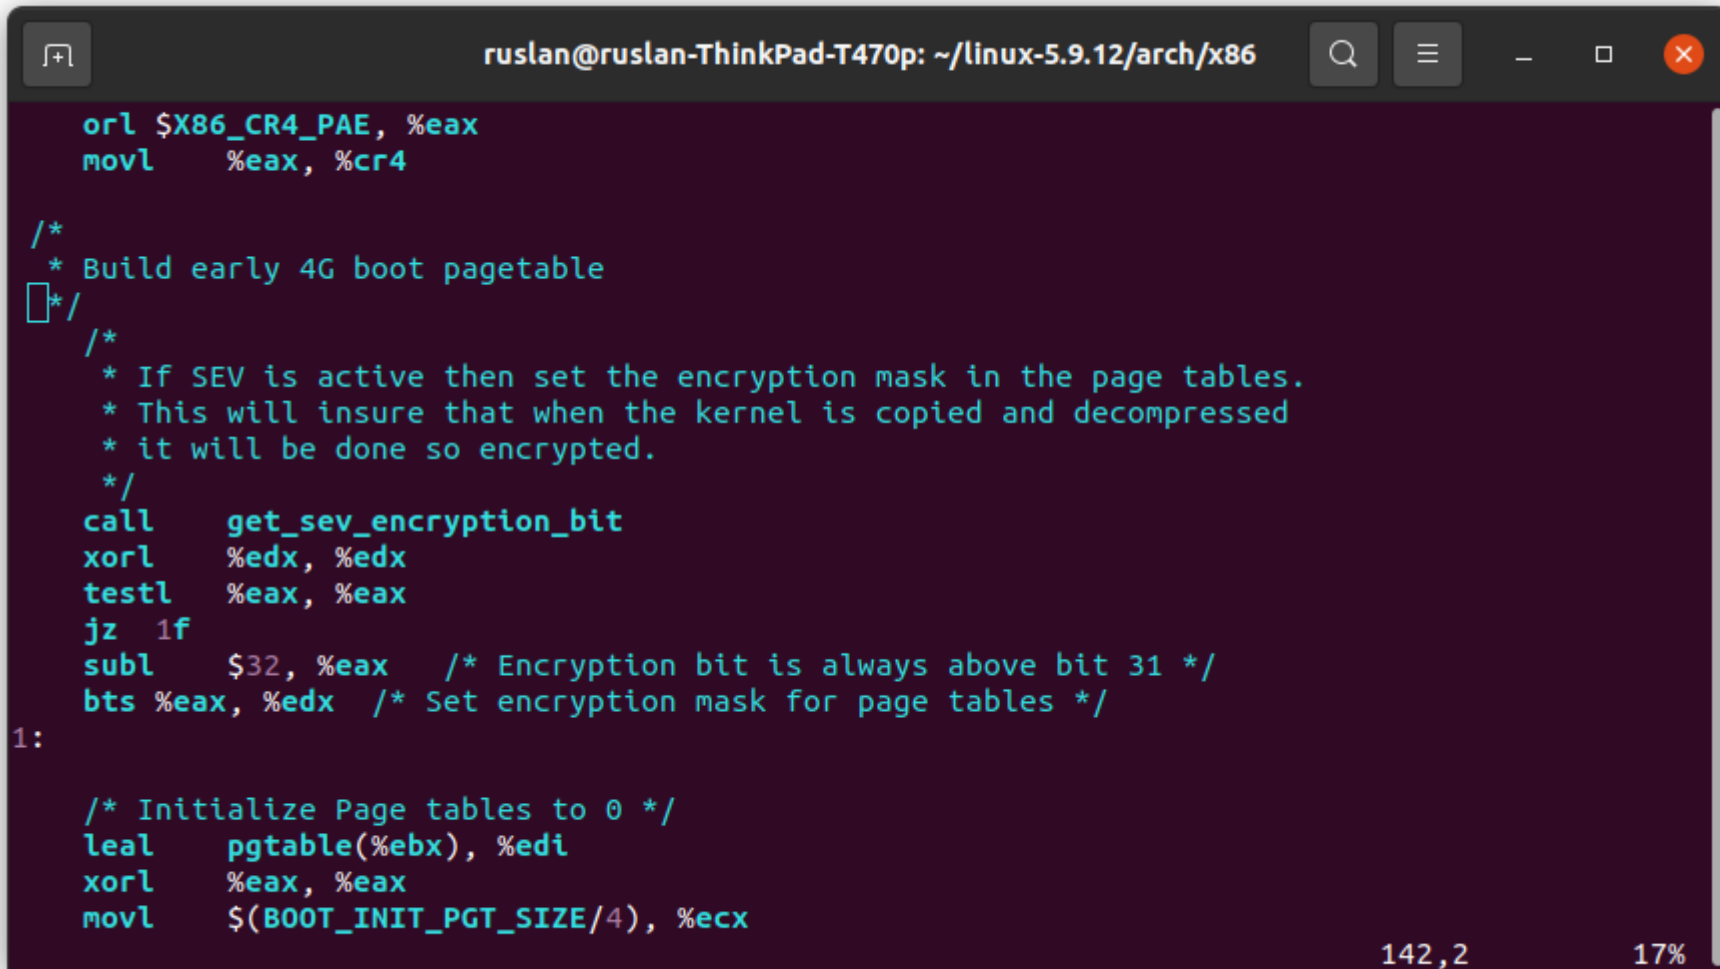
# Invalid %ss, make up a new stack
    movw    $_end, %dx
    testb   $CAN_USE_HEAP, loadflags
    jz      1f
    movw    heap_end_ptr, %dx
1:    addw    $STACK_SIZE, %dx
    jnc     2f
    xorw    %dx, %dx    # Prevent wraparound

2:    # Now %dx should point to the end of our stack space
```

587,1-4 92%

Linux Boot Process

- boot/compressed/head_64.S

A terminal window with a dark purple background and light green text. The window title is 'ruslan@ruslan-ThinkPad-T470p: ~/linux-5.9.12/arch/x86'. The code is assembly for the 'head_64.S' file. It starts with 'orl \$X86_CR4_PAE, %eax' and 'movl %eax, %cr4'. Then there is a comment block: '/* Build early 4G boot pagetable */'. Inside this block, there is another comment: '/* If SEV is active then set the encryption mask in the page tables. This will insure that when the kernel is copied and decompressed it will be done so encrypted. */'. This is followed by 'call get_sev_encryption_bit', 'xorl %edx, %edx', 'testl %eax, %eax', 'jz 1f', 'subl \$32, %eax /* Encryption bit is always above bit 31 */', and 'bts %eax, %edx /* Set encryption mask for page tables */'. Then there is a label '1:' followed by '/* Initialize Page tables to 0 */', 'leal pgtable(%ebx), %edi', 'xorl %eax, %eax', and 'movl \$(BOOT_INIT_PGT_SIZE/4), %ecx'. The bottom right of the terminal shows '142,2' and '17%'.

```
orl $X86_CR4_PAE, %eax
movl %eax, %cr4

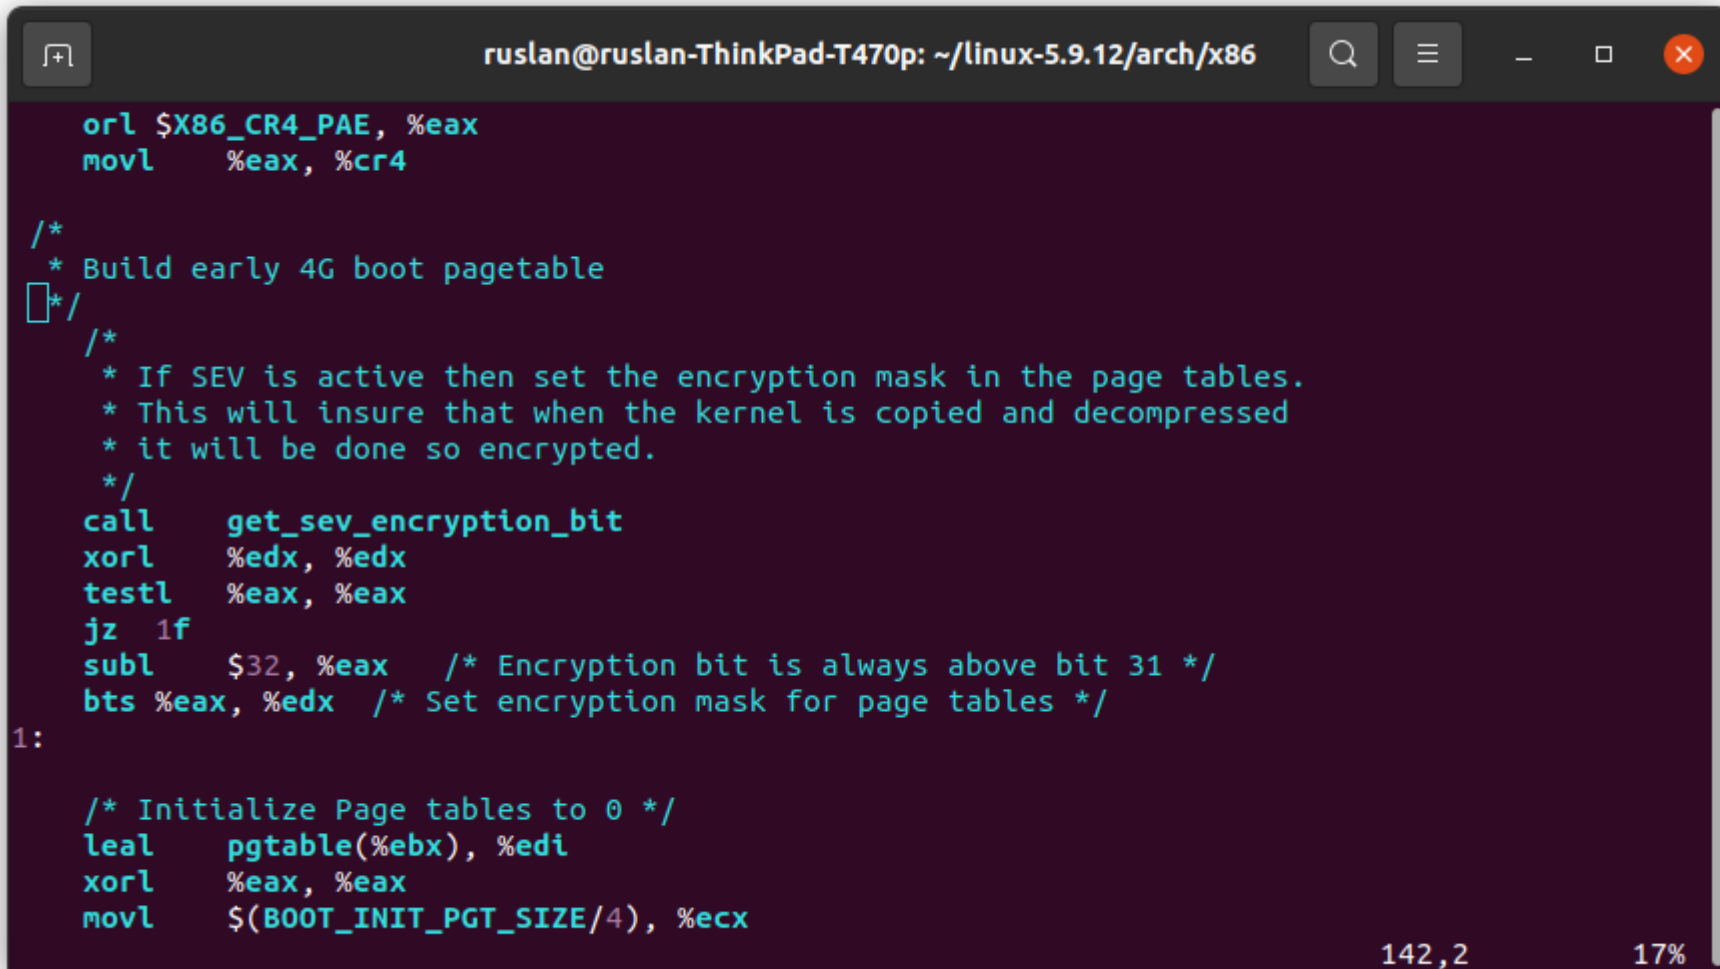
/*
 * Build early 4G boot pagetable
 */
/*
 * If SEV is active then set the encryption mask in the page tables.
 * This will insure that when the kernel is copied and decompressed
 * it will be done so encrypted.
 */
call    get_sev_encryption_bit
xorl    %edx, %edx
testl   %eax, %eax
jz      1f
subl    $32, %eax /* Encryption bit is always above bit 31 */
bts     %eax, %edx /* Set encryption mask for page tables */

1:

/* Initialize Page tables to 0 */
leal    pgtable(%ebx), %edi
xorl    %eax, %eax
movl    $(BOOT_INIT_PGT_SIZE/4), %ecx
```


Linux Boot Process

- boot/compressed/head_64.S

A terminal window with a dark purple background and light green text. The window title is 'ruslan@ruslan-ThinkPad-T470p: ~/linux-5.9.12/arch/x86'. The code is assembly for the 'head_64.S' file. It starts with 'orl \$X86_CR4_PAE, %eax' and 'movl %eax, %cr4'. Then there is a comment block: '/* Build early 4G boot pagetable */'. This is followed by another comment block: '/* If SEV is active then set the encryption mask in the page tables. This will insure that when the kernel is copied and decompressed it will be done so encrypted. */'. The code then calls 'get_sev_encryption_bit', XORs '%edx' with itself, tests '%eax', and jumps if zero to '1f'. It then subtracts 32 from '%eax' with a comment '/* Encryption bit is always above bit 31 */' and sets a bit in '%edx' with a comment '/* Set encryption mask for page tables */'. Finally, it initializes page tables to 0 by loading '%ebx' into '%edi', XORing '%eax' with itself, and moving '\$(BOOT_INIT_PGT_SIZE/4)' into '%ecx'. The bottom right of the terminal shows '142,2' and '17%'.

```
orl $X86_CR4_PAE, %eax
movl %eax, %cr4

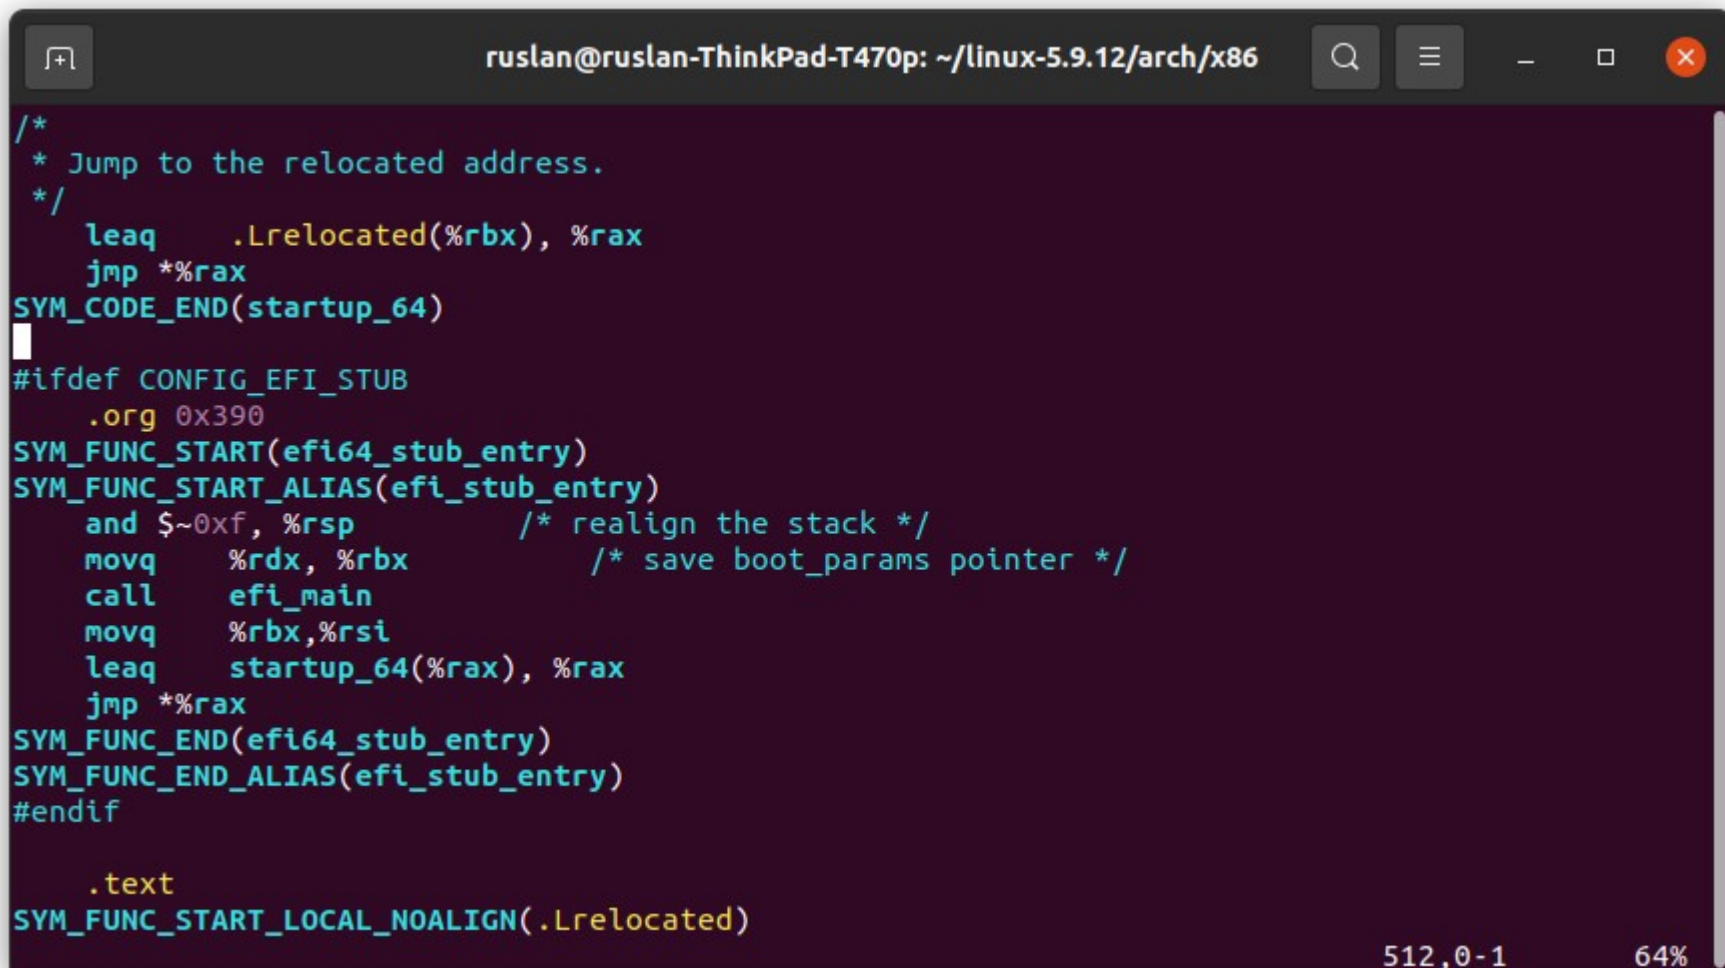
/*
 * Build early 4G boot pagetable
 */
/*
 * If SEV is active then set the encryption mask in the page tables.
 * This will insure that when the kernel is copied and decompressed
 * it will be done so encrypted.
 */
call    get_sev_encryption_bit
xorl    %edx, %edx
testl   %eax, %eax
jz      1f
subl    $32, %eax    /* Encryption bit is always above bit 31 */
bts     %eax, %edx   /* Set encryption mask for page tables */

1:

/* Initialize Page tables to 0 */
leal    pgtable(%ebx), %edi
xorl    %eax, %eax
movl    $(BOOT_INIT_PGT_SIZE/4), %ecx
```

Linux Boot Process (EFI)

- boot/compressed/head_64.S



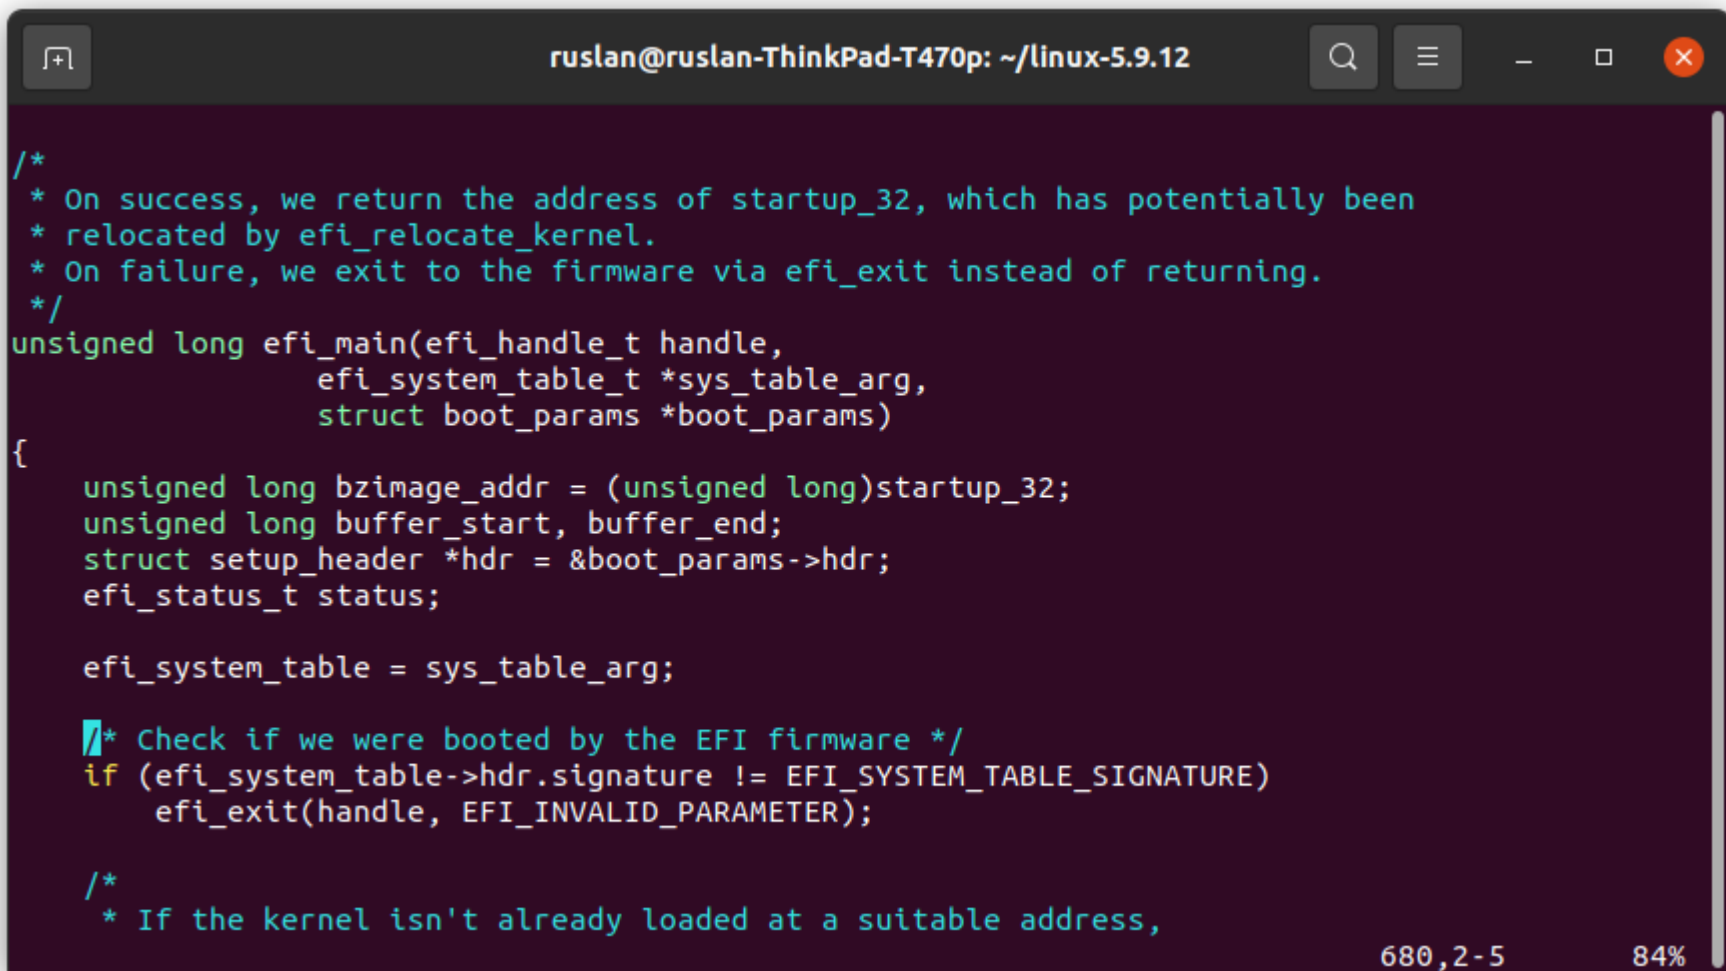
```
/*
 * Jump to the relocated address.
 */
    leaq    .Lrelocated(%rbx), %rax
    jmp *%rax
SYM_CODE_END(startup_64)
#ifdef CONFIG_EFI_STUB
    .org 0x390
SYM_FUNC_START(efi64_stub_entry)
SYM_FUNC_START_ALIAS(efi_stub_entry)
    and $~0xf, %rsp        /* realign the stack */
    movq    %rdx, %rbx      /* save boot_params pointer */
    call    efi_main
    movq    %rbx,%rsi
    leaq    startup_64(%rax), %rax
    jmp *%rax
SYM_FUNC_END(efi64_stub_entry)
SYM_FUNC_END_ALIAS(efi_stub_entry)
#endif

    .text
SYM_FUNC_START_LOCAL_NOALIGN(.Lrelocated)
```

512,0-1 64%

Linux Boot Process (EFI)

- drivers/firmware/efi/libstub/x86-stub.c



```
ruslan@ruslan-ThinkPad-T470p: ~/linux-5.9.12
/*
 * On success, we return the address of startup_32, which has potentially been
 * relocated by efi_relocate_kernel.
 * On failure, we exit to the firmware via efi_exit instead of returning.
 */
unsigned long efi_main(efi_handle_t handle,
                      efi_system_table_t *sys_table_arg,
                      struct boot_params *boot_params)
{
    unsigned long bzimage_addr = (unsigned long)startup_32;
    unsigned long buffer_start, buffer_end;
    struct setup_header *hdr = &boot_params->hdr;
    efi_status_t status;

    efi_system_table = sys_table_arg;

    /* Check if we were booted by the EFI firmware */
    if (efi_system_table->hdr.signature != EFI_SYSTEM_TABLE_SIGNATURE)
        efi_exit(handle, EFI_INVALID_PARAMETER);

    /*
     * If the kernel isn't already loaded at a suitable address,
```

680,2-5 84% 27 / 28

Linux Boot Process (EFI)

- drivers/firmware/efi/libstub/x86-stub.c

A terminal window with a dark purple background and light-colored text. The window title bar shows 'ruslan@ruslan-ThinkPad-T470p: ~/linux-5.9.12'. The code is C language, showing functions for EFI boot process. It includes calls to efi_random_get_seed, efi_retrieve_tpm2_eventlog, setup_graphics, setup_efi_pci, and setup_quirks. It then calls exit_boot and checks for success. If it fails, it calls efi_err and goto fail. The fail block calls efi_err and efi_exit. The window has standard Linux terminal window controls (search, menu, zoom, close) in the top right.

```
efi_random_get_seed();
efi_retrieve_tpm2_eventlog();
setup_graphics(boot_params);
setup_efi_pci(boot_params);
setup_quirks(boot_params);

status = exit_boot(boot_params, handle);
if (status != EFI_SUCCESS) {
    efi_err("exit_boot() failed!\n");
    goto fail;
}

return bimage_addr;
fail:
efi_err("efi_main() failed!\n");
efi_exit(handle, status);
}
```

794,0-1 Bot