


# CSE 511: Operating Systems Design


## Lectures 9-10



Thread Local Storage (TLS)  
In-kernel per-CPU variables  
Global Descriptor Table (GDT)  
Task State Segment (TSS)  
Interrupt Descriptor Table (IDT)

# Threads

- Linux systems use the 1:1 threading model
  - Each thread is assigned its own task
  - Other models are N:1 and M:N



# Threads

- N:1 – all threads are mapped to one task
  - Needs only minimal kernel support but it only allows the use of a single CPU for all threads
- M:N – M threads are mapped to N tasks
  - Reconciles the N:1 and 1:1 models
  - Normally  $M \geq N$ 
    - Needs a user-mode scheduler (e.g., in libpthread) to suspend some threads and resume others
  - Capable to context switch between different threads in a process without using the underlying OS kernel
    - Difficult to implement, e.g., blocking system calls can monopolize the underlying “task”

# Problems with system calls

- Can pollute TLB
  - User-space virtual addresses are different from kernel-space addresses
- Can pollute cache
  - Kernel data and code are different
- Direct cost of system calls
  - You need to switch CPU between different modes
  - Additional checks, different stacks in kernel space and user space, etc

# Thread Local Storage (TLS)

- All threads in one process share the **same** address space
- But how do we maintain per-thread variables?
  - e.g., **errno** in C must contain an error number when an operation/system call fails, operations can execute concurrently
- Thread Local Storage (TLS) is a special per-thread area which provides private memory address space for each thread
  - Should be accessed efficiently, i.e., no page table switches
  - Not supposed to be **fully** isolated from other threads if they somehow figure out the base address of this area, i.e., they can still corrupt this area since there is no isolation

# Thread Local Storage (TLS)

- TLS is architecture- and OS-specific, see “ELF Handling for Thread-Local Storage” for x86-64/Linux and other architectures (<https://akkadia.org/drepper/tls.pdf>)
- To support TLS, you typically need a special register
  - It specifies a thread-specific offset (base) in the memory
- x86 historically supported “segmentation”, where memory pointers would be bound to a specific segment
  - Segments can be more trivial, i.e., change the base address
  - x86-64 made segmentation obsolete but still supports base addresses for fs: and gs: segment registers

# Thread Local Storage (TLS)

- Many x86 instructions support segment override, by default ds: (data segment), es: (another data segment), or ss: (stack segment) are used depending on the instruction
  - Historically, it was supported for 16-bit and 32-bit CPUs
    - Segmentation was almost never used with 32-bit CPUs
  - ds:, es:, and ss: will now have their base=0 for x86-64
  - However, fs: and gs: can still change their base using WRMSR (FS.base is C000\_0100h, GS.base is C000\_0101h)
- fs: and gs: segments are still used
  - fs: is used by TLS in Linux
  - gs: is used by TLS in Windows

# Thread Local Storage (TLS)

## Create tls.c:

```
__thread int a = 5;
```

```
int func() {  
    return a;  
}
```

## Compile:

```
gcc -Wall -O2 -S tls.c
```

## tls.s:

```
func:
```

```
.LFB0:
```

```
.cfi_startproc
```

```
endbr64
```

```
movl  %fs:a@tpoff, %eax
```

```
ret
```

```
.cfi_endproc
```

```
...
```

```
.section .tdata,"awT",@progbits
```

```
.align 4
```

```
.type a, @object
```

```
.size a, 4
```

```
a:
```

```
.long 5
```



# Thread Local Storage (TLS)

## Create tls.c:

```
__thread int a = 5;
```

```
int func() {  
    return a;  
}
```

## Compile:

```
gcc -Wall -O2 -S tls.c
```

## tls.s:

```
func:
```

```
.LFB0:
```

```
.cfi_startproc
```

```
endbr64
```

```
movl  %fs:a@tpoff, %eax
```

```
ret
```

```
.cfi_endproc
```

```
...
```

```
.section .tdata,"awT",@progbits
```

```
.align 4
```

```
.type a, @object
```

```
.size a, 4
```

```
a:
```

```
.long 5
```

Segment override, use  
**fs:** instead of **ds:**

# Thread Local Storage (TLS)

## Create tls.c:

```
__thread int a = 5;
```

```
int func() {  
    return a;  
}
```

## Compile:

```
gcc -Wall -O2 -S tls.c
```

## tls.s:

```
func:
```

```
.LFB0:
```

```
.cfi_startproc
```

```
endbr64
```

```
movl  %fs:a@tpoff, %eax
```

```
ret
```

```
.cfi_endproc
```

```
...
```

```
.section .tdata,"awT",@progbits
```

```
.align 4
```

```
.type a, @object
```

```
.size a, 4
```

```
a:
```

```
.long 5
```

**a** is a symbol  
**@tpoff** – a relocation with an offset relative to the TLS block

# Thread Local Storage (TLS)

## Create `tls.c`:

```
__thread int a = 5;
```

```
int func() {  
    return a;  
}
```

## Compile:

```
gcc -Wall -O2 -S tls.c
```

## `tls.s`:

```
func:
```

```
.LFB0:
```

```
.cfi_startproc
```

```
endbr64
```

```
movl  %fs:a@tpoff, %eax
```

```
ret
```

```
.cfi_endproc
```

```
...
```

```
.section .tdata,"awT",@progbits
```

```
.align 4
```

```
.type a, @object
```

```
.size a, 4
```

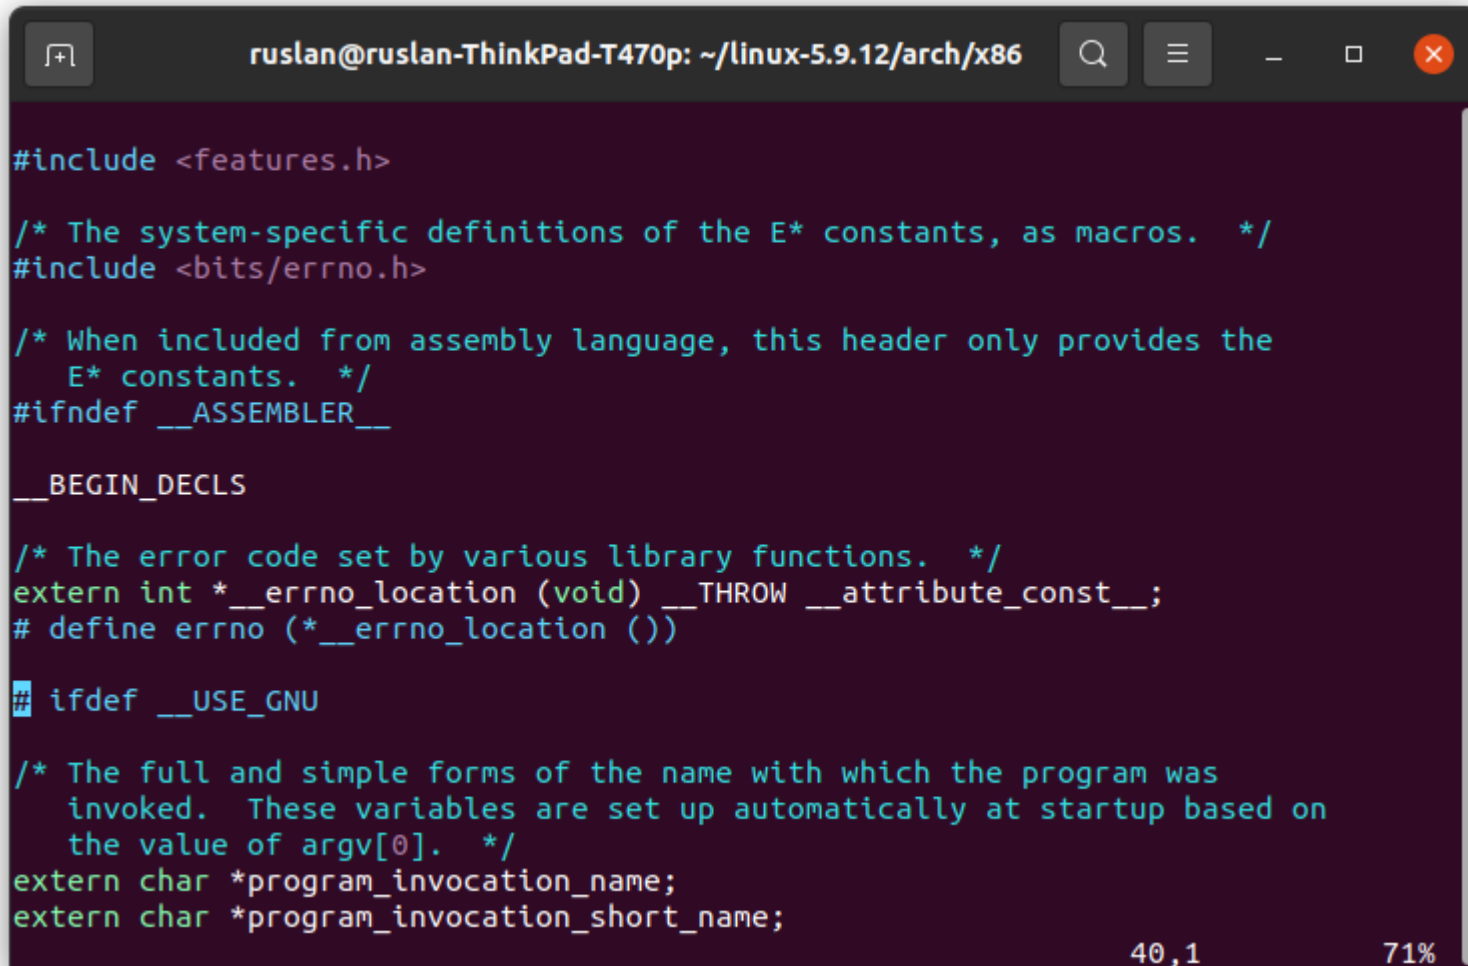
```
a:
```

```
.long 5
```

**.tdata** – a TLS .data section,  
works as an initialization image

# Thread Local Storage (TLS)

- What about errno? We can find the following trick in /usr/include/errno.h:



```
#include <features.h>

/* The system-specific definitions of the E* constants, as macros. */
#include <bits/errno.h>

/* When included from assembly language, this header only provides the
   E* constants. */
#ifndef __ASSEMBLER__

__BEGIN_DECLS

/* The error code set by various library functions. */
extern int *__errno_location (void) __THROW __attribute_const__;
# define errno (*__errno_location ())

#ifdef __USE_GNU


/* The full and simple forms of the name with which the program was
   invoked. These variables are set up automatically at startup based on
   the value of argv[0]. */
extern char *program_invocation_name;
extern char *program_invocation_short_name;
```

40,1 71%

# Thread Local Storage (TLS)

- Each address is calculated using a segment register
  - $[ds:] + \text{addr} = \text{addr}$ , since  $[ds:] = 0$
  - $[fs:] + \text{addr} = \text{addr} + \text{TLS\_block}$ , since  $[fs:]$  points to a TLS block
- `sym@tpoff` is `addr` in case of TLS
  - Non-negative addresses are used ***internally*** for a thread control block (TCB)
    - `fs:0` points to `TLS_block` itself
  - Negative addresses used for TLS variables
    - `@tpoff` is actually negative, -4 in the example above when linking is complete

# What is really happening here?



# Thread Local Storage (TLS)

## Create `tls.c`:

```
__thread int a[100];

int func() {
    int sum = 0;
    for (int i = 0; i < 100; i++)
        sum += a[i];
    return sum;
}
```

## Compile:

```
gcc -Wall -O2 -S tls.c
```

- Linux segment overrides are used sparingly, `fs:0` is sometimes preferred

## `tls.s`:

```
func:
.LFB0:
    .cfi_startproc
    endbr64
    movq    %fs:0, %rdx
    xorl    %r8d, %r8d
    leaq    a@tpoff(%rdx), %rax
    addq    $400+a@tpoff, %rdx
    .p2align 4,,10
    .p2align 3
.L2:
    addl    (%rax), %r8d
    addq    $4, %rax
    cmpq    %rdx, %rax
    jne .L2
    movl    %r8d, %eax
    ret
    .cfi_endproc
```

# Thread Local Storage (TLS)

## Create `tls.c`:

```
__thread int a[100];

int func() {
    int sum = 0;
    for (int i = 0; i < 100; i++)
        sum += a[i];
    return sum;
}
```

## Compile:

```
gcc -Wall -O2 -S tls.c
```

- Linux segment overrides are used sparingly, `fs:0` is sometimes preferred

## `tls.s`:

```
func:
.LFB0:
    .cfi_startproc
    endbr64
    movq    %fs:0, %rdx
    xorl    %r8d, %r8d
    leaq    a@tpoff(%rdx), %rax
    addq    $400+a@tpoff, %rdx
    .p2align 4,,10
    .p2align 3
.L2:
    addl    (%rax), %r8d
    addq    $4, %rax
    cmpq    %rdx, %rax
    jne .L2
    movl    %r8d, %eax
    ret
    .cfi_endproc
```

**%rdx – TLS\_block**



# Thread Local Storage (TLS)

## Create `tls.c`:

```
__thread int a[100];

int func() {
    int sum = 0;
    for (int i = 0; i < 100; i++)
        sum += a[i];
    return sum;
}
```

## Compile:

```
gcc -Wall -O2 -S tls.c
```

- Linux segment overrides are used sparingly, `fs:0` is sometimes preferred

## `tls.s`:

```
func:
.LFB0:
    .cfi_startproc
    endbr64
    movq    %fs:0, %rdx
    xorl    %r8d, %r8d
    leaq    a@tpoff(%rdx), %rax
    addq    $400+a@tpoff, %rdx
    .p2align 4,,10
    .p2align 3
.L2:
    addl    (%rax), %r8d
    addq    $4, %rax
    cmpq    %rdx, %rax
    jne .L2
    movl    %r8d, %eax
    ret
    .cfi_endproc
```

**%rax** – the beginning of the array

# Thread Local Storage (TLS)

## Create `tls.c`:

```
__thread int a[100];

int func() {
    int sum = 0;
    for (int i = 0; i < 100; i++)
        sum += a[i];
    return sum;
}
```

## Compile:

```
gcc -Wall -O2 -S tls.c
```

- Linux segment overrides are used sparingly, `fs:0` is sometimes preferred

## `tls.s`:

```
func:
.LFB0:
    .cfi_startproc
    endbr64
    movq    %fs:0, %rdx
    xorl    %r8d, %r8d
    leaq    a@tpoff(%rdx), %rax
    addq    $400+a@tpoff, %rdx
    .p2align 4,,10
    .p2align 3
.L2:
    addl    (%rax), %r8d
    addq    $4, %rax
    cmpq    %rdx, %rax
    jne .L2
    movl    %r8d, %eax
    ret
    .cfi_endproc
```

**%rdx** – the end  
of the array

# Thread Local Storage (TLS)

## Create `tls.c`:

```
__thread int a[100];

int func() {
    int sum = 0;
    for (int i = 0; i < 100; i++)
        sum += a[i];
    return sum;
}
```

## Compile:

```
gcc -Wall -O2 -S tls.c
```

- Linux segment overrides are used sparingly, `fs:0` is sometimes preferred

## `tls.s`:

```
func:
.LFB0:
    .cfi_startproc
    endbr64
    movq    %fs:0, %rdx
    xorl    %r8d, %r8d
    leaq    a@tpoff(%rdx), %rax
    addq    $400+a@tpoff, %rdx
    .p2align 4,,10
    .p2align 3
.L2:
    addl    (%rax), %r8d
    addq    $4, %rax
    cmpq    %rdx, %rax
    jne .L2
    movl    %r8d, %eax
    ret
    .cfi_endproc
```

the actual loop

# Other usages of [fs:0]

## Create tls.c:

```
__thread int a = 5;
```

```
int func() {  
    return a;  
}
```

## Compile:

```
gcc -Wall -O2 -mno-tls-direct-seg-refs -S tls.c
```

## tls.s:

```
func:
```

```
.LFB0:
```

```
.cfi_startproc
```

```
endbr64
```

```
movq  %fs:0, %rax
```

```
movl  a@tpoff(%rax), %eax
```


```
ret
```

```
.cfi_endproc
```

- It is even possible to disable direct access through segment registers, and **always** use fs:0
- Why?
  - Until recently fs: and gs: could only be updated in privileged mode through WRMSR

# Other usages of [fs:0]

- But what if you need to context switch in unprivileged mode?
  - Not in Linux since the 1:1 threading model is used
  - Consider M:N threading with a user-space scheduler
- Also useful in other places
  - Microkernels, hypervisors that use paravirtualization
- **Example:** an OS can use indirection by creating a “dummy” TCB, which points to a real TCB
  - Still works in a regular case (Linux), fs:0 points to itself



# In-kernel usage

- We have a similar (but not identical) problem in the kernel
  - We want to have per-CPU variables
  - There is a similar set of variables but each CPU will have its own value (e.g., interrupt stack context)
- Why not use the other register (gs:) for that purpose?
  - Linux implements exactly that in percpu.h

# Global Descriptor Table (GDT)

- Historically, x86 provided a very elaborated virtual memory system
  - Paging (e.g., Windows, Linux, and most other OSes)
  - Segmentation
  - Or even both (e.g., OS/2)
- Even though segmentation is now obsolete in x86-64, certain pieces have to be initialized correctly
  - Thread Local Storage (TLS) uses fs: or gs: segment registers
  - All other segment registers have to be initialized appropriately (cs, ds, es, ss)

# Global Descriptor Table (GDT)

- A segment is an **isolated** contiguous piece of memory space, which has a **base** address and **limit** (size - 1)
  - You can restrict access to segments (e.g., each process can access just its own segments)
  - Compare that with a per-process page table (and TLB), though no complex translation is needed
    - An address will be `segment_base + offset`, offset is your program's (segment's) virtual address
- A segment register refers to a segment **selector**
  - Cannot be arbitrarily replaced without proper permissions
- Except TLS, you want to select segments with `base = 0` and `limit = MAX_ADDRESS` for x86-64 ("flat memory")



# Global Descriptor Table (GDT)

- GDT is a table of segment *descriptors*
  - Each entry is 8 bytes
  - The first entry is reserved (0)
  - Other entries can define segments for code and data
  - For x86-64, base and limit are always 0x0000000000000000 and 0xFFFFFFFFFFFFFFFF, respectively

31				16				15				0			
Base 0:15								Limit 0:15							
63		56		55	52	51	48	47		40		39	32		
Base 24:31		Flags		Limit 16:19		Access Byte		Base 16:23							



# Rumprun-SMP

- A unikernel based on NetBSD
- Do not copy any code but you can check it as an example to get some idea
- A version that supports SMP is available at
  - <https://github.com/ssrg-vt/rumprun-smp>

# Example of GDT (rumprun-smp)

```
ruslan@ruslan-ThinkPad-T470p: ~/test/rumprun-smp

/*
 * amd64 programmer's manual:
 *
 * "In long mode, segmentation is not used ... except for a few exceptions."
 *
 * Uuuyea, exceptions.
 */

.data
.align 64
.globl cpu_gdt64
cpu_gdt64:
    .quad 0x0000000000000000
    .quad 0x00af9b0000000fff /* 64bit CS */
    .quad 0x00cf9b0000000fff /* 32bit CS */
    .quad 0x00cf930000000fff /* DS */
    .rept BMK_MAXCPUS
    .quad 0x0000000000000000 /* TSS part 1 (via C) */
    .quad 0x0000000000000000 /* TSS part 2 (via C) */
    .endr
gdt64_end:
.align 64

.type gdt64_ptr, @object
gdt64_ptr:
    .word gdt64_end-cpu_gdt64-1
    .quad cpu_gdt64
```

225,1 85%

# Example of GDT (rump-run-smp)

```
ruslan@ruslan-ThinkPad-T470p: ~/test/rump-run-smp

/*
 * amd64 programmer's manual:
 *
 * "In long mode, segmentation is not used ... except for a few exceptions."
 *
 * Uuueea, exceptions.
 */

.data
.align 64
.globl cpu_gdt64
cpu_gdt64:
    .quad 0x0000000000000000
    .quad 0x00af9b0000000fff /* 64bit CS */
    .quad 0x00cf9b0000000fff /* 32bit CS */
    .quad 0x00cf930000000fff /* DS */
    .rept BMK_MAXCPUS
    .quad 0x0000000000000000 /* TSS part 1 (via C) */
    .quad 0x0000000000000000 /* TSS part 2 (via C) */
    .endr
gdt64_end:
.align 64

.type gdt64_ptr, @object
gdt64_ptr:
    .word gdt64_end-cpu_gdt64-1
    .quad cpu_gdt64
```

GDT entries  
(8 bytes each)

225,1 85%

# Example of GDT (rumprun-smp)

```
ruslan@ruslan-ThinkPad-T470p: ~/test/rumprun-smp

/*
 * amd64 programmer's manual:
 *
 * "In long mode, segmentation is not used ... except for a few exceptions."
 *
 * Uuuyea, exceptions.
 */

.data
.align 64
.globl cpu_gdt64
cpu_gdt64:
    .quad 0x0000000000000000
    .quad 0x00af9b0000000fff /* 64bit CS */
    .quad 0x00cf9b0000000fff /* 32bit CS */
    .quad 0x00cf930000000fff /* DS */
    .rept BMK_MAXCPUS
    .quad 0x0000000000000000 /* TSS part 1 (via C) */
    .quad 0x0000000000000000 /* TSS part 2 (via C) */
    .endr
gdt64_end:
.align 64

.type gdt64_ptr, @object
gdt64_ptr:
    .word gdt64_end-cpu_gdt64-1
    .quad cpu_gdt64
```

80-bit "pointer" to GDT:  
16-bit (size-1)  
64-bit (pointer)

Loading GDT  
with:

**lgdt (gdt64\_ptr)**

85%

# Segment Registers

- Need to initialize all data-related segment registers
  - Use (byte) offsets in GDT
  - FS and GS are handled differently (TLS), you can change their 64-bit base

```
ruslan@ruslan-ThinkPad-T470p: ~/test/rumprun-smp
lgdt (gdt64_ptr)
pushl $0x0
pushw $0x10
pushl $1f
lret
1: movl $0x18, %eax
   movl %eax, %ds
   movl %eax, %es
   movl %eax, %ss
   xorl %eax, %eax
   movl %eax, %fs
   movl %eax, %gs
```

"DS" in GDT,  
0x18

# Segment Registers

- Need to initialize the code segment register
  - You cannot write to %cs, use “long/far jump” or “long/far return”, i.e., change both the instruction pointer and %cs
  - For long mode (your code is already in 64-bit mode), both *pushw* and *pushl* here **must be replaced** with *pushq*



```
ruslan@ruslan-ThinkPad-T470p: ~/test/rumprun-smp
/* 5: poetically longjump to longmode */
pushw $0x08
pushl $_start64
lret

/* NOTREACHED */
jmp haltme

nomultiboot:
/* we don't have printf available yet just output manually */
mov $nomultibootmsg, %ebx
mov $(CONS_ADDRESS), %ecx
```

"64-bit CS"  
in GDT, 0x08

# Task State Segment (TSS)

- Historically, **hardware** context switches were supported
  - But this approach is no longer preferred, **software** context switches can be more optimized (e.g., in Linux)
  - Hardware “tasks” use TSS
- Even though they are mostly obsolete in x86-64, you still need to set up TSS, at least one TSS per each CPU
- Why do you still need TSS?
  - If an interrupt happens while you are in user mode, you must set up a correct stack pointer for privileged mode
    - Remember that you cannot trust user-space applications! (And their stack pointer values.)



# Task State Segment (TSS)

```
/*
 * amd64 programmer's manual:
 *
 * "In long mode, segmentation is not used ... except for a few exceptions."
 *
 * Uuuyea, exceptions.
 */

.data
.align 64
.globl cpu_gdt64
cpu_gdt64:
    .quad 0x0000000000000000
    .quad 0x00af9b0000000fff /* 64bit CS */
    .quad 0x00cf9b0000000fff /* 32bit CS */
    .quad 0x00cf930000000fff /* DS */
    .rept BMK_MAXCPUS
    .quad 0x0000000000000000 /* TSS part 1 (via C) */
    .quad 0x0000000000000000 /* TSS part 2 (via C) */
    .endr
gdt64_end:
.align 64

.type gdt64_ptr, @object
gdt64_ptr:
    .word gdt64_end-cpu_gdt64-1
    .quad cpu_gdt64
```

TSS needs  
two descriptors!  
(16 bytes)

225,1

85%

33 / 48

# Task State Segment (TSS)

- In the previous example, TSS descriptors are not yet initialized
- They are initialized later in the code

See `platform/hw/arch/amd64/machdep.c`:

```
struct taskgate_descriptor {  
    unsigned long td_lolimit:16; /* segment extent (lsb) */  
    unsigned long td_lobase:24; /* segment base address (lsb) */  
    unsigned long td_type:5; /* segment type */  
    unsigned long td_dpl:2; /* segment descriptor priority level */  
    unsigned long td_p:1; /* segment descriptor present */  
    unsigned long td_hilimit:4; /* segment extent (msb) */  
    unsigned long td_xx1:3; /* avl, long and def32 (not used) */  
    unsigned long td_gran:1; /* limit granularity (byte/page) */  
    unsigned long td_hibase:40; /* segment base address (msb) */  
    unsigned long td_xx2:8; /* reserved */  
    unsigned long td_zero:5; /* must be zero */  
    unsigned long td_xx3:19; /* reserved */  
} __packed;
```

# Task State Segment (TSS)

See `platform/hw/arch/amd64/machdep.c`:

```
static inline void
_init_taskgate(unsigned long cpu)
{
    unsigned long d = 4 + cpu * 2; /* TSS descriptor offset. */
    struct taskgate_descriptor *td = (void *)&cpu_gdt64[d];

    td->td_lolimit = 0;
    td->td_lobase = 0;
    td->td_type = 0x9;
    td->td_dpl = 0;
    td->td_p = 1;
    td->td_hilimit = 0xf;
    td->td_gran = 0;
    td->td_hibase = 0xffffffffUL;
    td->td_zero = 0;
    amd64_ltr(d*8);
}
```

Assume just one CPU (cpu=0)  
for simplicity, i.e. d=4

`amd64_ltr(0x20)`, i.e.  $4*8=32$

# Global Descriptor Table (GDT)

- Historically, x86 provided a very elaborated virtual memory system
  - Paging (e.g., Windows, Linux, and most other OSes)
  - Segmentation
  - Or even both (e.g., OS/2)
- Even though segmentation is now obsolete in x86-64, certain pieces have to be initialized correctly
  - Thread Local Storage (TLS) uses fs: or gs: segment registers
  - All other segment registers have to be initialized appropriately (cs, ds, es, ss)

# Global Descriptor Table (GDT)

- A segment is an **isolated** contiguous piece of memory space, which has a **base** address and **limit** (size - 1)
  - You can restrict access to segments (e.g., each process can access just its own segments)
  - Compare that with a per-process page table (and TLB), though no complex translation is needed
    - An address will be `segment_base + offset`, offset is your program's (segment's) virtual address
- A segment register refers to a segment **selector**
  - Cannot be arbitrarily replaced without proper permissions
- Except TLS, you want to select segments with `base = 0` and `limit = MAX_ADDRESS` for x86-64 ("flat memory")

# Global Descriptor Table (GDT)

- GDT is a table of segment *descriptors*
  - Each entry is 8 bytes
  - The first entry is reserved (0)
  - Other entries can define segments for code and data
  - For x86-64, base and limit are always 0x0000000000000000 and 0xFFFFFFFFFFFFFFFF, respectively

31				16				15				0			
Base 0:15								Limit 0:15							
63		56		55	52	51	48	47		40		39	32		
Base 24:31		Flags		Limit 16:19		Access Byte		Base 16:23							

# Example of GDT (rumprun-smp)

```
ruslan@ruslan-ThinkPad-T470p: ~/test/rumprun-smp

/*
 * amd64 programmer's manual:
 *
 * "In long mode, segmentation is not used ... except for a few exceptions."
 *
 * Uuuyea, exceptions.
 */

.data
.align 64
.globl cpu_gdt64
cpu_gdt64:
    .quad 0x0000000000000000
    .quad 0x00af9b0000000fff /* 64bit CS */
    .quad 0x00cf9b0000000fff /* 32bit CS */
    .quad 0x00cf930000000fff /* DS */
    .rept BMK_MAXCPUS
    .quad 0x0000000000000000 /* TSS part 1 (via C) */
    .quad 0x0000000000000000 /* TSS part 2 (via C) */
    .endr
gdt64_end:
.align 64

.type gdt64_ptr, @object
gdt64_ptr:
    .word gdt64_end-cpu_gdt64-1
    .quad cpu_gdt64
```

225,1 85%

# Example of GDT (rumprun-smp)

```
ruslan@ruslan-ThinkPad-T470p: ~/test/rumprun-smp

/*
 * amd64 programmer's manual:
 *
 * "In long mode, segmentation is not used ... except for a few exceptions."
 *
 * Uuuyea, exceptions.
 */

.data
.align 64
.globl cpu_gdt64
cpu_gdt64:
    .quad 0x0000000000000000
    .quad 0x00af9b0000000fff /* 64bit CS */
    .quad 0x00cf9b0000000fff /* 32bit CS */
    .quad 0x00cf930000000fff /* DS */
    .rept BMK_MAXCPUS
    .quad 0x0000000000000000 /* TSS part 1 (via C) */
    .quad 0x0000000000000000 /* TSS part 2 (via C) */
    .endr
gdt64_end:
.align 64

.type gdt64_ptr, @object
gdt64_ptr:
    .word gdt64_end-cpu_gdt64-1
    .quad cpu_gdt64
```

GDT entries  
(8 bytes each)

225,1 85%



# Example of GDT (rumprun-smp)

```
ruslan@ruslan-ThinkPad-T470p: ~/test/rumprun-smp

/*
 * amd64 programmer's manual:
 *
 * "In long mode, segmentation is not used ... except for a few exceptions."
 *
 * Uuuyea, exceptions.
 */

.data
.align 64
.globl cpu_gdt64
cpu_gdt64:
    .quad 0x0000000000000000
    .quad 0x00af9b0000000fff /* 64bit CS */
    .quad 0x00cf9b0000000fff /* 32bit CS */
    .quad 0x00cf930000000fff /* DS */
    .rept BMK_MAXCPUS
    .quad 0x0000000000000000 /* TSS part 1 (via C) */
    .quad 0x0000000000000000 /* TSS part 2 (via C) */
    .endr
gdt64_end:
.align 64

.type gdt64_ptr, @object
gdt64_ptr:
    .word gdt64_end-cpu_gdt64-1
    .quad cpu_gdt64
```

80-bit "pointer" to GDT:  
16-bit (size-1)  
64-bit (pointer)

Loading GDT  
with:

**lgdt (gdt64\_ptr)**

85%

# Segment Registers

- Need to initialize all data-related segment registers
  - Use (byte) offsets in GDT
  - FS and GS are handled differently (TLS), you can change their 64-bit base

```
ruslan@ruslan-ThinkPad-T470p: ~/test/rumprun-smp
lgdt (gdt64_ptr)
pushl $0x0
pushw $0x10
pushl $1f
lret
1: movl $0x18, %eax
   movl %eax, %ds
   movl %eax, %es
   movl %eax, %ss
   xorl %eax, %eax
   movl %eax, %fs
   movl %eax, %gs
```

"DS" in GDT,  
0x18

# Segment Registers

- Need to initialize the code segment register
  - You cannot write to %cs, use “long/far jump” or “long/far return”, i.e., change both the instruction pointer and %cs
  - For long mode (your code is already in 64-bit mode), both *pushw* and *pushl* here **must be replaced** with *pushq*



```
ruslan@ruslan-ThinkPad-T470p: ~/test/rumprun-smp
/* 5: poetically longjump to longmode */
pushw $0x08
pushl $_start64
lret

/* NOTREACHED */
jmp haltme

nomultiboot:
/* we don't have printf available yet just output manually */
mov $nomultibootmsg, %ebx
mov $(CONS_ADDRESS), %ecx
```

"64-bit CS"  
in GDT, 0x08

# Task State Segment (TSS)

- Historically, **hardware** context switches were supported
  - But this approach is no longer preferred, **software** context switches can be more optimized (e.g., in Linux)
  - Hardware “tasks” use TSS
- Even though they are mostly obsolete in x86-64, you still need to set up TSS, at least one TSS per each CPU
- Why do you still need TSS?
  - If an interrupt happens while you are in user mode, you must set up a correct stack pointer for privileged mode
    - Remember that you cannot trust user-space applications! (And their stack pointer values.)

# Task State Segment (TSS)

```
/*
 * amd64 programmer's manual:
 *
 * "In long mode, segmentation is not used ... except for a few exceptions."
 *
 * Uuuyea, exceptions.
 */

.data
.align 64
.globl cpu_gdt64
cpu_gdt64:
    .quad 0x0000000000000000
    .quad 0x00af9b0000000fff /* 64bit CS */
    .quad 0x00cf9b0000000fff /* 32bit CS */
    .quad 0x00cf930000000fff /* DS */
    .rept BMK_MAXCPUS
    .quad 0x0000000000000000 /* TSS part 1 (via C) */
    .quad 0x0000000000000000 /* TSS part 2 (via C) */
    .endr
gdt64_end:
.align 64

.type gdt64_ptr, @object
gdt64_ptr:
    .word gdt64_end-cpu_gdt64-1
    .quad cpu_gdt64
```

TSS needs  
two descriptors!  
(16 bytes)

225,1

85%

45 / 48

# Task State Segment (TSS)

- In the previous example, TSS descriptors are not yet initialized
- They are initialized later in the code

See `platform/hw/arch/amd64/machdep.c`:

```
struct taskgate_descriptor {  
    unsigned long td_lolimit:16; /* segment extent (lsb) */  
    unsigned long td_lobase:24; /* segment base address (lsb) */  
    unsigned long td_type:5; /* segment type */  
    unsigned long td_dpl:2; /* segment descriptor priority level */  
    unsigned long td_p:1; /* segment descriptor present */  
    unsigned long td_hilimit:4; /* segment extent (msb) */  
    unsigned long td_xx1:3; /* avl, long and def32 (not used) */  
    unsigned long td_gran:1; /* limit granularity (byte/page) */  
    unsigned long td_hibase:40; /* segment base address (msb) */  
    unsigned long td_xx2:8; /* reserved */  
    unsigned long td_zero:5; /* must be zero */  
    unsigned long td_xx3:19; /* reserved */  
} __packed;
```

# Task State Segment (TSS)

See `platform/hw/arch/amd64/machdep.c`:

```
static inline void
_init_taskgate(unsigned long cpu)
{
    unsigned long d = 4 + cpu * 2; /* TSS descriptor offset. */
    struct taskgate_descriptor *td = (void *)&cpu_gdt64[d];

    td->td_lolimit = 0;
    td->td_lobase = 0;
    td->td_type = 0x9;
    td->td_dpl = 0;
    td->td_p = 1;
    td->td_hilimit = 0xf;
    td->td_gran = 0;
    td->td_hibase = 0xffffffffUL;
    td->td_zero = 0;
    amd64_ltr(d*8);
}
```

Assume just one CPU (cpu=0)  
for simplicity, i.e. d=4

`amd64_ltr(0x20)`, i.e.  $4*8=32$

# Interrupt Descriptor Table (IDT)

- Loading IDT is similar to that of GDT. You need to have an 80-bit “pointer”. Below is an example for this in C.

From `platform/hw/arch/amd64/machdep.c`:

```
struct region_descriptor {  
    unsigned short rd_limit;  
    unsigned long rd_base;  
} __attribute__((__packed__)); <-- “packed” is important here to avoid padding
```

```
struct region_descriptor region;  
...  
region.rd_limit = sizeof(idt)-1;  
region.rd_base = (uintptr_t)(void *)idt;  
amd64_lidt(&region);
```

Where `amd64_lidt` is just the **lidt** instruction

```
ENTRY(amd64_lidt)  
    lidt (%rdi)  
    ret  
END(amd64_lidt)
```