



CSE 511: Operating Systems Design

Lectures 12-13
Virtualization

Virtualization

- Allows multiple instances of an operating system to run on a single computer
 - Originally introduced for IBM VM/370
 - Has later been revitalized for modern platforms
- A **hypervisor** is a software layer that
 - Separates the **virtual** hardware an OS sees from the **actual** hardware
 - Arbitrates access to physical resources such as CPU or memory

Virtualization

- Widely known hypervisors
 - Xen, KVM, VMware ESX, Hyper V, VirtualBox, qemu etc
- Virtualization
 - Improves **isolation** and **reliability** because each OS runs independently from the others on its own **virtual** processors so that OS failures are **contained**
 - Increases **resource utilization** as the same hardware can be used for multiple purpose
 - Leads to better **productivity** as large pieces of software can be preconfigured and installed very easily



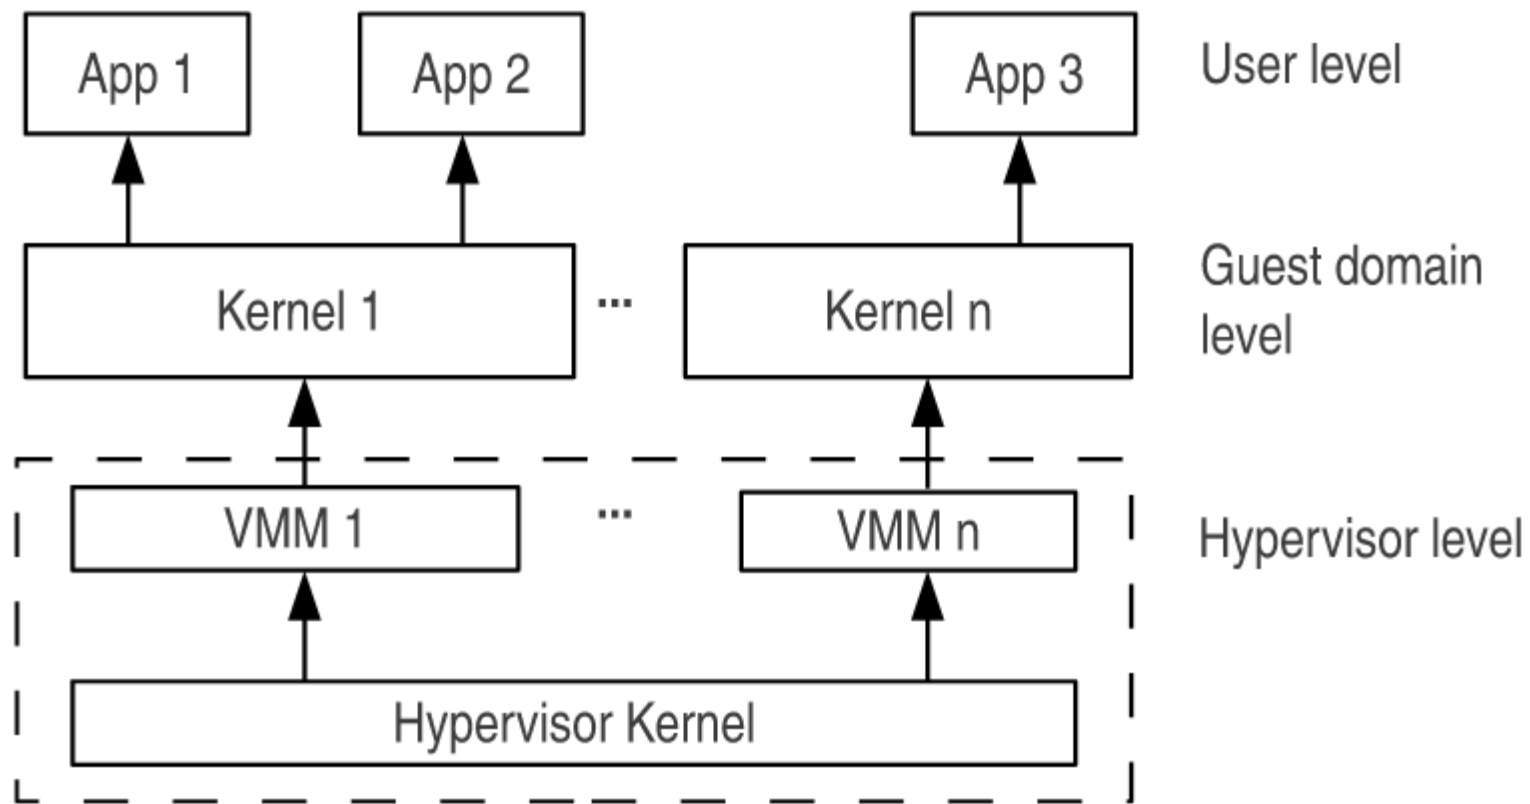
Virtualization

- Virtual appliances
 - Software bundles containing preconfigured packages with their own OS aimed at simplified distribution and deployment
 - They run along with other virtual appliances and general purpose OS on a single machine
- Infrastructure as a service (IaaS) solutions
 - Rely on virtualization to provide business solutions for server consolidation
- Hypervisors can be utilized to provide better security and robustness for operating systems

Hypervisor Types

- Type I
 - Allows execution of virtual machines in **guest** domains
- Type II
 - Runs on top of some OS, known as a **host** OS
- Type I hypervisors avoid unnecessary overheads, latencies; they are also not subject to host OS vulnerabilities
 - The hypervisor is at the lowest layer: the hypervisor kernel and Virtual Machine Monitors (VMMs)
 - The hypervisor kernel has direct access to hardware, resource allocation
 - The VMM layer is responsible for virtualizing resources

Hypervisor Types



CPU Virtualization

- A hypervisor needs to provide guests with ***virtual*** CPUs
 - Entities that allow guests to have access to CPU resources while allowing the hypervisor to schedule different operating systems on the available ***physical*** CPUs or cores
- Complete and transparent CPU virtualization traditionally required ***architectural support***

CPU Virtualization

- Popek and Goldberg formulated necessary and sufficient conditions for virtualization
 - ***Fidelity***: requires that the virtual environment in which programs run should be indistinguishable from the real hardware
 - ***Performance***: a substantial amount of instructions need to be executed directly by hardware and without any additional hypervisor handling
 - ***Safety***: requires that the hypervisor has full control of all system resources

CPU Virtualization

- ***Sensitive instructions*** used by guests need to be trapped and emulated if they could affect the correctness of the hypervisor's behavior
 - ***Trap-and-emulate*** is a technique for intercepting CPU instructions and executing them using special (hypervisor) handlers
- The original x86 architecture does not meet described criteria, as it does not allow to intercept ***all*** necessary instructions when executed in ***unprivileged*** mode

CPU Virtualization

- Techniques to overcome this problem
 - **Binary translation:** machine code is being processed by a translator program, which allows prefetching, inspecting and special treatment of all relevant instructions (e.g., VMware)
 - **Paravirtualization:** adaptations of the guest OS kernel code that avoids the use of untrappable instructions, reducing emulation and management costs and yielding better performance (e.g., Xen)

CPU Virtualization

- Xen's original paravirtualization relied on the 4-ring protection hierarchy present in 32-bit x86 processors
 - Ring 0: privileged hypervisor code
 - Ring 1: guest kernel
 - Ring 3: user processes
- Segmentation protected guest kernel code and data from user processes, but guest kernels had to be adapted to successfully run in this mode
 - x86-64 effectively removed segmentation

CPU Virtualization

- The original paravirtualization approach is not so practical for x86-64
 - Need to run the kernel in Ring 3 and switch page tables every single system call
- But CPU vendors introduced **hardware extensions** (Intel VT-x, AMD-V) that allowed the execution of unchanged guest kernels
 - Safe virtualization of the CPU by introducing a VMM mode distinct from the mode in which privileged guest kernel code executes

CPU Virtualization

- Later generations added support for MMU virtualization via ***nested paging*** without resorting to ***shadow page tables***
- Well known hypervisors such as Xen and VMware now provide support for hardware virtualization
- Xen modes
 - ***PV***: paravirtualization
 - ***HVM, PVHVM***: using hardware extensions
 - ***PVH***: in between

CPU Virtualization

Poor Performance
 Scope for Improvement
 Optimal Performance

PV = Paravirtualized
 VS = Software Virtualized (QEMU)
 VH = Hardware Virtualized
 HA = Hardware Accelerated



x86 Shortcut	Mode	With					
HVM / Fully Virtualized	HVM		VS	VS ¹	VS	VH	Yes
HVM + PV drivers	HVM	PV Drivers Installed	PV	VS ¹	VS	VH	Yes
PVHVM	HVM	PVHVM Capable Guest	PV	PV ²	VS	VH	Yes
PVH	PVH	PVH Capable Guest	PV	HA ³	PV ⁴	VH	No
PV	PV		PV	PV	PV ⁵	PV	No
ARM							
N/A	N/A		PV	VH	PV ⁶	VH	No

* The picture is taken from
https://wiki.xenproject.org/wiki/Xen_Project_Software_Overview

Memory Virtualization

- Conventional OS manage memory assuming that they have unrestricted control over physical memory
- Hypervisors introduce special adaptations to guest OS allowing to run multiple guest domains at the same time
 - *Virtual*, *physical*, and *machine* memory address space
 - *Machine* memory is the actual memory exposed by hardware



Memory Virtualization

- Physical and virtual memory are guest specific and have their conventional meaning
 - Physical addresses correspond to memory exposed by the VMM rather than hardware
- Hypervisors and CPUs provide support for virtual-to-physical, physical-to-machine and virtual-to-machine address translation
 - Page granularity: hypervisors and guests refer to the physical and machine addresses using their physical and machine frame numbers

Memory Virtualization

- Fully-virtualized guests are completely unaware of the underlying hypervisor
 - They know nothing about machine frames and treat physical frames as if they represented actual hardware exposed memory
- Guests create page tables containing virtual-to-physical memory mappings
 - Cannot be used by the hardware since physical frames do not correspond to actual memory location

Memory Virtualization

- The hypervisor traps any attempt to load page tables and replace them with ***shadow page tables***
 - ***Shadow page tables*** are created and maintained by the hypervisor
 - Entries are created on demand
 - Contain virtual-to-machine memory mappings generated by the hypervisor from guest page tables and the physical-to-machine (P2M) hypervisor translation tables
 - Transparent (not exposed to guests)

Memory Virtualization

- CPU vendors also introduced *nested page tables*
 - Eliminates shadow page tables by providing an additional hardware translation layer
 - The first layer is for regular OS page tables containing virtual-to-physical mappings
 - The second layer is for physical-to-machine mappings and is managed by the hypervisor

Memory Virtualization

- Paravirtualization
 - Guests are adapted to assist the hypervisor with memory management
 - They have direct access to the P2M translation table, allowing them to perform physical-to-machine translation
 - No shadow page tables: guests create page tables with virtual-to-machine mappings
 - The hypervisor will verify the consistency of page tables