

CSE 511: Operating Systems Design



Lectures 17,19

Concurrent Objects

The slides are partially based on Maurice Herlihy's slides from "The Art of Multiprocessing"

Concurrency



Thread





Thread





Concurrent Computation





FIFO Queue: Enqueue Method





FIFO Queue: Dequeue Method



Lock-Based Queue




Lock-Based Queue



A Lock-Based Queue


```
int Head, Tail;  
T items[capacity];  
Lock lock;  
  
mutex_init(&lock);
```



A Lock-Based Queue



```
int Head, Tail;  
T items[capacity];  
Lock lock;  
  
mutex_init(&lock);
```

Variables protected by
single shared lock



Lock-Based Queue



Initially `head = tail`





A Lock-Based Queue

```
int Head = 0, Tail = 0;  
T items[capacity],  
Lock lock;  
mutex_init(&lock);
```



Initially head = tail



Lock-Based `deq()`




Acquire Lock





Implementation: `deq()`

```
T deq() {  
    mutex_lock(&lock);  
    if (tail == head) {  
        mutex_unlock(&lock);  
        /* Error! */  
    }  
    T x = items[Head % capacity];  
    Head++;  
    mutex_unlock(&lock);  
    return x;  
}
```

Acquire lock




Check if Non-Empty





Implementation: deq()

```
T deq() {  
    mutex_lock(&lock);  
    if (tail == head) {  
        mutex_unlock(&lock);  
        /* Error! */  
    }  
    T x = items[Head % capacity];  
    Head++;  
    mutex_unlock(&lock);  
    return x;  
}
```

Return error
if queue is empty




Modify the Queue




Implementation: `deq()`

```
T deq() {  
    mutex_lock(&lock);  
    if (tail == head) {  
        mutex_unlock(&lock);  
        /* Error! */  
    }  
    T x = items[Head % capacity];  
    Head++;  
    mutex_unlock(&lock);  
  
    return x;  
}
```





Queue not empty?
Remove item and update head



Implementation: `deq()`

```
T deq() {  
    mutex_lock(&lock);  
    if (tail == head) {  
        mutex_unlock(&lock);  
        /* Error! */  
    }  
    T x = items[Head % capacity];  
    Head++;  
    mutex_unlock(&lock);  
    return x;  
}
```



Return result





Implementation: `deq()`

```
T deq() {  
    mutex_lock(&lock);  
    if (tail == head) {  
        mutex_unlock(&lock);  
        /* Error! */  
    }  
    T x = items[Head % capacity];  
    Head++;  
    mutex_unlock(&lock);  
    return x;  
}
```



Return result



Release the Lock




Release the Lock




Implementation: `deq()`

```
T deq() {  
    mutex_lock(&lock);  
    if (tail == head) {  
        mutex_unlock(&lock);  
        /* Error! */  
    }  
    T x = items[Head % capacity];  
    Head++;  
    mutex_unlock(&lock);  
    return x;  
}
```




Release lock no matter what!



Implementation: `deq()`

```
T deq() {  
    mutex_lock(&lock);  
    if (tail == head) {  
        mutex_unlock(&lock);  
        /* Error! */  
    }  
    T x = items[Head % capacity];  
    Head++;  
    mutex_unlock(&lock);  
  
    return x;  
}
```

Should be correct because
modifications mutually exclusive...




An alternative implementation

The same code without mutual exclusion

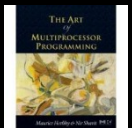

For simplicity, only *two* threads

One thread enq only



The other deq only





Wait-free 2-Thread Queue



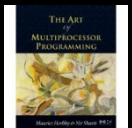

Wait-free 2-Thread Queue



Wait-free 2-Thread Queue



Wait-free 2-Thread Queue




Wait-free 2-Thread Queue

```
int head = 0, tail = 0;
Item items[capacity];

void enq(Item x) {
    if (tail-head == capacity) {
        /* Error */
    }
    items[tail % capacity] = x; tail++;
}

Item deq() {
    if (tail == head) {
        /* Error */
    }
    Item item = items[head % capacity]; head++;
    return item;
}
```

No lock needed!




Wait-free 2-Thread Queue

```
int head = 0, tail = 0;
Item items[capacity];

void enq(Item x) {
    if (tail-head == capacity) {
        /* Error */
    }
    items[tail % capacity] = x;
    tail++;
}

Item deq() {
    if (tail-head == 0) {
        /* Error */
    }
    Item item = items[head % capacity]; head++;
    return item;
}
```


How do we define “correct” when modifications are not mutually exclusive?



What *is* a Concurrent Queue?

Need a way to *specify* a concurrent queue object

Need a way to prove that an algorithm implements the object's specification




Correctness and Progress

In a concurrent setting, must specify both *safety* and *liveness* properties

Need a way to define

when an implementation is *correct*

When it guarantees *progress*



Sequential Objects

Each object has a *state*


Usually given by a set of *fields*

Queue example: sequence of items

Each object has a set of *methods*

Only way to manipulate state

Queue example: *enq* and *deq* methods



Sequential Specifications

If (precondition)

the object is in such-and-such a state

before you call the method,


Then (postcondition)

the method will return a particular value

or return an error.

and (also postcondition)

The object will be in some other state



Pre and Post Conditions for Dequeue

Precondition

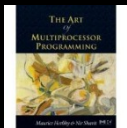
Queue is non-empty

Postcondition

Returns first item in queue

Postcondition

Removes first item in queue



Pre and Post Conditions for Dequeue

Precondition


Queue is empty

Postcondition

Returns an error

Postcondition

Queue state unchanged




Why Are Sequential Specifications Useful?

Interactions among methods captured by side-effects on object state

Documentation size linear in number of methods

Can add new methods without combinatorial blow-up




What About Concurrent Specifications ?



Methods?

Documentation?



Adding new methods?



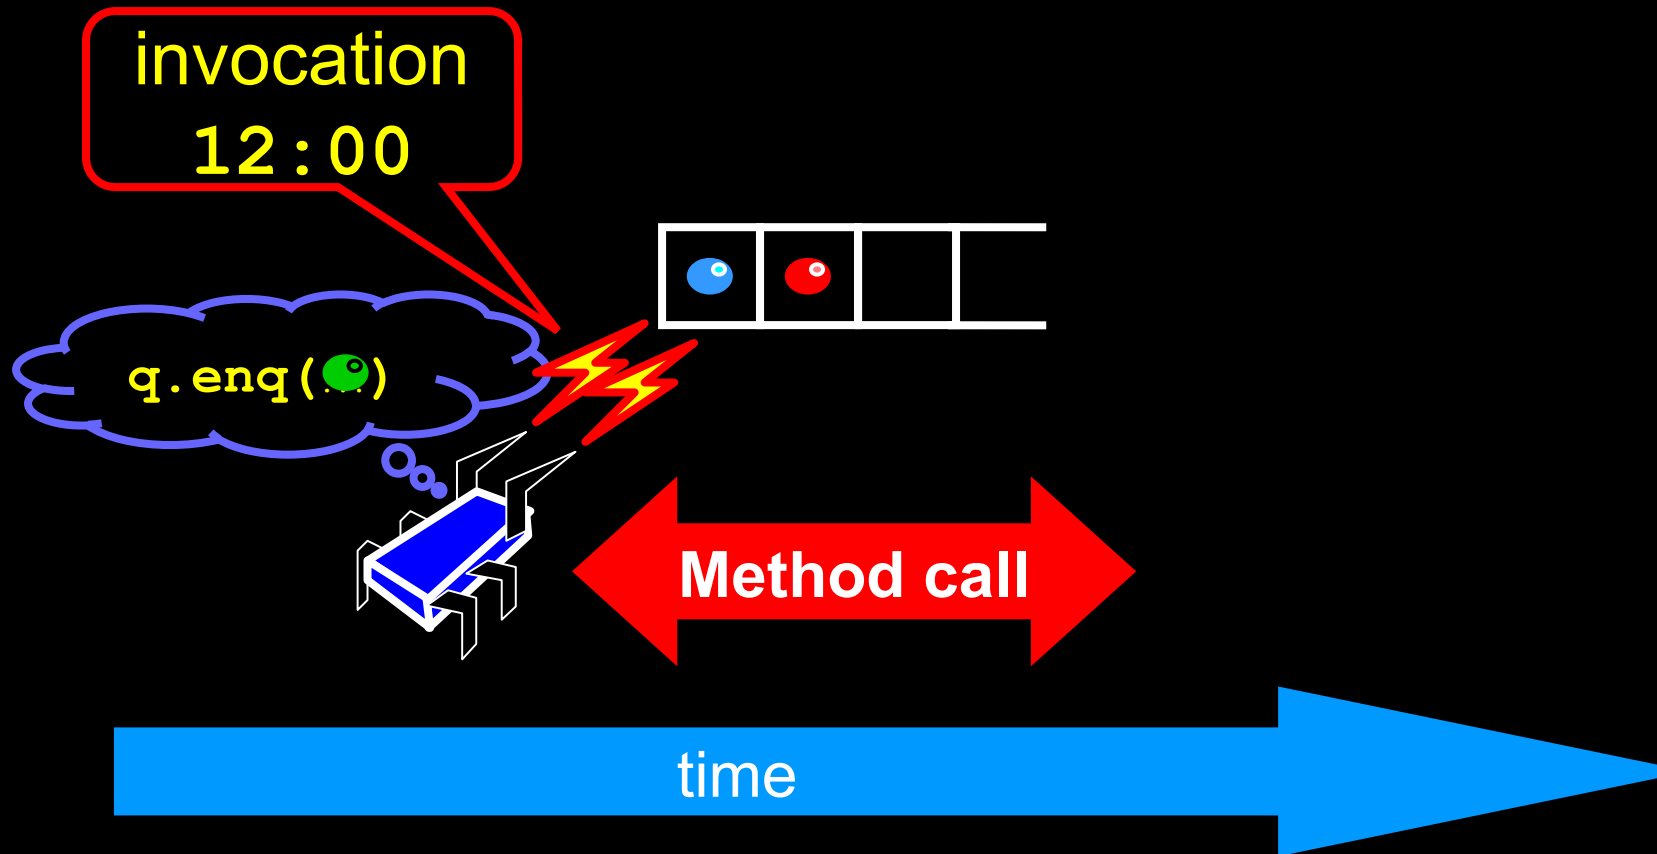
Methods Take Time





Methods Take Time





Methods Take Time



Methods Take Time



Methods Take Time



Sequential vs Concurrent


Sequential

Method calls take time? Who knew?


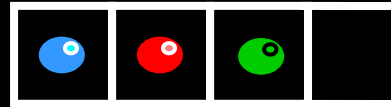
Concurrent

Method call is not an *event*



Method call is an *interval*.



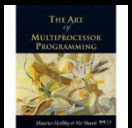

Concurrent Methods Take *Overlapping* Time




Concurrent Methods Take *Overlapping* Time



Concurrent Methods Take *Overlapping* Time



Concurrent Methods Take *Overlapping* Time




Sequential vs Concurrent

Sequential

Object needs meaningful state only
between method calls

Concurrent

Because method calls overlap, object
might *never* be between method calls



Sequential vs Concurrent

Sequential


Each method described in isolation

Concurrent

Must characterize *all* possible interactions with concurrent calls

What if two **enq** () calls overlap?

Two **deq** () calls? **enq** () and **deq** () ? ...



Sequential vs Concurrent


Sequential

Can add new methods without affecting older methods' specs

Concurrent

Adding a method can potentially rewrite every other method's spec

Panic!




Today's Big Question

What does it even *mean* for a concurrent object to be correct?

What *is* a concurrent FIFO queue?


FIFO means strict temporal order

Concurrent means ambiguous temporal order



Intuitively...


```
void deq()
{
    mutex_lock(&lock);
    if (tail == head) {
        /* Error */
        mutex_unlock(&lock);
    }
    T x = items[head % capacity];
    head++;
    mutex_unlock(&lock);
    return x;
}
```



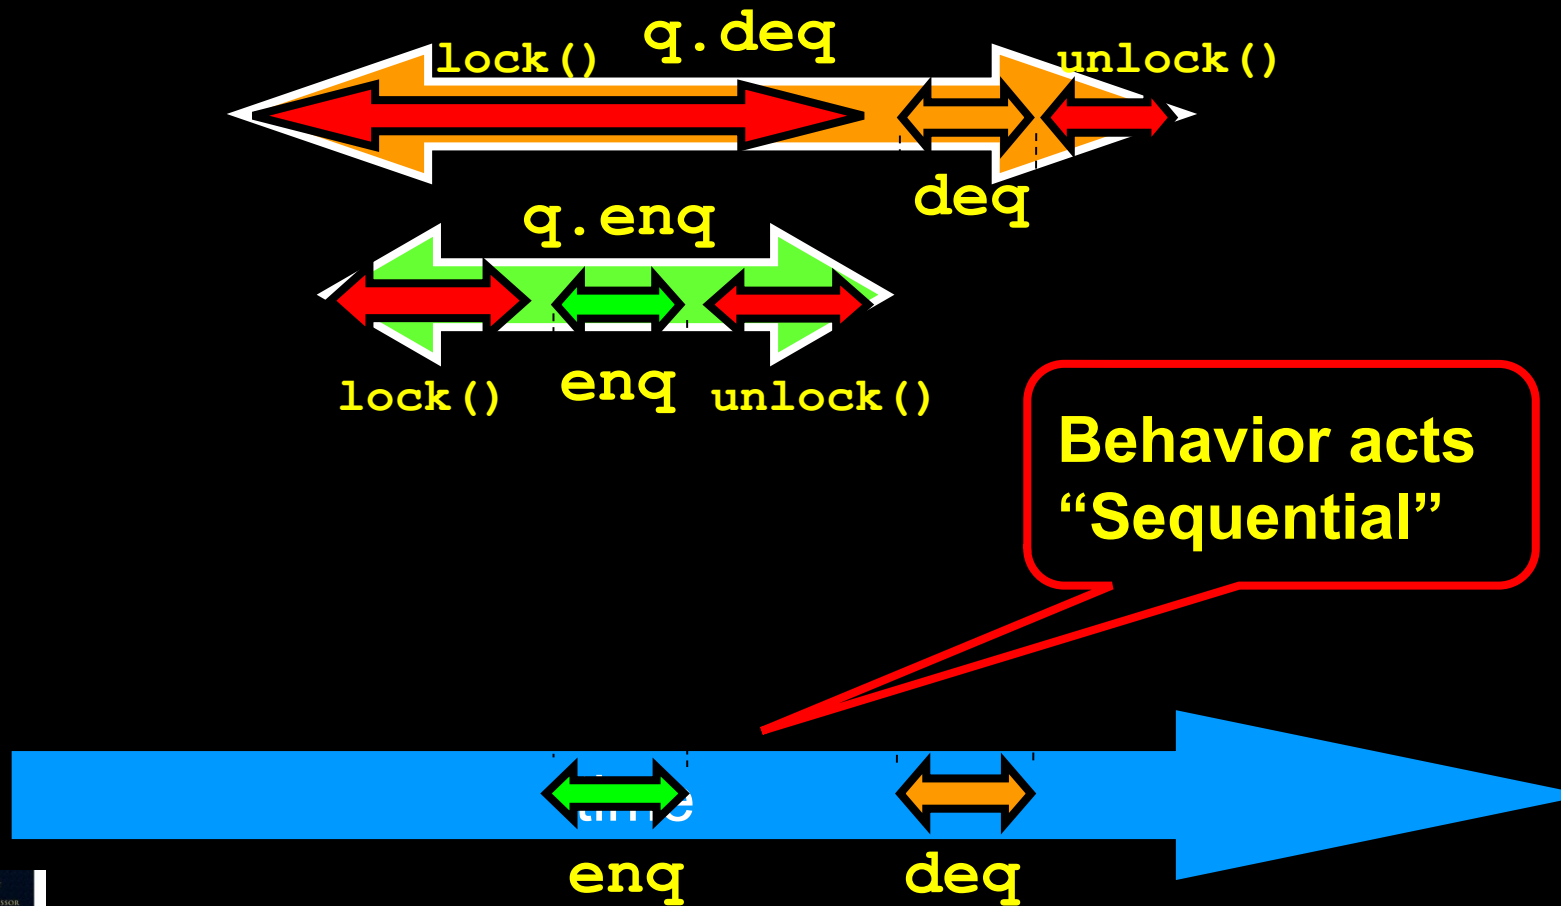
Intuitively...

```
void deq()
{
    mutex_lock(&lock);
    if (tail == head) {
        /* Error */
        mutex_unlock(&lock);
    }
    T x = items[head % capacity];
    head++;
    mutex_unlock(&lock);
    return x;
}
```

All queue modifications
are mutually exclusive



Intuitively



Linearizability

Each method call should


“take effect”

Instantaneously

Between invocation and response events

Object is correct if “sequential” behavior is correct

Any such concurrent object is *linearizable*



Linearizability

Each method call should

“take effect”

Instantaneously

Between invocation and response events

Actually a property of an execution

A linearizable object: one all of whose possible executions are linearizable

