

CSE 511: Operating Systems Design

Lectures 21

Linearizability, Sequential Consistency Progress Conditions

Sequential Histories

A q.enq(3)

A q:void

B p.enq(4)

B p:void

B q.deq()

B q:3


A q:enq(5)

Method calls of different
threads do not interleave

match

match

Final pending
invocation OK



Well-Formed Histories

Per-thread projections sequential

$H =$


```
A q.enqueue(3)
B p.enqueue(4)
B p:void
B q.dequeue()
A q:void
B q:3
```

$H | B =$

```
B p.enqueue(4)
B p:void
B q.dequeue()
B q:3
```

$H | A =$

```
A q.enqueue(3)
A q:void
```



Equivalent Histories


Threads see the same thing in both $\left\{ \begin{array}{l} H|A = G|A \\ H|B = G|B \end{array} \right.$

H=

```
A q.enq(3)
B p.enq(4)
B p:void
B q.deq()
A q:void
B q:3
```

G=

```
A q.enq(3)
A q:void
B p.enq(4)
B p:void
B q.deq()
B q:3
```



Sequential Specifications

A sequential specification is some way of telling whether


A single-thread, single-object history

Is legal

For example:

Pre and post-conditions

But any style will do ...




Legal Histories

A sequential (multi-object) history H is *legal* if



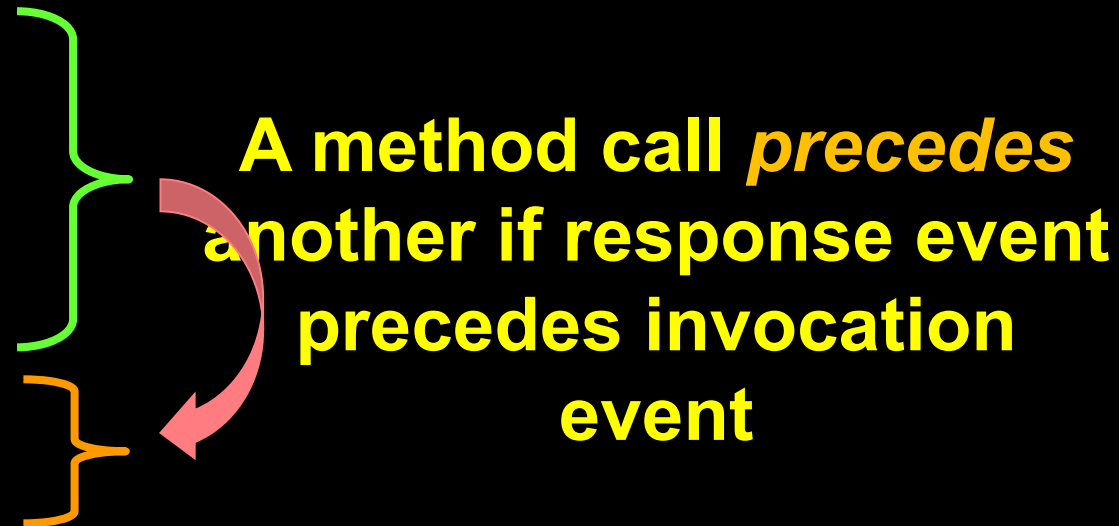
For every object x

$H|x$ is in the sequential spec for x




Precedence

```
A q.enq(3)
B p.enq(4)
B p.void
A q:void
B q.deq()
B q:3
```




Concurrency

```
A q.enq(3)
B p.enq(4)
B p.void
B q.deq()
A q:void
B q:3
```



**Some method calls
overlap one another**



Linearizability

History H is *linearizable* if it can be extended to G by


Appending zero or more responses to pending invocations

Discarding other pending invocations

So that G is equivalent to

Legal sequential history S

where S respects “real-time order” of G

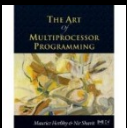


Remarks

Some pending invocations

Took effect, so keep them


Discard the rest




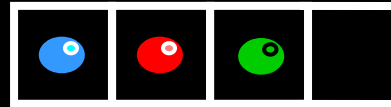
Linearizability: Summary

Standard notion of correctness



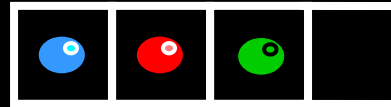
Captures the notion of objects being “atomic”





Sequential Consistency: Example



Example


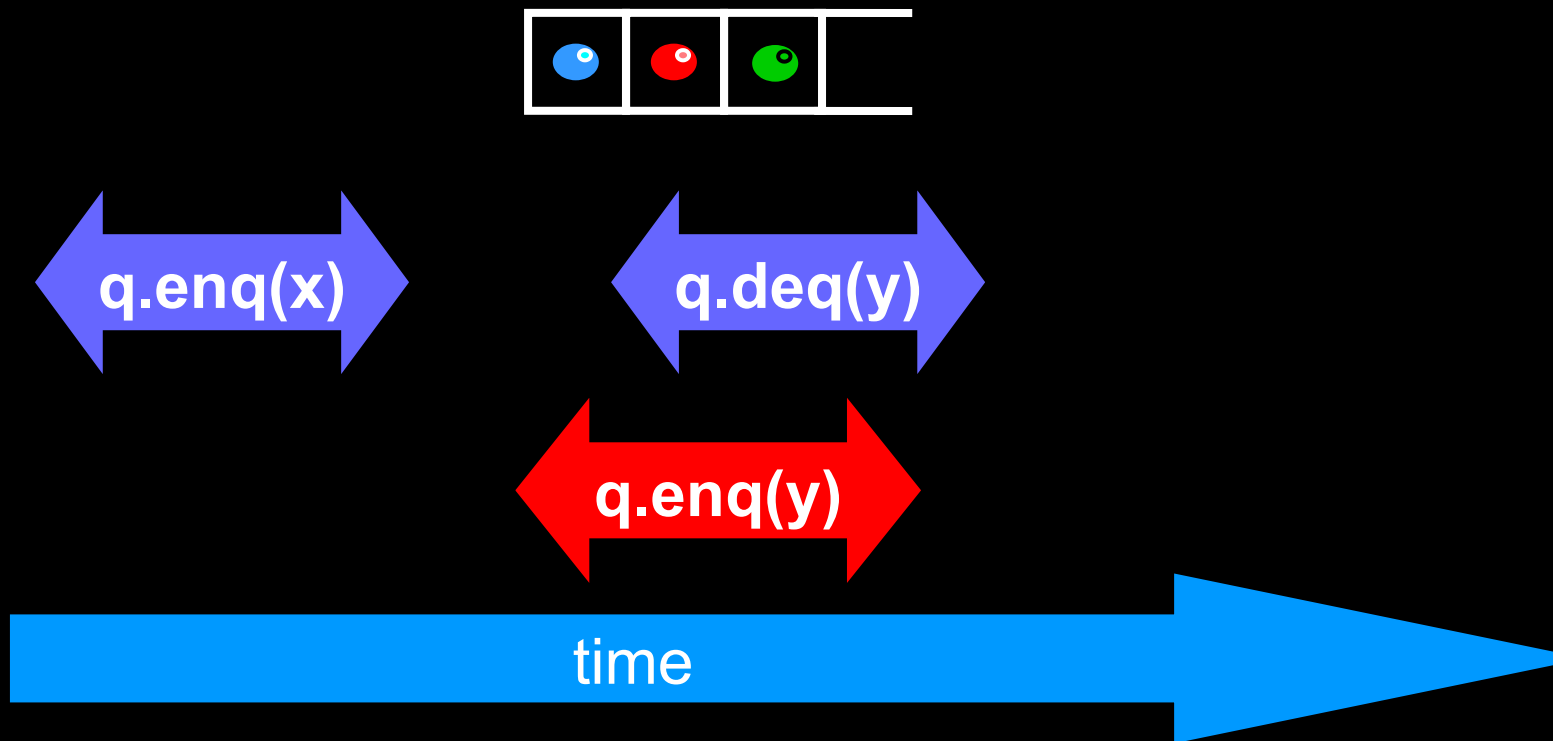


Example




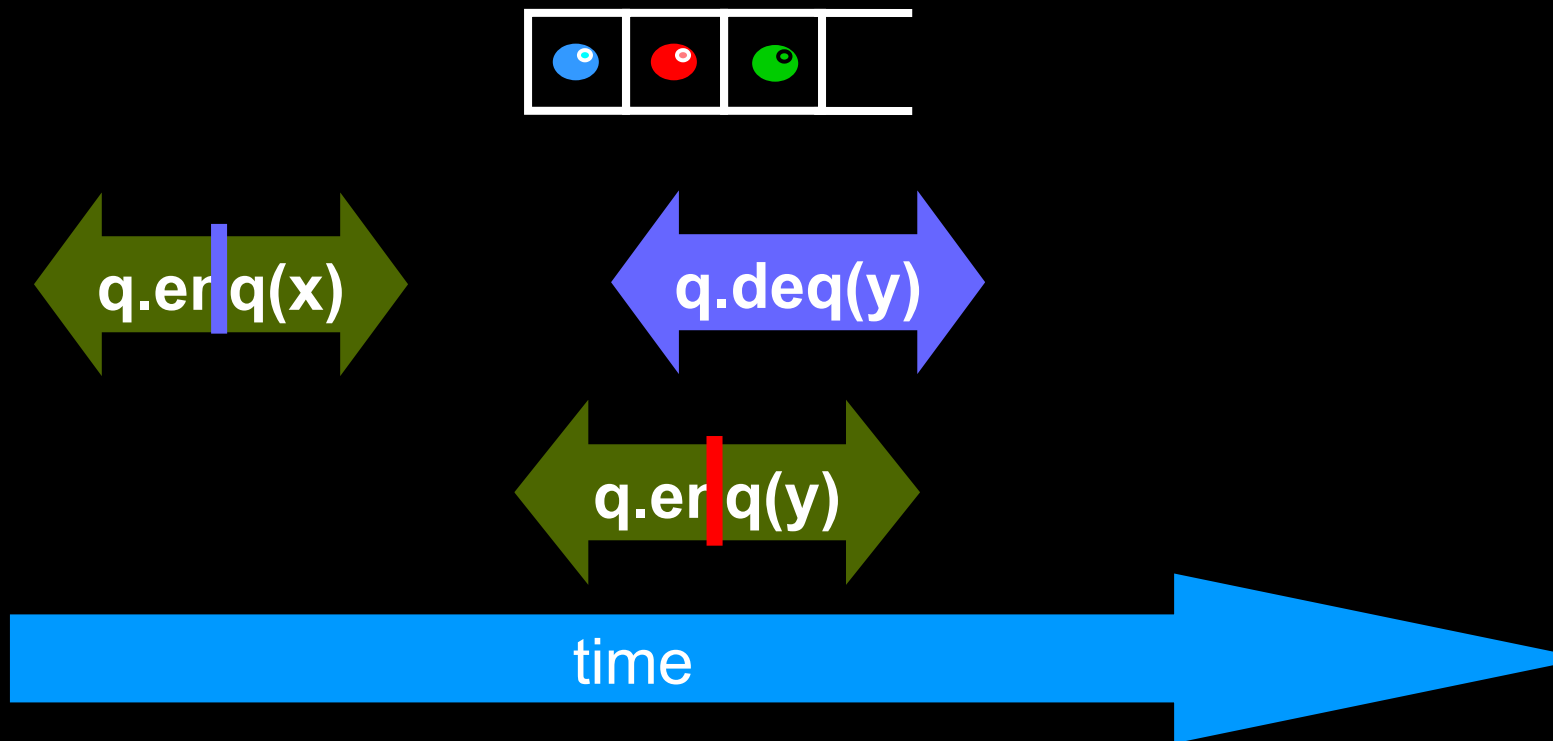


Example









Example





Example





not linearizable





Exa

Yet Sequentially
Consistent




Sequential consistency

Most hardware does not even support that, architects think that sequential consistency is *too strong*

Too expensive for modern hardware

OK if violated by *default*

Honored by *explicit request*



Hardware Consistency

Initially, $a = b = 0$.

Core 0

```
mov 1, a    ;Store  
mov b, %ebx ;Load
```

Core 1

```
mov 1, b    ;Store  
mov a, %eax ;Load
```

What are the possible values of `%eax` and `%ebx` registers after both processors have executed?

Sequential consistency implies that no execution ends with $\%eax = \%ebx = 0$



Hardware Consistency

No modern-day processor implements sequential consistency

Hardware actively reorders instructions.

Compilers too

Because performance is dominated by single-thread unsynchronized execution



Instruction Reordering

```
mov 1, a    ;write  
mov b, %ebx ;read
```

```
mov b, %ebx ;read  
mov 1, a    ;write
```

Program Order

Execution Order

Q. Why might the hardware or compiler decide to reorder these instructions?

A. Higher performance by covering load latency — *instruction-level parallelism*.



Slide used with permission of
Charles E. Leiserson

Instruction Reordering

```
mov 1, a    ;write  
mov b, %ebx ;read
```

```
mov b, %ebx ;read  
mov 1, a    ;write
```

Program Order

Execution Order

Q. When is it safe for the hardware or compiler to perform this reordering?

A. When $a \neq b$.

A'. And there's no concurrency.



Slide used with permission of
Charles E. Leiserson

X86: Memory Consistency – TSO (Total Store Ordering)

Thread
Code

~~Store1~~
~~Store2~~
Load1
~~Load2~~
~~Store3~~
Store4
Load3
~~Load4~~
~~Load5~~

Loads are *not* reordered with loads.

Stores are *not* reordered with stores.

Stores are *not* reordered with prior loads.

A load *may* be reordered with a prior store to a different location but *not* with a prior store to the same location.

Stores to the same location respect a global total order.



Memory Barriers (Fences)

A *memory barrier* (or *fence*) is a hardware operation that enforces ordering between the instructions before and after the fence

A memory barrier can be an explicit instruction (x86: mfence)

The typical cost of a memory fence is comparable to that of an L2-cache access



X86: Memory Consistency

Thread's Code

Store1
Store2
Load1
Load2
Store3
Store4
Barrier
Load3
Load4
Load5

Loads are *not* reordered with loads.

Store to the same location but not
with a prior store to the same
location.

**Total Store Ordering +
properly placed memory
barriers = sequential
consistency**

Stores to the same location respect a
global total order.



Memory Barriers


Memory barrier instruction will

Flush write buffer

Bring caches up to date

Compilers often do this for you

Entering and leaving critical sections




Summary: Real-World

Hardware *weaker* than sequential consistency

Can get sequential consistency at a price

Linearizability better fit for high-level software




Linearizability

Operation takes effect instantaneously between invocation and response

Based on sequential spec

“gold standard” for concurrent data structures




Progress Conditions

Deadlock-free: some thread trying to acquire the lock eventually succeeds

Starvation-free: every thread trying to acquire the lock eventually succeeds.


Lock-free: some thread calling a method eventually returns.

Wait-free: every thread calling a method eventually returns.



Progress Conditions

	Non-Blocking	Blocking
Everyone makes progress	Wait-free	Starvation-free
Someone makes progress	Lock-free	Deadlock-free



Kinds of Architectures

SISD (Uniprocessor)

Single instruction stream

Single data stream

SIMD (GPU)

Single instruction

Multiple data

MIMD (Multiprocessors)

Multiple instruction

Multiple data



Kinds of Architectures

SISD (Uniprocessor)

Single instruction stream

Single data stream

SIMD (GPU)

Our space

Single instruction

Multiple data

MIMD (Multiprocessors)

Multiple instruction

Multiple data



MIMD Architectures



Shared Bus

Distributed

Memory Contention

Communication Contention

Communication Latency

