

CSE 511: Operating Systems Design

Lectures 24

Lock-free data structures

Exercise

- Modify the lock-free stack implementation using ABA tags

- Can use something like:

```
struct aba_ptr_s {  
    void *ptr;  
    size_t tag;  
};
```

```
aba_ptr_new = { .ptr = new_ptr, .tag = aba_ptr_old.tag + 1 };
```

```
CAS(&aba_ptr_global, aba_ptr_old, aba_ptr_new)
```

- Use two stacks
 - free_stack contains deallocated/free nodes
 - alloc_stack contains allocated nodes (normal stack)

Memory Reclamation

- Still **unsafe** to return memory from `free_stack` back to the OS when only using the ABA tags but can recycle past nodes
 - Why?
- Special approaches exist to solve the problem:
 - Epoch-based reclamation
 - Hazard pointers
 - Hazard eras (combines the above two schemes)
 - etc

Memory Reclamation

- Will use **epoch-based reclamation**

Enclose the code which access shared memory (it is very similar to lock() and unlock() but does not block)

start_op(thread_id); // Creates a thread reservation

...

end_op(thread_id); // Clears a thread reservation

Instead of free(), call **retire()**. free() will be called automatically when it is safe to do so.

Michael-Scott's Lock-Free Queue

```
struct node_s { // Queue Element
    struct node_s *next; // Next
    void *obj; // Associated Object
};

typedef struct node_s node_t;

// Head and Tail pointers
// Create one sentinel node initially
node_t* head = malloc(sizeof(node_t));
node_t* tail = head;
head->next = NULL;
```

Michael-Scott's Lock-Free Queue

```
void enqueue(void *obj) {
    node_t *node = malloc(sizeof(node_t);
    node->obj = obj;
    node->next = NULL;
    while (true) {
        node_t *t = LOAD(tail);
        node_t *next = t->next;
        if (t == LOAD(tail)) {
            if (next == NULL) {
                if (CAS(&t->next, next, node)) break;
            } else {
                CAS(&tail, t, next); // Help to move tail
            }
        }
    }
    CAS(&tail, t, next); // Try to move tail (one time)
}
```

Michael-Scott's Lock-Free Queue

```
void *dequeue() {  
    while (true) {  
        node_t *h = LOAD(head);  
        node_t *t = LOAD(tail);  
        node_t *next = h->next;  
I    if (h == LOAD(head)) {  
        if (h == t) {  
            if (next == NULL) return NULL; // Empty  
            CAS(&tail, t, next);  
        } else {  
            void *obj = next->value; // Read before the update  
            if (CAS(&head, h, next))  
                return obj; // Return the object  
        }  
    }  
}
```