
Operating Systems

CSE 511

Assignment 3 Directions

File Systems

Instructor: Ruslan Nikolaev

Assignment 3

- The document will be released tonight
 - You can already get started with the code <https://classroom.github.com/a/uBT6S-5T>
 - Academic integrity is very important! (do not use any online implementations)
 - Use our slides and official manuals!
 - You need to implement lock-based and lock-free queues and stacks
 - Use `pthread_mutex` for lock-based
 - Use C11 atomics for lock-free

Assignment 3

- Review pthread_mutex documentation
- Review C11 atomics (stdatomic.h)
 - atomic_compare_exchange_strong (CAS)
 - atomic_load (LOAD)
 - atomic_store (STORE)
 - You do not need atomic_compare_exchange_weak!

Assignment 3

- For ABA queues/stacks
 - Create two stacks/queues (e.g., store two head and two tail pointers for queues)
 - e.g., can create a separate structure with head/tail and then create 'aq' and 'fq' instances inside the queue object
 - Allows to use the same code for aq and fq
 - Recycle elements through 'fq'

Assignment 3

- [Hint] Stacks: **need tagging only for ‘top’**
 - *Not needed for ‘next’ (think why)*
- [Hint] Queues: **need tagging everywhere (head, tail, next)**
- When you need tagging, create a separate structure like this

```
typedef struct aba_ptr {  
    void *ptr; // or any other pointer type  
    size_t tag;  
} aba_ptr;
```

Assignment 3

- Just pass pointers/values of `aba_ptr` directly to `atomic_compare_exchange_strong`, `atomic_load`, `atomic_store` like you pass pointers/values to regular pointers
- Use `_Atomic(...)` to encapsulate custom atomic types, e.g., `_Atomic(queue_node *) head` or `_Atomic(aba_ptr) aba_head`.

Filesystem interface

- What is a file? **Named storage**
- File operations:
 - read(fd, buffer, length)
 - write(fd, buffer, length)
 - fd = open(pathname, flags)
 - close(fd)

**syscalls
(unbuffered)**
- fread(buffer, size, num, file)
 - fwrite(buffer, size, num, file)
 - file = fopen(pathname, mode)
 - fclose(file)

**Buffered
C library**

File descriptor (fd)

- Purpose: identify file being accessed
- What actually is a file descriptor?
 - 0-indexed number – index of file descriptor table

File descriptor	Referencing
0 (stdin)	Console
1 (stdout)	Console
2 (stderr)	Console
3	My_open_file.txt
...	...

- Important points:
 - Each process has its own table
 - There is a limit (default 1024)

Filesystems

- Purpose: making storage usable
 1. Tracking file information (metadata)
 2. Mapping files to blocks (i.e., LBN)
 3. Tracking/allocating free blocks

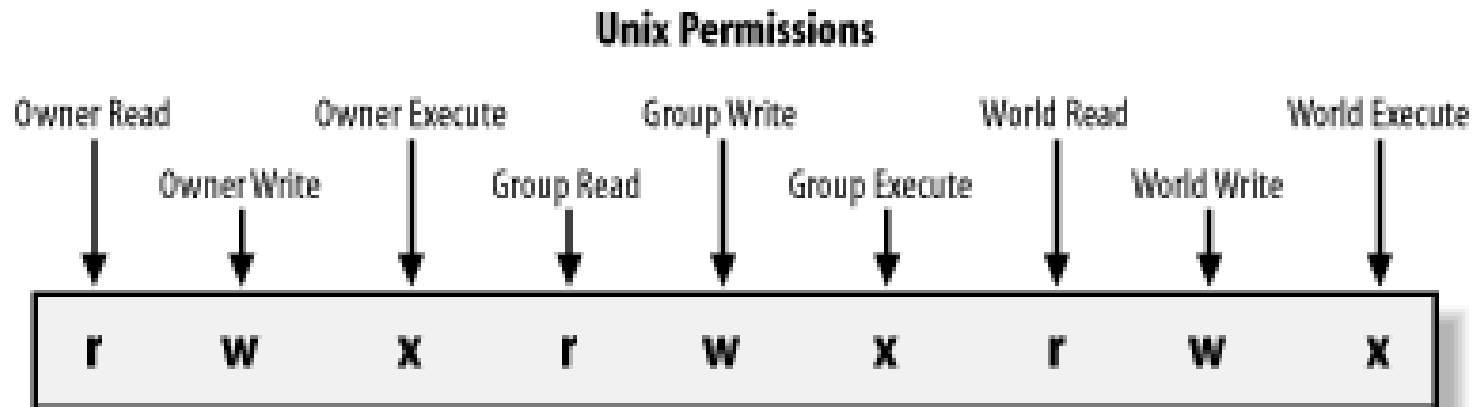
File metadata

- Metadata: data that's not the actual data
- Examples:
 - Filename
 - Size of file
 - Last access timestamp
 - Last modification timestamp
 - File owner
 - Permission
 - Location of data

} **inode**

File permission

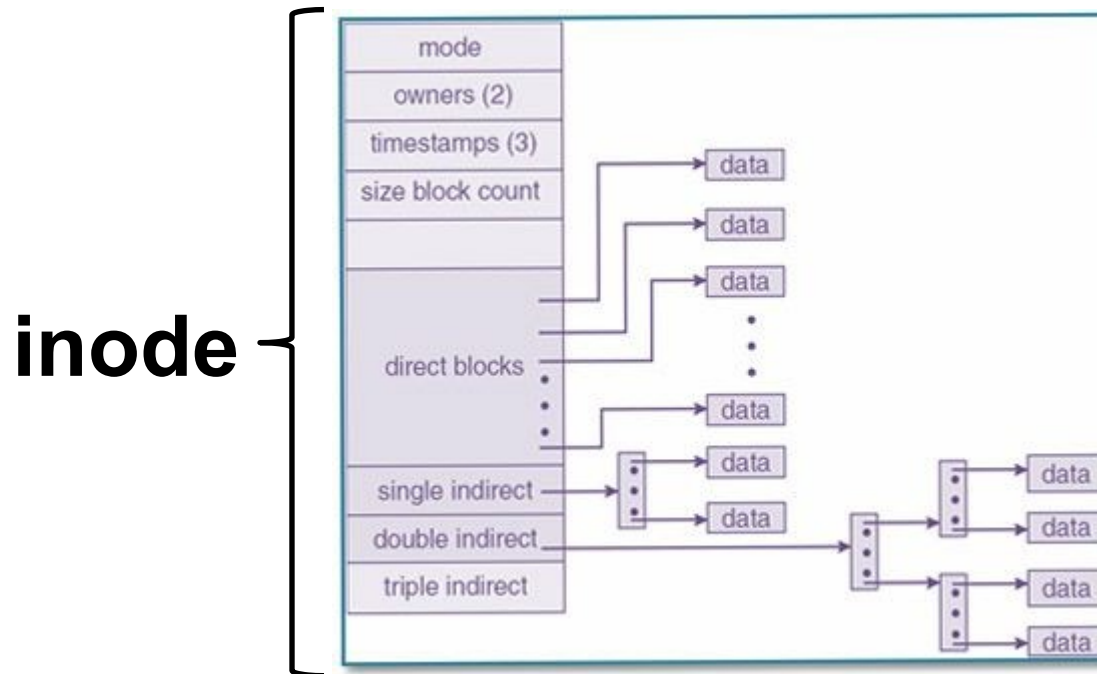
- Files have owners/groups
 - Use chown to modify owner/group
- File permission
 - Use chmod to modify permission



Mapping files to blocks (i.e., LBN)

- Simple: linked-list of data blocks (FAT)
- Better: indexed table (ext2/ext3)
- Problem: how big of a table?
- Solution: multi-level index

ext4 uses an
“extent tree” instead



Tracking/allocating free blocks

- Simple: free list (slow)
- Better: bitmap (easy to scan)
- Complex: list of extents (i.e., ranges of free blocks)