

CSE 511: Operating Systems

Spinlocks on Shared Memory Multiprocessors

Lectures 29-30

Lock Synchronization

Two common techniques for implementing synchronization

- Blocking
- Busy waiting (spin locks)

When do you choose one approach over another?

Recall: Assumed Hardware Support

- Atomic Instructions (e.g. test&set, fetch&add,...)
- Atomic test&set(x) {
 temp = x;
 x = BUSY;
 Return(temp);
 }
- Atomic fetch&Add(x,y) {
 temp = x;
 X = x + y;
 return(temp);
 }


Bus-based Shared Memory Multiprocessors

Load A, R0

Load A, R0

Store R1, A

Load A, R0




What about non-bus based systems?

1. Spin on test&set

```
Lock() {
    while (test&set(x) == BUSY)
        ;
}
```

```
Unlock() {
    x = CLEAR;
}
```



Problems:

- Each attempt creates network traffic
- Unlock also contends with trying for a lock


2. Spin on Read (test-test&set)

```

Lock() {
    while (x == BUSY || test&set(x) == BUSY)
        ;
}

Unlock() {
    x = CLEAR;
}

```



Problems:

- Gap between detecting the lock is clear and trying for it allows more people to try for it than required
- Cached copies get invalidated with further (possibly unsuccessful attempts)
- Time for network activity to settle down after lock is released (Quiesce Time)
- For reasonable performance: Critical Section \gg Quiesce Time

3. Delay after Spinning Processor Notices lock is released

```

Lock() {
    while (x == BUSY || test&set(x) == BUSY) {
        while (x == BUSY)
            ;
        Delay();
    }
}

```

```

Unlock() { x = CLEAR}

```

Problems:

- Delay based solutions are not necessarily the best solutions.
- How do you set the delays?
- In invalidation-based schemes, the local busy wait may still incur misses (even if unsuccessful) because of other attempts.

4. Delay between each memory reference (for invalidation based schemes)

```
Lock() {  
    while (x == BUSY || test&set(x) == BUSY) Delay();  
}
```

```
Unlock() { x = CLEAR; }
```

Problems:

- Most of the problems of previous solution persist.

5. Ticket Lock

```
int next_ticket, now_serving = 0;
```

```
Lock() {  
    int myticket;  
  
    myticket = Fetch&Add(next_ticket,1);  
    while (myticket != now_serving)  
        ;  
}
```

```
Unlock() { now_serving++ }
```

Problems:

- Everyone reading and writing next_ticket, now_serving (particularly the latter)

6. Array-based Queueing (Anderson Lock)

```
int Flags[0..P-1]; // Flags[0] = HAS_LOCK and rest are MUST_WAIT
Int tail = 0;
```

```
Lock() {
    int myplace = Fetch&Add(tail,1);
    while (Flags[myplace % P] == MUST_WAIT)
        ;
    Flags[myplace % P] = MUST_WAIT;
}
Unlock() { Flags[myplace+1 % P] = HAS_LOCK; }
```

Problems:

- Need to use arrays
- Need to worry about false sharing
- You have no control which position you will wait on. On a system without caches (which is not really the case anymore), this can be bad