

CSE 511: Operating Systems Design

Lectures 3,4

OS Boot Process

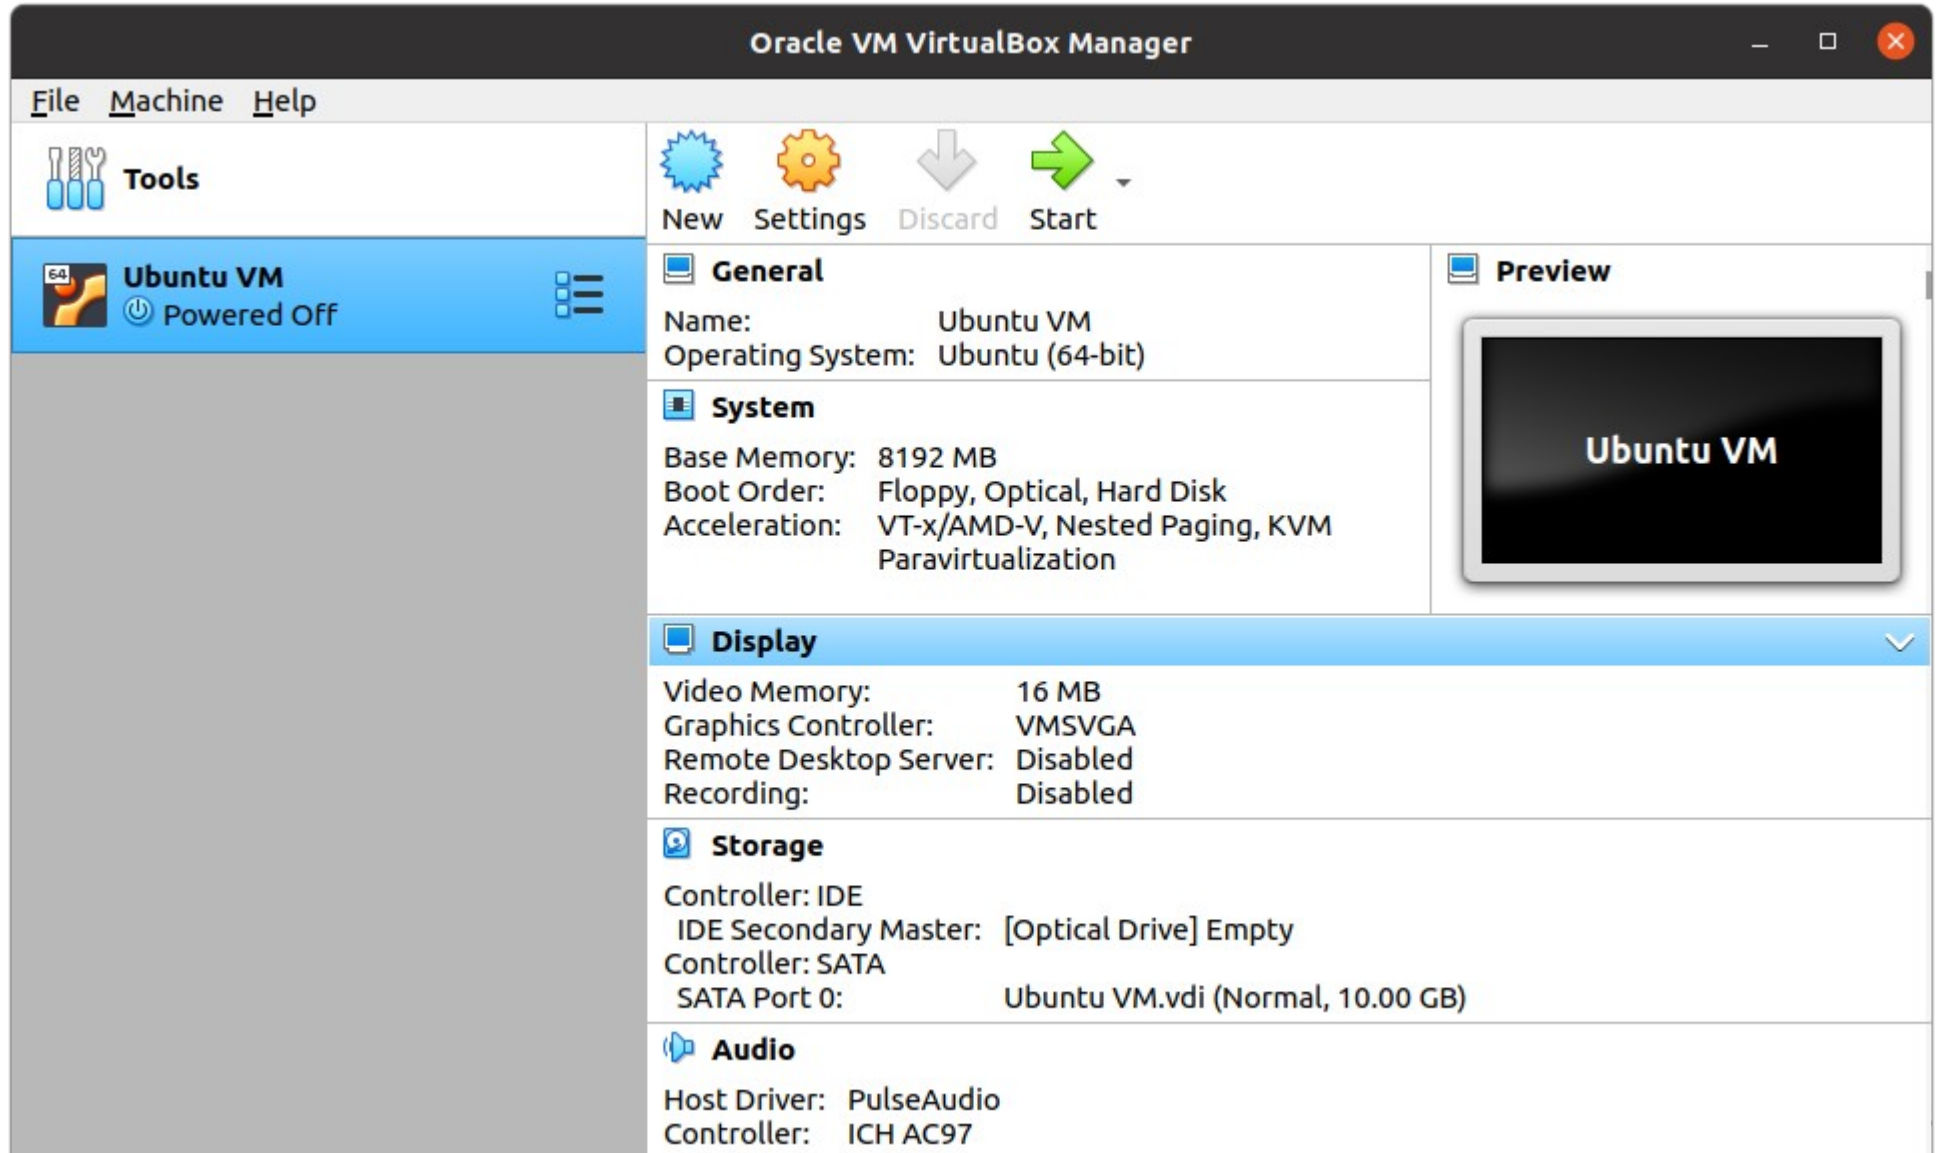
UEFI Programming

Assignment 1

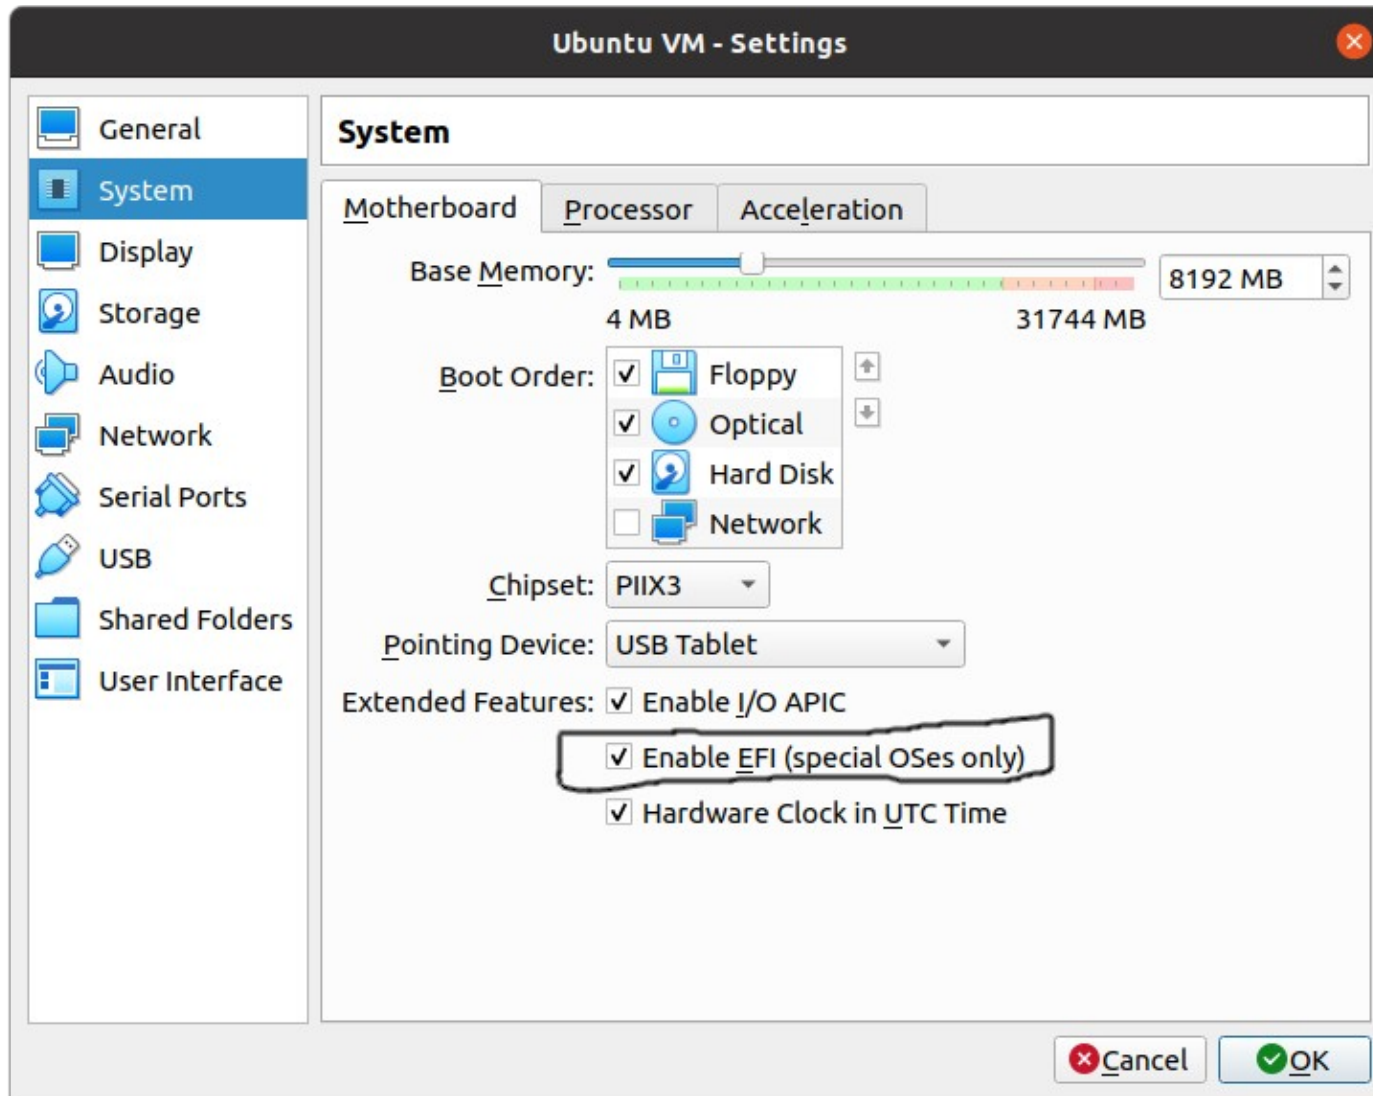
UEFI Firmware

- Used by all newer hardware but legacy (BIOS) boot is still supported
 - Legacy BIOS boot was supposed to be phased out completely in 2020. Cannot boot MS-DOS anymore.
 - Your machine may still use legacy boot!
- **Since all machines vary, and BIOS legacy boot is still widely used, we will use VirtualBox and qemu for our assignments/projects to make things simpler**
 - UEFI boot will work even for legacy machines
 - Ubuntu: `sudo apt-get install virtualbox`

UEFI Firmware



UEFI Firmware



UEFI is disabled by default.

Enable it in Settings!

UEFI Shell

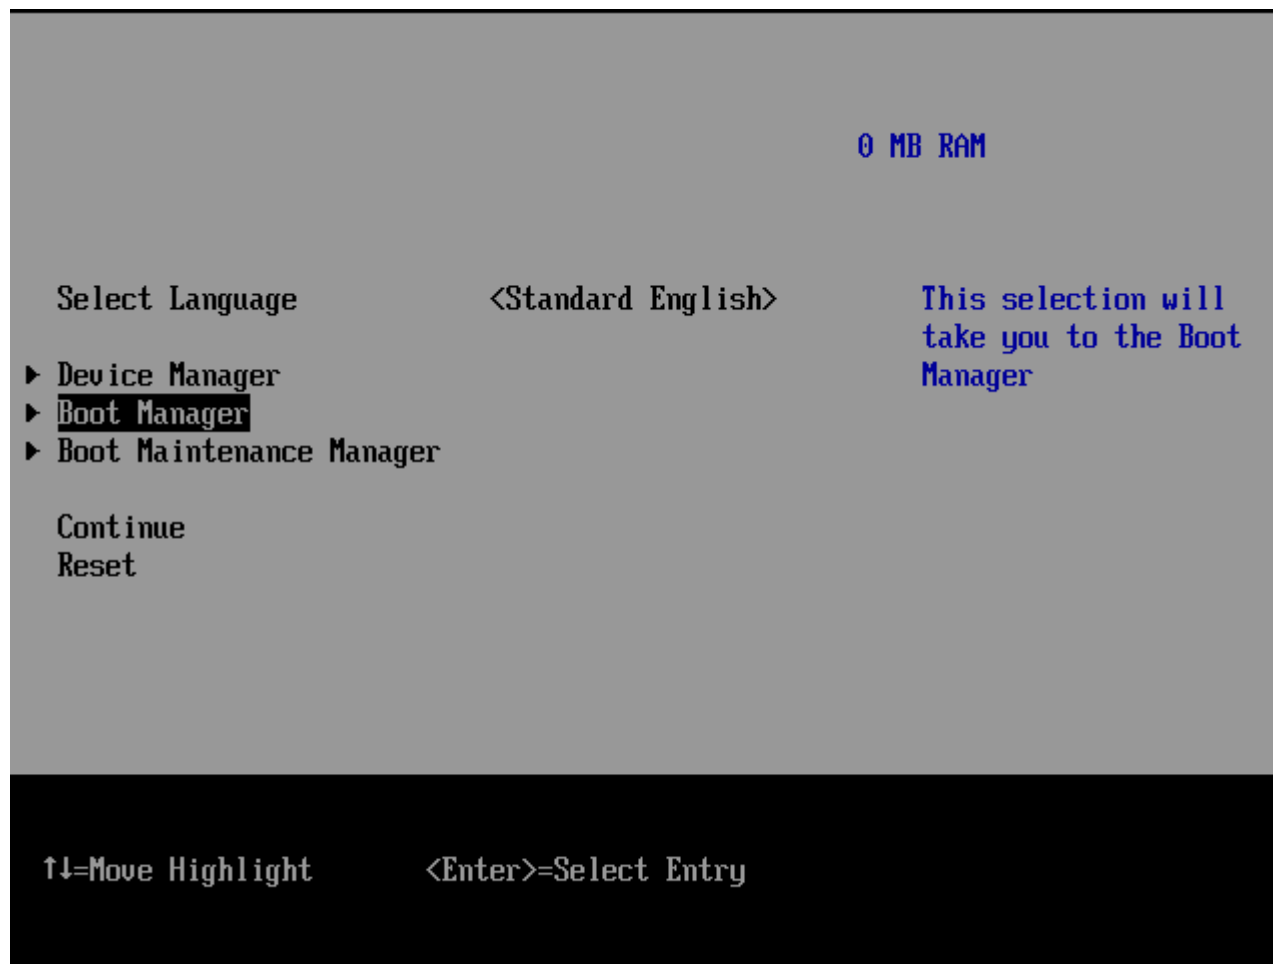
GNU GRUB version 2.04

```
Ubuntu
Ubuntu (safe graphics)
OEM install (for manufacturers)
Boot from next volume
*UEFI Firmware Settings
```

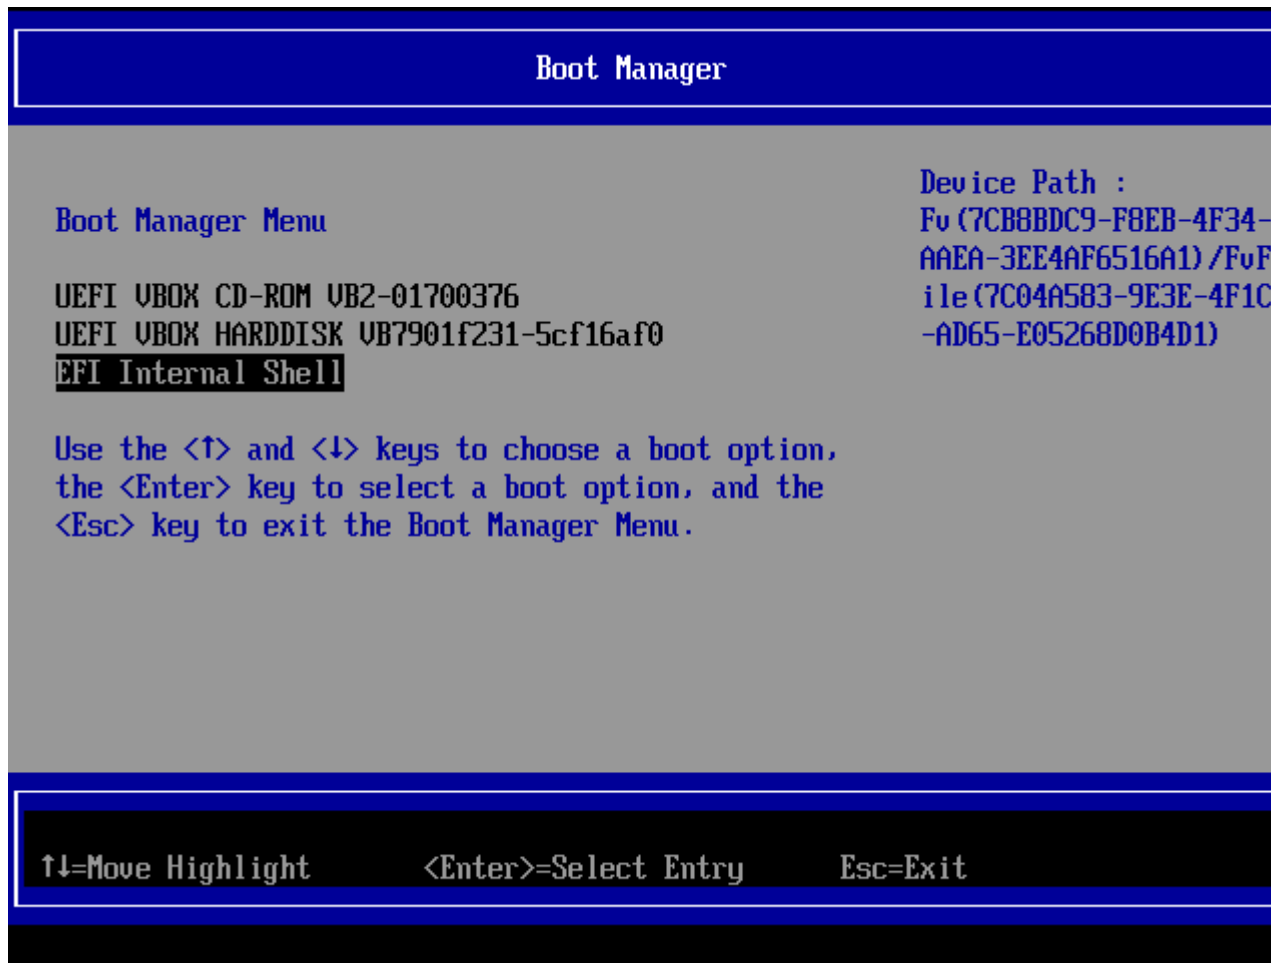
Use the ↑ and ↓ keys to select which entry is highlighted.
Press enter to boot the selected OS, `e` to edit the commands before booting or `c` for a command-line. ESC to return previous menu.

**Booting from
Ubuntu 20.04
ISO image**

UEFI Shell



UEFI Shell



List of bootable devices

UEFI Shell

```
FS0: Alias(s):F0c::BLK0:
PciRoot (0x0)/Pci (0x1,0x1)/Ata (0x0)
FS1: Alias(s):CD0c0::BLK2:
PciRoot (0x0)/Pci (0x1,0x1)/Ata (0x0)/CDROM (0x0)
FS2: Alias(s):CD0c1::BLK4:
PciRoot (0x0)/Pci (0x1,0x1)/Ata (0x0)/CDROM (0x1)
BLK5: Alias(s):
PciRoot (0x0)/Pci (0xD,0x0)/Sata (0x0,0xFFFF,0x0)
BLK1: Alias(s):
PciRoot (0x0)/Pci (0x1,0x1)/Ata (0x0)/CDROM (0x0)
BLK3: Alias(s):
PciRoot (0x0)/Pci (0x1,0x1)/Ata (0x0)/CDROM (0x0)
Press ESC in 4 seconds to skip startup.nsh or any other key to continue.
```

```
Shell>
Shell> FS0:
FS0:\> ls
Directory of: FS0:\
00/00/0000 00:00 <DIR>          2,048 .disk
00/00/0000 00:00 <DIR>          2,048 boot
00/00/0000 00:00 <DIR>          2,048 casper
00/00/0000 00:00 <DIR>          2,048 dists
00/00/0000 00:00 <DIR>          2,048 EFI
00/00/0000 00:00 <DIR>          2,048 install
00/00/0000 00:00 <DIR>         34,816 isolinux
00/00/0000 00:00             42,622 md5sum.txt
00/00/0000 00:00 <DIR>          2,048 pics
00/00/0000 00:00 <DIR>          2,048 pool
00/00/0000 00:00 <DIR>          2,048 preseed
00/00/0000 00:00             231 README.diskdefines
00/00/0000 00:00              0 ubuntu
      3 File(s)      42,853 bytes
      10 Dir(s)
FS0:\> ls EFI\BOOT
Directory of: FS0:\EFI\BOOT\
00/00/0000 00:00             1,334,816 BOOTx64.EFI
00/00/0000 00:00             1,419,136 grubx64.efi
00/00/0000 00:00             1,269,496 mmx64.efi
      3 File(s)     4,023,448 bytes
      0 Dir(s)
FS0:\> _
```

BOOTx64.EFI
is the
primary ISO
boot file

UEFI Shell: Booting example (mmx64)

```
UEFI Interactive Shell v2.2
EDK II
UEFI v2.70 (EDK II, 0x00010000)
Mapping table
FS0: Alias(s) :F0c::BLK0:
    PciRoot(0x0)/Pci(0x1,0x1)/Ata(0x0)
FS1: Alias(s) :CD0c0::BLK2:
    PciRoot(0x0)/Pci(0x1,0x1)/Ata(0x0)/CDROM(0x0)
FS2: Alias(s) :CD0c1::BLK4:
    PciRoot(0x0)/Pci(0x1,0x1)/Ata(0x0)/CDROM(0x1)
BLK5: Alias(s) :
    PciRoot(0x0)/Pci(0xD,0x0)/Sata(0x0,0xFFFF,0x0)
BLK1: Alias(s) :
    PciRoot(0x0)/Pci(0x1,0x1)/Ata(0x0)/CDROM(0x0)
BLK3: Alias(s) :
    PciRoot(0x0)/Pci(0x1,0x1)/Ata(0x0)/CDROM(0x0)
Press ESC in 5 seconds to skip startup.nsh or any other key to continue.
Shell> FS0:
FS0:\> EFI\BOOT\mmx64.efi_
```

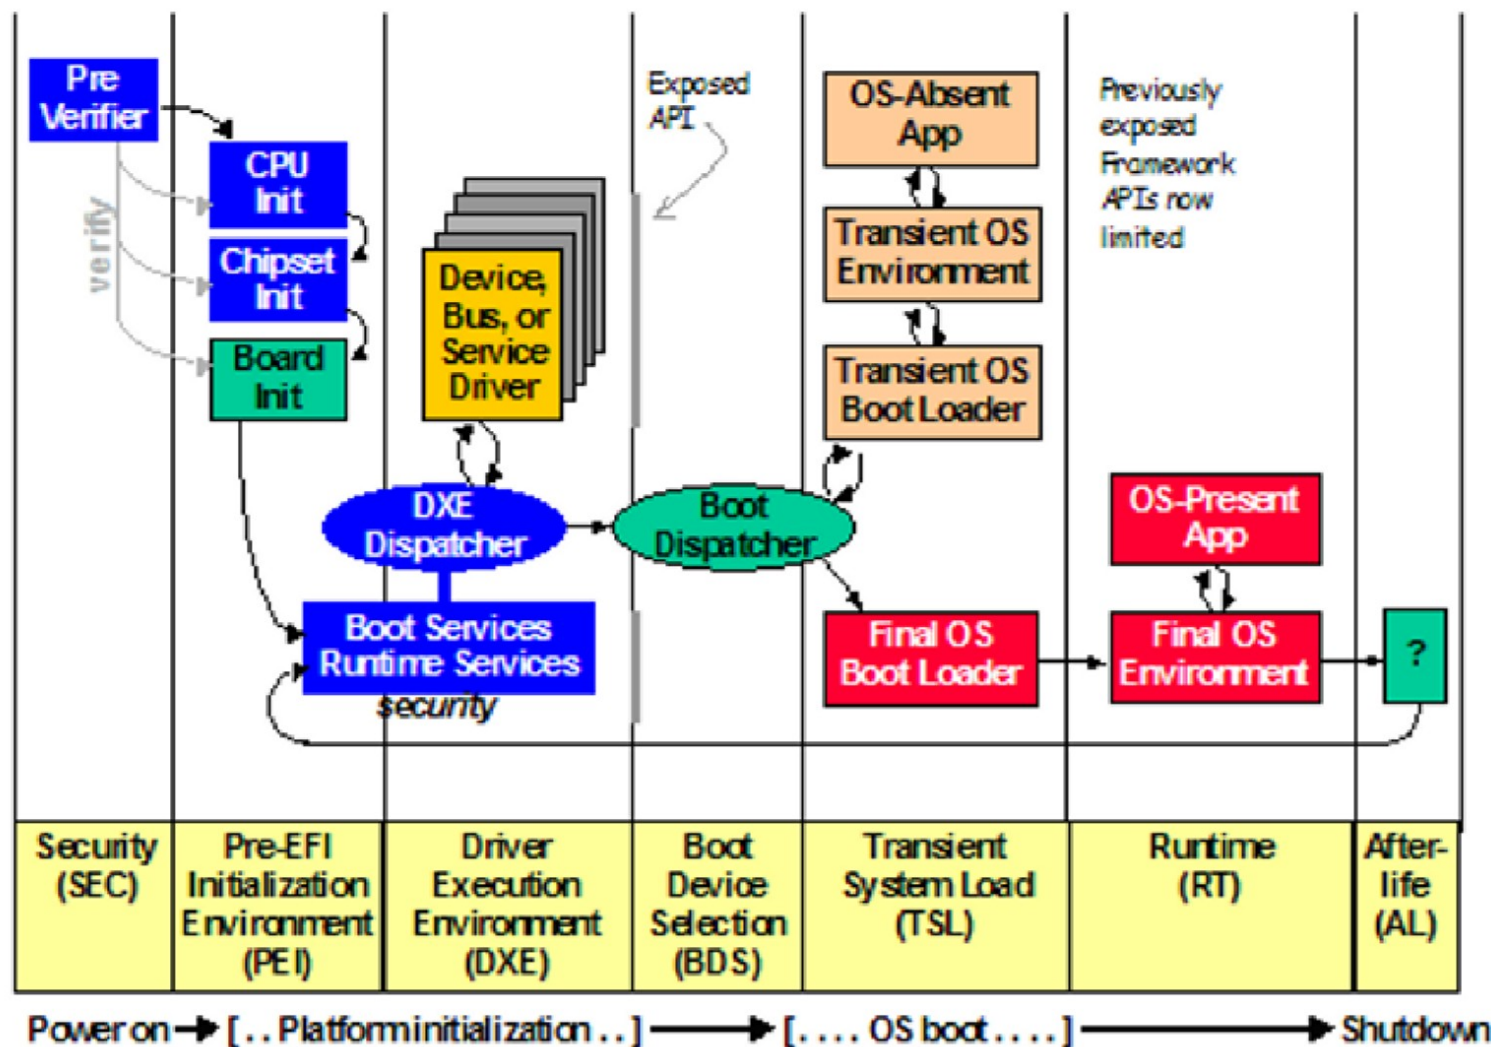
UEFI Shell: Booting example (mmx64)

Shim UEFI key management

Press any key to perform MOK management

Booting in 4 seconds

UEFI Boot Process



* The picture is taken from EDK2 specification: https://edk2-docs.gitbook.io/edk-ii-build-specification/2_design_discussion/23_boot_sequence

UEFI Services

- Native code
 - No legacy 16-bit BIOS code
 - 32-bit and 64-bit depending on an OS [For x86-64, 32-bit support is optional]
- UEFI boot time services
 - Destroyed when an OS is booted up
 - Cannot access corresponding methods afterwards
- UEFI run time service
 - Can be used within the booted OS
 - Some similarity to what BIOS used to be before its 16-bit mode became obsolete

Legacy Master Boot Record (MBR)

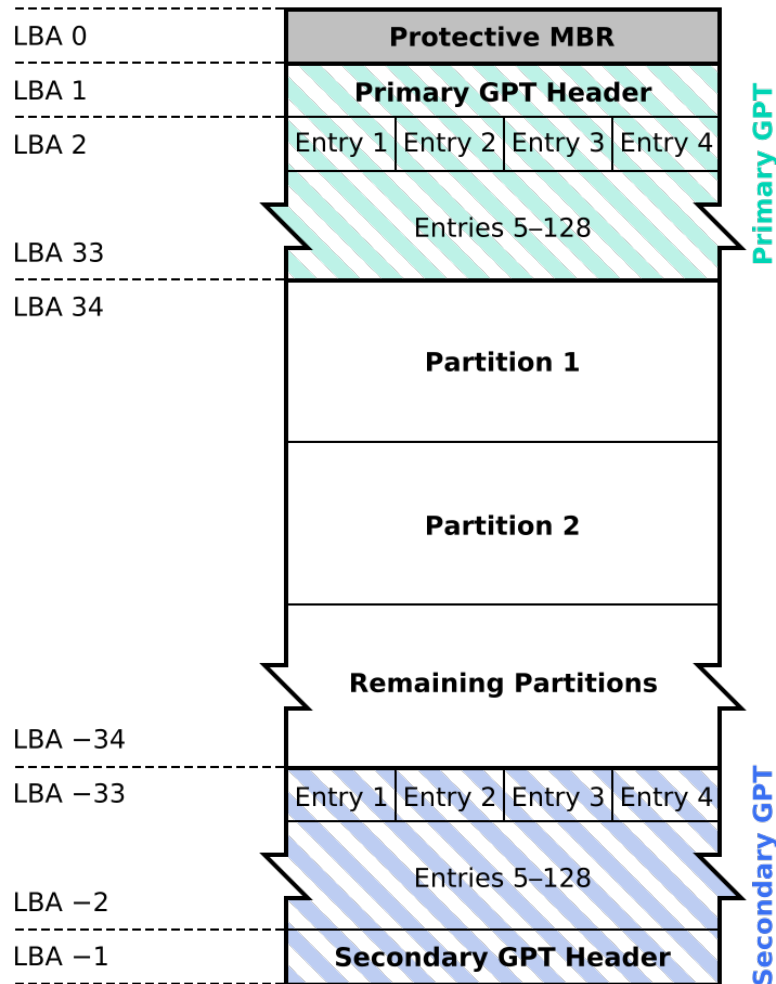
- Resides in disk sector #0 and is dual-purposed (sectors are typically 512 bytes for HDD)
 - The first-stage boot loader code that BIOS loads to memory address 0x7C00 – 446 bytes
 - Too small but it is feasible to create multiple stages, etc. Boot sector writing now is a lost art
 - A table of 4 partitions (max), more can be added by extra chaining – 4x16 bytes
 - Due to 32-bit logical block addresses (LBA), the disk space is limited to $(2^{32} * 512) = 2 \text{ TB}$
- Not used for ISO 9660 (CD/DVD), their sectors are 2048 bytes
 - Another “El-Torito” protocol is used instead

GUID Partition Table (GPT)

- MBR partition table is now obsolete due to the 2 TB limit
 - A dummy, “protective”, MBR sector is created which has no real partition table or boot code
- The protective MBR is followed by GPT (sector #1), a new partition table
 - A dedicated “EFI system partition” stores bootable files
 - Other partitions can be created as necessary
 - All partitions have unique 128-bit GUID identifiers
 - Can be really useful for an OS kernel to map these partitions (especially the bootable partition) to its internal names, e.g., the /dev namespace in Linux

GUID Partition Table (GPT)

GUID Partition Table Scheme



* The picture is taken from https://en.wikipedia.org/wiki/GUID_Partition_Table

EFI System Partition

- Typically initialized to FAT32
 - Widely supported across different OSs
- Can have multiple bootable files
 - Multiple OSs can be installed, one .efi program (“boot selector”) can load another .efi program, etc
 - grubx64.efi for the GRUB boot loader (Linux)
 - bootmgfw.efi for the Windows boot manager

Executable Format

- .EFI files are essentially .DLL files (with some minor differences)
 - Microsoft ABI and calling conventions
- .DLL is a *relocatable* binary (similar to ELF relocatable object files)
 - Can be loaded anywhere in the address space
 - UEFI creates a “default page table” already
 - No symbols are imported or exported, one function is the entry point, no symbols

Executable Format

```
ruslan@ruslan-ThinkPad-T470p: ~/examples/shim/build-x64
ruslan@ruslan-ThinkPad-T470p:~/examples/shim/build-x64$ objdump -x shimx64.efi

shimx64.efi:      file format pei-x86-64
shimx64.efi
architecture: i386:x86-64, flags 0x0000012f:
HAS_RELOC, EXEC_P, HAS_LINENO, HAS_DEBUG, HAS_LOCALS, D_PAGED
start address 0x0000000018002774a

Characteristics 0x2022
  executable
  large address aware
  DLL

Time/Date      Mon Jan 25 15:32:58 2021
Magic          020b (PE32+)
MajorLinkerVersion
0
MinorLinkerVersion
0
SizeOfCode     000000000004a000
SizeOfInitializedData 000000000002da00
SizeOfUninitializedData 0000000000000000
AddressOfEntryPoint 000000000002774a
BaseOfCode     0000000000001000
ImageBase      0000000018000000
SectionAlignment 00001000
FileAlignment  00000200
MajorOSVersion 0
MinorOSVersion 0
MajorImageVersion 0
MinorImageVersion 0
MajorSubsystemVersion 0
MinorSubsystemVersion 0
Win32Version   00000000
SizeOfImage    0007a000
SizeOfHeaders  00000400
Checksum       00000000
Subsystem      0000000a (EFI application)
DllCharacteristics 00000160
SizeOfStackReserve 0000000000000000
SizeOfStackCommit 0000000000000000
SizeOfHeapReserve 0000000000000000
SizeOfHeapCommit 0000000000000000
LoaderFlags    00000000
NumberOfRvaAndSizes 00000010

The Data Directory
Entry 0 0000000000000000 00000000 Export Directory [.edata (or where ever we found it)]
Entry 1 0000000000000000 00000000 Import Directory [parts of .idata]
Entry 2 0000000000000000 00000000 Resource Directory [.rsrc]
Entry 3 0000000000000000 00000000 Exception Directory [.pdata]
Entry 4 0000000000000000 00000000 Security Directory
Entry 5 0000000000077000 00002858 Base Relocation Directory [.reloc]
Entry 6 0000000000000000 00000000 Debug Directory
Entry 7 0000000000000000 00000000 Description Directory
Entry 8 0000000000000000 00000000 Special Directory
```

```
ruslan@ruslan-ThinkPad-T470p: ~/examples/shim/build-x64

reloc 125 offset b50 [75b50] DIR64
reloc 126 offset b58 [75b58] DIR64
reloc 127 offset b78 [75b78] DIR64
reloc 128 offset b80 [75b80] DIR64
reloc 129 offset ba0 [75ba0] DIR64
reloc 130 offset ba8 [75ba8] DIR64
reloc 131 offset bd0 [75bd0] DIR64
reloc 132 offset bd8 [75bd8] DIR64
reloc 133 offset be0 [75be0] DIR64
reloc 134 offset be8 [75be8] DIR64
reloc 135 offset bf0 [75bf0] DIR64
reloc 136 offset bf8 [75bf8] DIR64
reloc 137 offset c10 [75c10] DIR64
reloc 138 offset c18 [75c18] DIR64
reloc 139 offset c20 [75c20] DIR64
reloc 140 offset c40 [75c40] DIR64
reloc 141 offset c48 [75c48] DIR64
reloc 142 offset c50 [75c50] DIR64
reloc 143 offset c70 [75c70] DIR64
reloc 144 offset c78 [75c78] DIR64
reloc 145 offset c80 [75c80] DIR64
reloc 146 offset ca0 [75ca0] DIR64
reloc 147 offset ca8 [75ca8] DIR64
reloc 148 offset cb0 [75cb0] DIR64
reloc 149 offset cd0 [75cd0] DIR64
reloc 150 offset cd8 [75cd8] DIR64
reloc 151 offset ce0 [75ce0] DIR64
reloc 152 offset d00 [75d00] DIR64
reloc 153 offset d08 [75d08] DIR64
reloc 154 offset d10 [75d10] DIR64
reloc 155 offset d30 [75d30] DIR64
reloc 156 offset d38 [75d38] DIR64
reloc 157 offset d40 [75d40] DIR64
reloc 158 offset d60 [75d60] DIR64
reloc 159 offset d68 [75d68] DIR64
reloc 160 offset d70 [75d70] DIR64
reloc 161 offset d90 [75d90] DIR64
reloc 162 offset d98 [75d98] DIR64
reloc 163 offset da0 [75da0] DIR64

Sections:
Idx Name          Size      VMA               LMA               File off  Algn
 0 .text          00049e48 00000000180001000 00000000180001000 00000400 2**4
CONTENTS, ALLOC, LOAD, READONLY, CODE
 1 .rdata          00029e30 0000000018004b000 0000000018004b000 0004a400 2**4
CONTENTS, ALLOC, LOAD, READONLY, DATA
 2 .data           00001000 00000000180075000 00000000180075000 00074400 2**4
CONTENTS, ALLOC, LOAD, DATA
 3 .reloc           00002858 00000000180077000 00000000180077000 00075400 2**2
CONTENTS, ALLOC, LOAD, READONLY, DATA

SYMBOL TABLE:
no symbols

ruslan@ruslan-ThinkPad-T470p:~/examples/shim/build-x64$
```

Executable Format

```
ruslan@ruslan-ThinkPad-T470p: ~/examples/shim
}

extern EFI_STATUS EFI_API
efi_main(EFI_HANDLE passed_image_handle, EFI_SYSTEM_TABLE *passed_systab);

EFI_STATUS EFI_API
efi_main (EFI_HANDLE passed_image_handle, EFI_SYSTEM_TABLE *passed_systab)
{
    EFI_STATUS efi_status;
    EFI_HANDLE image_handle;

    systab = passed_systab;
    image_handle = global_image_handle = passed_image_handle;

    InitializeLib(image_handle, systab);

    verification_method = VERIFIED_BY_NOHING;

    vendor_authorized_size = cert_table.vendor_authorized_size;
}
```

2607,1 97%

```
ruslan@ruslan-ThinkPad-T470p: ~/examples/shim
endif

%.efi : $(BUILDDIR)/%.efi

src/%.o : $(BUILDDIR)/src/%.o

%.crt : $(BUILDDIR)/certdb/%.crt
%.cer : $(BUILDDIR)/certdb/%.cer

EFI_LDFLAGS = $(EFI_DEBUG_LDFLAGS) /dll /entry:efi_main /safeseh:no /nodefaultlib $(EFI_ARCH_LDFLAGS)

%.efi: %.efi.dll fwimage/fwimage
    ./fwimage/fwimage app $< $@
    @chmod 755 $@

ifneq ($(origin ENABLE_SBSIGN),undefined)
%.efi.signed: %.efi shim.key shim.crt
    @$(SBSIGN) \
        --key $(BUILDDIR)/certdb/shim.key \
        --cert $(BUILDDIR)/certdb/shim.crt \
search hit BOTTOM, continuing at TOP
```

89,0-1 77%

UEFI Secure Boot

- Malware can hijack the boot process
- One way to alleviate this concern is to use “digital signing” of .efi files with a private key
 - A trusted party (e.g., Microsoft, Red Hat, etc) can sign with their private key
 - The private key is unknown to the public
- The .efi file is then verified with a public key
 - The public key is installed in the UEFI database



EFI_SYSTEM_TABLE

- The topmost data structure which contains function pointers to
 - UEFI boot time services
 - UEFI run time services
 - Protocol services
 - Each protocol is identified with its own GUID

EFI_SYSTEM_TABLE

```
ruslan@ruslan-ThinkPad-T470p: ~/examples/shim/edk2
typedef struct {
    ///
    /// The table header for the EFI System Table.
    ///
    EFI_TABLE_HEADER          Hdr;
    ///
    /// A pointer to a null terminated string that identifies the vendor
    /// that produces the system firmware for the platform.
    ///
    CHAR16                    *FirmwareVendor;
    ///
    /// A firmware vendor specific value that identifies the revision
    /// of the system firmware for the platform.
    ///
    UINT32                     FirmwareRevision;
    ///
    /// The handle for the active console input device. This handle must support
    /// EFI_SIMPLE_TEXT_INPUT_PROTOCOL and EFI_SIMPLE_TEXT_INPUT_EX_PROTOCOL.
    ///
    EFI_HANDLE                 ConsoleInHandle;
    ///
    /// A pointer to the EFI_SIMPLE_TEXT_INPUT_PROTOCOL interface that is
    /// associated with ConsoleInHandle.
    ///
    EFI_SIMPLE_TEXT_INPUT_PROTOCOL *ConIn;
    ///
    /// The handle for the active console output device.
    ///
    EFI_HANDLE                 ConsoleOutHandle;
    ///
    /// A pointer to the EFI_SIMPLE_TEXT_OUTPUT_PROTOCOL interface
    /// that is associated with ConsoleOutHandle.
    ///
    EFI_SIMPLE_TEXT_OUTPUT_PROTOCOL *ConOut;
    ///
    /// The handle for the active standard error console device.
    /// This handle must support the EFI_SIMPLE_TEXT_OUTPUT_PROTOCOL.
    ///
    EFI_HANDLE                 StandardErrorHandle;
    ///
    /// A pointer to the EFI_SIMPLE_TEXT_OUTPUT_PROTOCOL interface
    /// that is associated with StandardErrorHandle.
    ///
    EFI_SIMPLE_TEXT_OUTPUT_PROTOCOL *StdErr;
    ///
    /// A pointer to the EFI Runtime Services Table.
    ///
    EFI_RUNTIME_SERVICES        *RuntimeServices;
    ///
    /// A pointer to the EFI Boot Services Table.
    ///
    EFI_BOOT_SERVICES           *BootServices;
    ///
    /// The number of system configuration tables in the buffer ConfigurationTable.
    ///
    UINTN                       NumberOfTableEntries;
    ///
    /// A pointer to the system configuration tables.
    /// The number of entries in the table is NumberOfTableEntries.
    ///
    EFI_CONFIGURATION_TABLE      *ConfigurationTable;
} EFI_SYSTEM_TABLE;
```

2010,5 91%

EFI_BOOT_SERVICES

```
ruslan@ruslan-ThinkPad-T470p: ~/examples/shim/edk2
//
typedef struct {
    ///
    /// The table header for the EFI Boot Services Table.
    ///
    EFI_TABLE_HEADER          Hdr;

    ///
    /// Task Priority Services
    ///
    EFI_RAISE_TPL              RaiseTPL;
    EFI_RESTORE_TPL            RestoreTPL;

    ///
    /// Memory Services
    ///
    EFI_ALLOCATE_PAGES          AllocatePages;
    EFI_FREE_PAGES              FreePages;
    EFI_GET_MEMORY_MAP           GetMemoryMap;
    EFI_ALLOCATE_POOL            AllocatePool;
    EFI_FREE_POOL                FreePool;

    ///
    /// Event & Timer Services
    ///
    EFI_CREATE_EVENT             CreateEvent;
    EFI_SET_TIMER                SetTimer;
    EFI_WAIT_FOR_EVENT           WaitForEvent;
    EFI_SIGNAL_EVENT             SignalEvent;
    EFI_CLOSE_EVENT              CloseEvent;
    EFI_CHECK_EVENT              CheckEvent;

    ///
    /// Protocol Handler Services
    ///
    EFI_INSTALL_PROTOCOL_INTERFACE InstallProtocolInterface;
    EFI_REINSTALL_PROTOCOL_INTERFACE ReinstallProtocolInterface;
    EFI_UNINSTALL_PROTOCOL_INTERFACE UninstallProtocolInterface;
    EFI_HANDLE_PROTOCOL           HandleProtocol;
    VOID                          *Reserved;
    EFI_REGISTER_PROTOCOL_NOTIFY RegisterProtocolNotify;
    EFI_LOCATE_HANDLE             LocateHandle;
    EFI_LOCATE_DEVICE_PATH         LocateDevicePath;
    EFI_INSTALL_CONFIGURATION_TABLE InstallConfigurationTable;

    ///
    /// Image Services
    ///
    EFI_IMAGE_LOAD               LoadImage;
    EFI_IMAGE_START               StartImage;
    EFI_EXIT                      Exit;
    EFI_IMAGE_UNLOAD              UnloadImage;
    EFI_EXIT_BOOT_SERVICES        ExitBootServices;

    ///
    /// Miscellaneous Services
    ///
    EFI_GET_NEXT_MONOTONIC_COUNT GetNextMonotonicCount;
    EFI_STALL                     Stall;
} EFI_BOOT_SERVICES;
```

1885,3 85%

```
ruslan@ruslan-ThinkPad-T470p: ~/examples/shim/edk2
EFI_UNINSTALL_PROTOCOL_INTERFACE UninstallProtocolInterface;
EFI_HANDLE_PROTOCOL              HandleProtocol;
VOID                             *Reserved;
EFI_REGISTER_PROTOCOL_NOTIFY      RegisterProtocolNotify;
EFI_LOCATE_HANDLE                 LocateHandle;
EFI_LOCATE_DEVICE_PATH            LocateDevicePath;
EFI_INSTALL_CONFIGURATION_TABLE    InstallConfigurationTable;

///
/// Image Services
///
EFI_IMAGE_LOAD                    LoadImage;
EFI_IMAGE_START                  StartImage;
EFI_EXIT                          Exit;
EFI_IMAGE_UNLOAD                  UnloadImage;
EFI_EXIT_BOOT_SERVICES            ExitBootServices;

///
/// Miscellaneous Services
///
EFI_GET_NEXT_MONOTONIC_COUNT       GetNextMonotonicCount;
EFI_STALL                         Stall;
EFI_SET_WATCHDOG_TIMER             SetWatchdogTimer;

///
/// DriverSupport Services
///
EFI_CONNECT_CONTROLLER             ConnectController;
EFI_DISCONNECT_CONTROLLER          DisconnectController;

///
/// Open and Close Protocol Services
///
EFI_OPEN_PROTOCOL                 OpenProtocol;
EFI_CLOSE_PROTOCOL                CloseProtocol;
EFI_OPEN_PROTOCOL_INFORMATION      OpenProtocolInformation;

///
/// Library Services
///
EFI_PROTOCOLS_PER_HANDLE           ProtocolsPerHandle;
EFI_LOCATE_HANDLE_BUFFER           LocateHandleBuffer;
EFI_LOCATE_PROTOCOL                LocateProtocol;
EFI_INSTALL_MULTIPLE_PROTOCOL_INTERFACES InstallMultipleProtocolInterfaces;
EFI_UNINSTALL_MULTIPLE_PROTOCOL_INTERFACES UninstallMultipleProtocolInterfaces;

///
/// 32-bit CRC Services
///
EFI_CALCULATE_CRC32                CalculateCrc32;

///
/// Miscellaneous Services
///
EFI_COPY_MEM                       CopyMem;
EFI_SET_MEM                        SetMem;
EFI_CREATE_EVENT_EX                CreateEventEx;
} EFI_BOOT_SERVICES;
```

1882,3 87%



UEFI Example Program

```
#include <Uefi/UefiSpec.h>
```

```
EFI_STATUS EFIAPI
```

```
efi_main(EFI_HANDLE ImageHandle, EFI_SYSTEM_TABLE *SystemTable) {  
    SystemTable->ConOut->OutputString(SystemTable->ConOut, L"Hello World!\n");  
    return EFI_SUCCESS;  
}
```


UEFI Example Program

```
ruslan@ruslan-ThinkPad-T470p: ~/examples/shim/edk2
/// The number of modes supported by QueryMode () and SetMode ().
///
INT32  MaxMode;

///
/// current settings
///

///
/// The text mode of the output device(s).
///
INT32  Mode;
///
/// The current character output attribute.
///
INT32  Attribute;
///
/// The cursor's column.
///
INT32  CursorColumn;
///
/// The cursor's row.
///
INT32  CursorRow;
///
/// The cursor is currently visible or not.
///
BOOLEAN CursorVisible;
} EFI_SIMPLE_TEXT_OUTPUT_MODE;

///
/// The SIMPLE_TEXT_OUTPUT protocol is used to control text-based output devices.
/// It is the minimum required protocol for any handle supplied as the ConsoleOut
/// or StandardError device. In addition, the minimum supported text mode of such
/// devices is at least 80 x 25 characters.
///
struct _EFI_SIMPLE_TEXT_OUTPUT_PROTOCOL {
    EFI_TEXT_RESET      Reset;

    EFI_TEXT_STRING      OutputString;
    EFI_TEXT_TEST_STRING TestString;

    EFI_TEXT_QUERY_MODE  QueryMode;
    EFI_TEXT_SET_MODE    SetMode;
    EFI_TEXT_SET_ATTRIBUTE SetAttribute;

    EFI_TEXT_CLEAR_SCREEN ClearScreen;
    EFI_TEXT_SET_CURSOR_POSITION SetCursorPosition;
    EFI_TEXT_ENABLE_CURSOR EnableCursor;

    ///
    /// Pointer to SIMPLE_TEXT_OUTPUT_MODE data.
    ///
    EFI_SIMPLE_TEXT_OUTPUT_MODE *Mode;
};

extern EFI_GUID gEfiSimpleTextOutProtocolGuid;

#endif
```

409,1 Bot

```
ruslan@ruslan-ThinkPad-T470p: ~/examples/shim/edk2
/**
typedef
EFI_STATUS
(EFIAPI *EFI_TEXT_RESET)(
    IN EFI_SIMPLE_TEXT_OUTPUT_PROTOCOL *This,
    IN BOOLEAN ExtendedVerification
);

/**
Write a string to the output device.

@param This The protocol instance pointer.
@param String The NULL-terminated string to be displayed on the output
device(s). All output devices must also support the Unicode
drawing character codes defined in this file.

@retval EFI_SUCCESS The string was output to the device.
@retval EFI_DEVICE_ERROR The device reported an error while attempting to o
utput
the text.
@retval EFI_UNSUPPORTED The output device's mode is not currently in a
defined text mode.
@retval EFI_WARN_UNKNOWN_GLYPH This warning code indicates that some of the
characters in the string could not be
rendered and were skipped.

/**
typedef
EFI_STATUS
(EFIAPI *EFI_TEXT_STRING)(
    IN EFI_SIMPLE_TEXT_OUTPUT_PROTOCOL *This,
    IN CHAR16 *String
);

/**
Verifies that all characters in a string can be output to the
target device.

@param This The protocol instance pointer.
@param String The NULL-terminated string to be examined for the output
device(s).

@retval EFI_SUCCESS The device(s) are capable of rendering the output string.
@retval EFI_UNSUPPORTED Some of the characters in the string cannot be
rendered by one or more of the output devices mapped
by the EFI handle.

/**
typedef
EFI_STATUS
(EFIAPI *EFI_TEXT_TEST_STRING)(
    IN EFI_SIMPLE_TEXT_OUTPUT_PROTOCOL *This,
    IN CHAR16 *String
);

/**
Returns information for an available text mode that the output device(s)
supports.

search hit BOTTOM, continuing at TOP
```

171,0-1 46%

Toolkits

- We are going to use clang and lld-link
 - `sudo apt-get install clang lld`
 - clang/lld 7.0.0+ should be OK
- Why clang?
 - UEFI requires to compile PE DLL executables (but they do not use anything from Windows)
 - Clang, unlike gcc, provides both ELF and PE support by default
 - Likewise, lld can link PE in addition to ELF

Toolkits

- In addition, you need to download and compile a tool which converts DLL to EFI
 - Use my version which I ported from Windows to Linux
 - git clone
<https://github.com/rusnikola/fwimage>
 - Run 'make'
 - It will generate the 'fwimage' program

Toolkits

- LLD is a tricky program which behaves differently depending on the program file name
 - “lld.link” for ELF and “lld-link” for PE
 - Ubuntu (and possibly other distributions) does not create lld-link to lld for some reason
 - Do the following (after you install lld)

```
$ cd /usr/bin
```

```
$ sudo ln -s lld lld-link
```

[will create lld-link if it does not exist already]

Toolkits

- Download Intel Tianocore/Edk2 headers
 - Create “uefi_projects” directory and make it current
 - wget
<https://github.com/tianocore/edk2/archive/refs/tags/edk2-stable202208.tar.gz>
 - tar xvpf ./edk2-stable202208.tar.gz **CORRECTION:
command to unpack**
 - mv edk2-edk2-stable202208/MdePkg/Include ./
 - rm -r edk2-edk2-stable202208
- “Include” in your current directory will have all necessary Edk2 headers

UEFI Example Program

Create hello.c in “uefi_projects”:

```
#include <Uefi.h>
```

```
EFI_STATUS EFIAPI
```

```
efi_main(EFI_HANDLE ImageHandle, EFI_SYSTEM_TABLE *SystemTable) {  
    SystemTable->ConOut->OutputString(SystemTable->ConOut, L"Hello World!\r\n");  
    SystemTable->BootServices->Stall(5 * 1000000); // 5 seconds  
    return EFI_SUCCESS;  
}
```

What is \r\n?

It is a Windows/DOS style for new lines (also adopted by UEFI):

\r (carriage return, CR)

\n (line feed, LF)

Linux uses \n, macOS uses \r, Windows/DOS uses both

UEFI Example Program

Compile:

```
clang -m64 -O2 -fshort-wchar -I ./Include -I ./Include/X64 -mcmmodel=small -mno-red-zone  
-mno-stack-arg-probe -target x86_64-pc-mingw32 -c hello.c
```

Link:

```
lld-link /dll /nodefaultlib /safeseh:no /machine:AMD64 /entry:efi_main hello.o /out:hello.dll
```

Convert DLL to EFI:

```
./fwimage/fwimage app hello.dll hello.efi
```

Create a bootable FAT image (10 MB):

```
mkdir ./uefi_fat_mnt  
dd if=/dev/zero of=boot.img bs=1M count=10  
mkfs.vfat ./boot.img  
sudo mount -o loop ./boot.img ./uefi_fat_mnt  
sudo mkdir -p ./uefi_fat_mnt/EFI/BOOT  
sudo cp hello.efi uefi_fat_mnt/EFI/BOOT/BOOTx64.EFI  
sudo umount ./uefi_fat_mnt
```

CORRECTION: there are some problems with ISO9660 when using recent versions of VirtualBox, so we will use FAT!

UEFI Example Program

Instructions for VirtualBox:

Download boot.vmdk (wrapper for boot.img) from Canvas (Modules->VirtualBox Files) and place it in the same directory as boot.img!

Note: If you have multiple copies of boot.vmdk, please change ddb.uuid.image=... in each 'boot.vmdk' copy:

Install UUID: `sudo apt-get install uuid`

Run 'uuid' and copy a unique UUID string to boot.vmdk

Instructions for qemu:

Download bios.zip (UEFI image for qemu) from Canvas (Modules->Qemu Files) and unpack it in the same directory as boot.img:

`unzip ./bios.zip`


What are those parameters?

- -fshort-wchar to use UTF-16 (2-byte) wide chars, by default Linux assumes UTF-32 (4-byte) wide chars
 - Note that we have L"Hello world!" rather than "Hello world!" since UEFI uses UTF-16, not UTF-8!
- -target x86_64-pc-mingw32 to use Windows executable format rather than ELF
- -I ./Include -I ./Include/X64 to use headers from your "Include" and "Include/X64" (architecture-specific) directories
- Other parameters to disable default system libraries, UEFI-incompatible options, etc

Why does UEFI use UTF-16?

- Probably due to Windows executables and because EFI was developed in the late 90s
- Legacy encodings only required 1-byte characters
- Unicode initially used 2-byte characters
 - Windows adopted Unicode (UTF-16)
 - Unicode later expanded beyond 64K characters, so we really need UTF-32 now. UTF-16 may now require “surrogate pairs.”
 - Why not 1-byte pairs then? This is what UTF-8 does. Linux adopted UTF-8 for simplicity.

VirtualBox Setup





Create Virtual Machine

Name and operating system

Please choose a descriptive name and destination folder for the new virtual machine and select the type of operating system you intend to install on it. The name you choose will be used throughout VirtualBox to identify this machine.

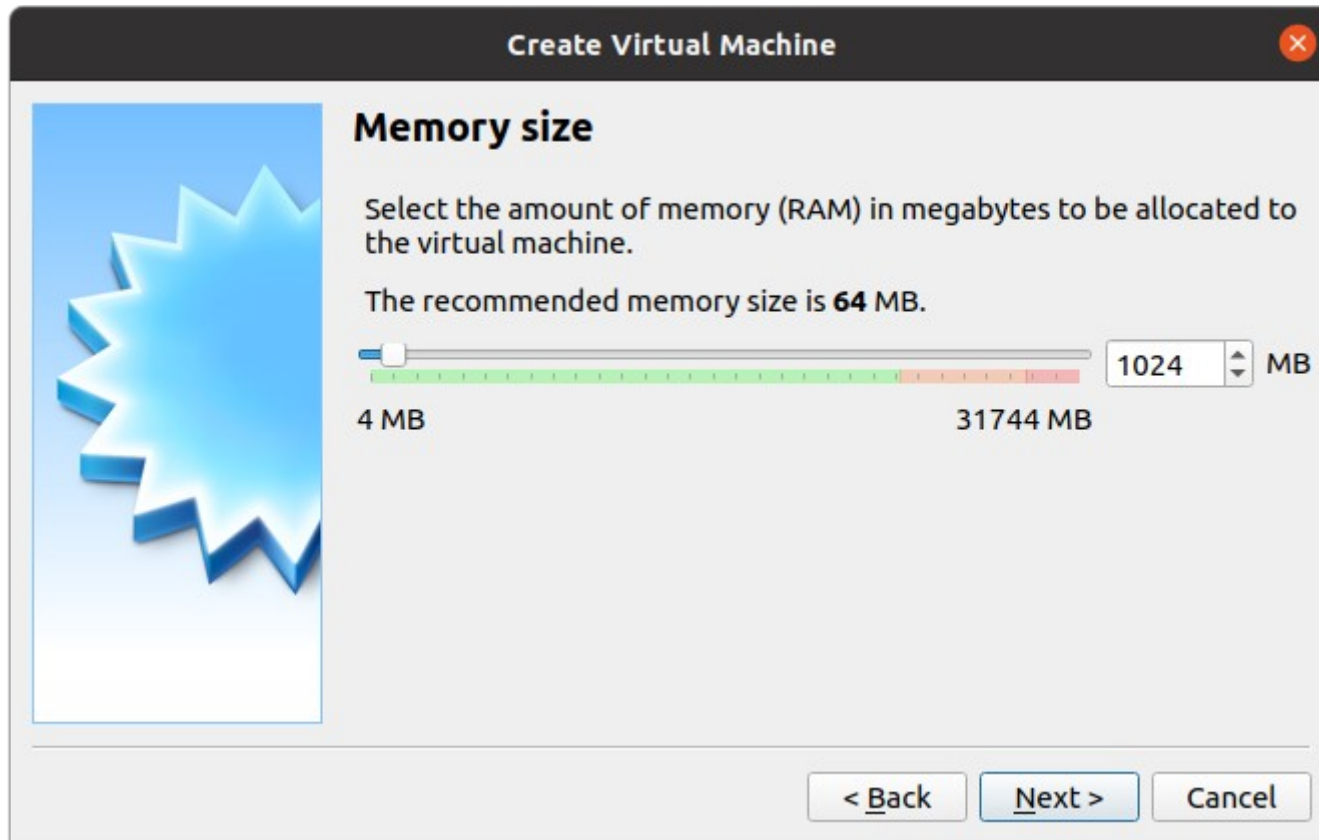
Name:

Machine Folder: 

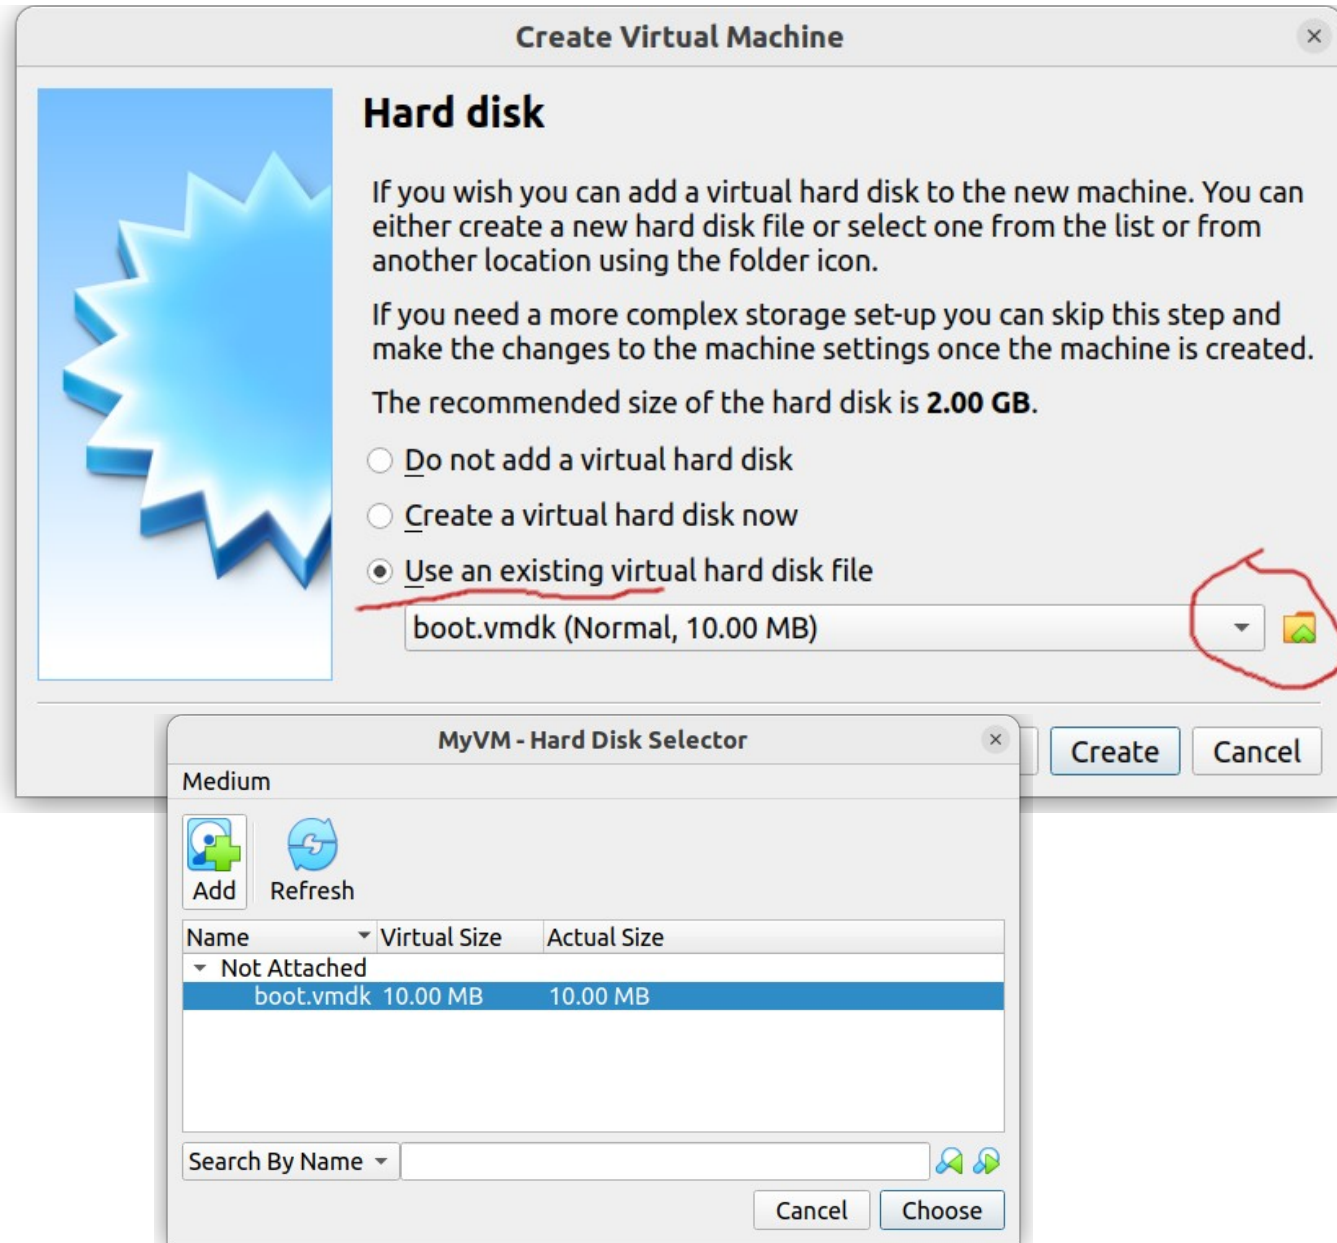
Type: 

Version:

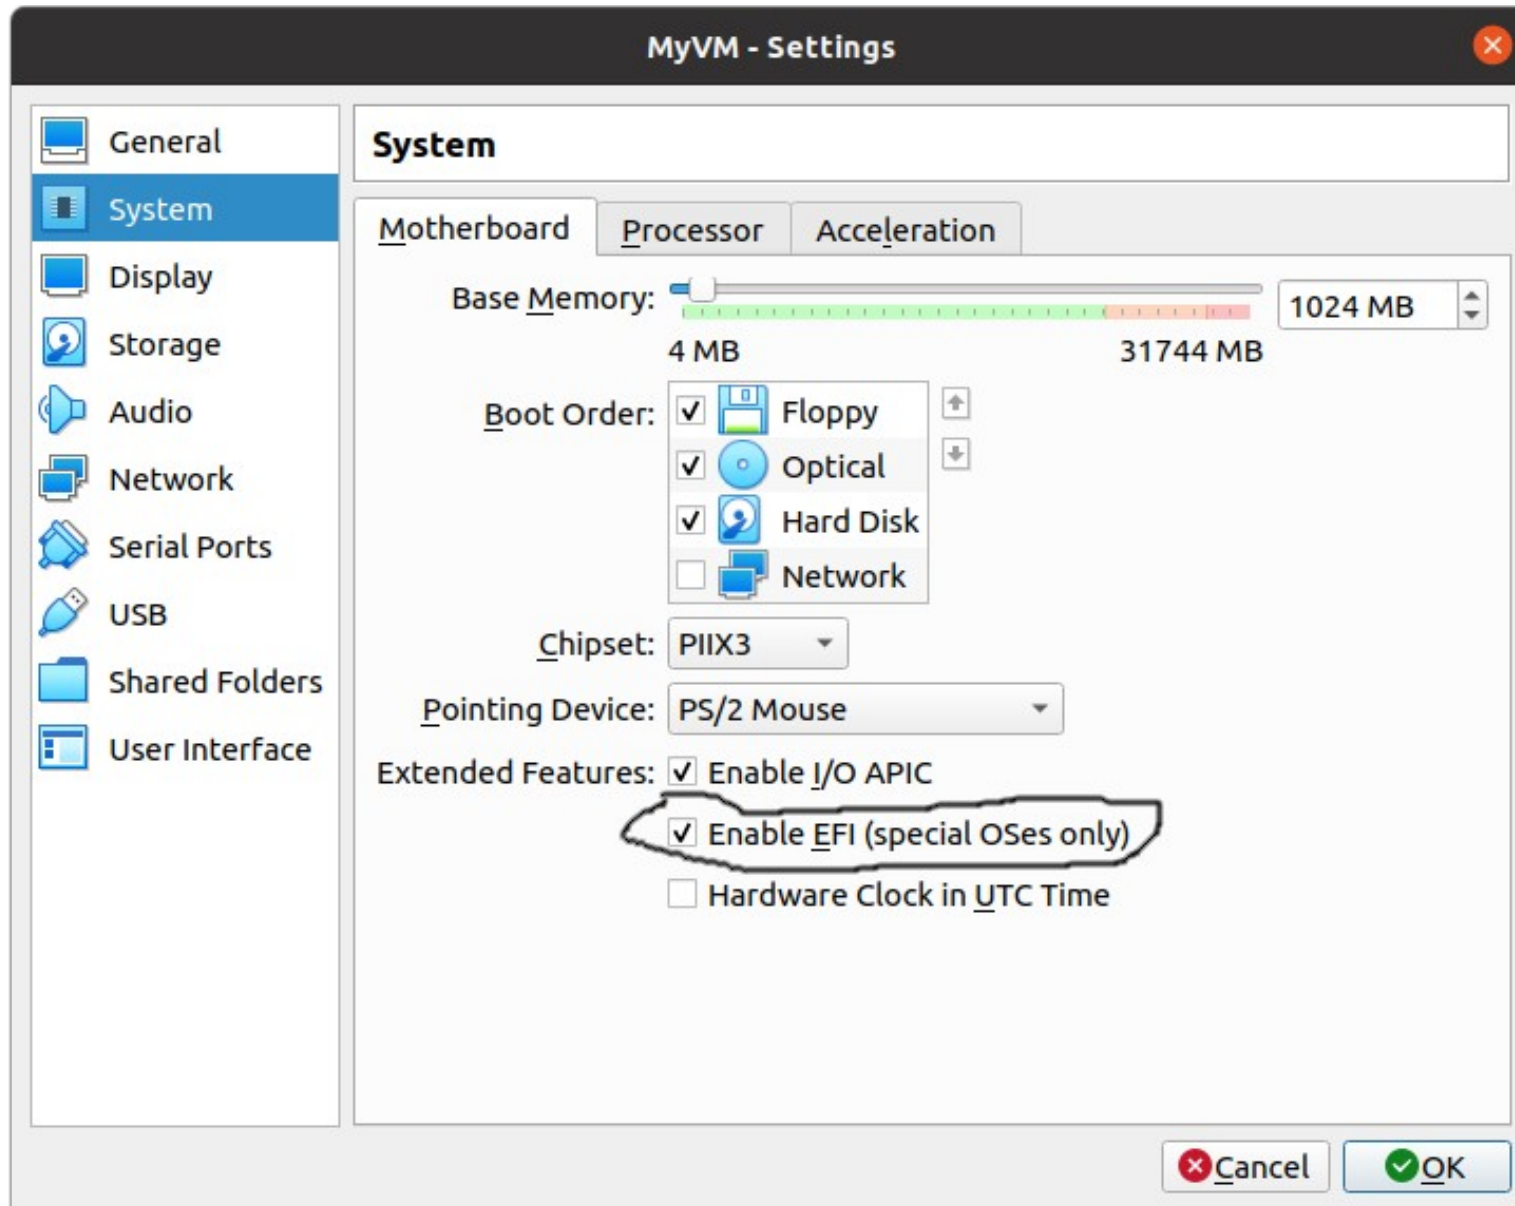
VirtualBox Setup



VirtualBox Setup



VirtualBox Setup



VirtualBox Setup: Execute

```
BdsDxe: failed to load Boot0001 "UEFI VBox CD-ROM VB2-01700376 " from PciRoot(0x0)/Pci(0x1,0x1)/Ata(Secondary,Master,0x0) : Not Found
BdsDxe: loading Boot0002 "UEFI VBox HARDDISK VB10a4524e-91b31f9a " from PciRoot(0x0)/Pci(0x1,0x1)/Ata(Primary,Master,0x0)
BdsDxe: starting Boot0002 "UEFI VBox HARDDISK VB10a4524e-91b31f9a " from PciRoot(0x0)/Pci(0x1,0x1)/Ata(Primary,Master,0x0)
Hello World!
```

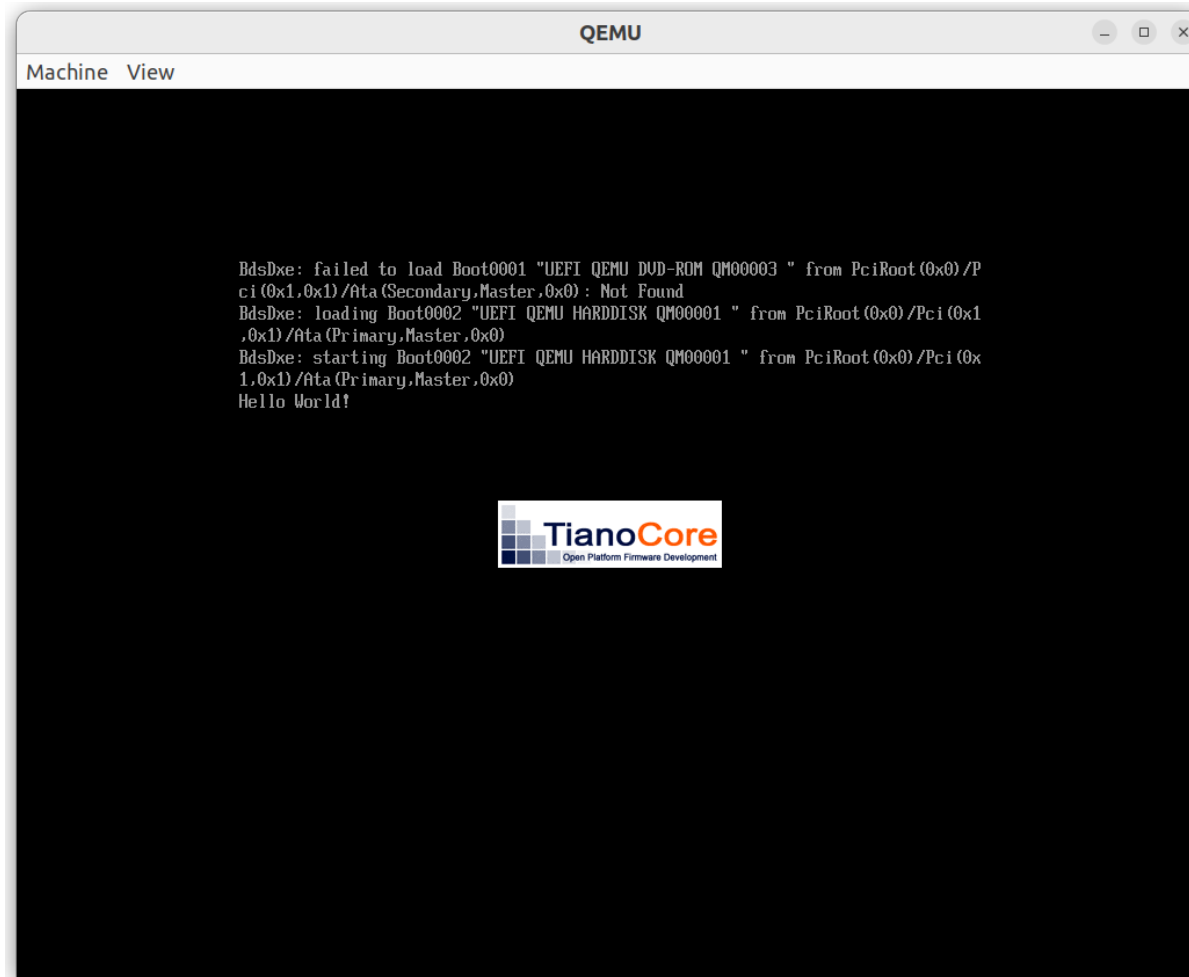


VirtualBox

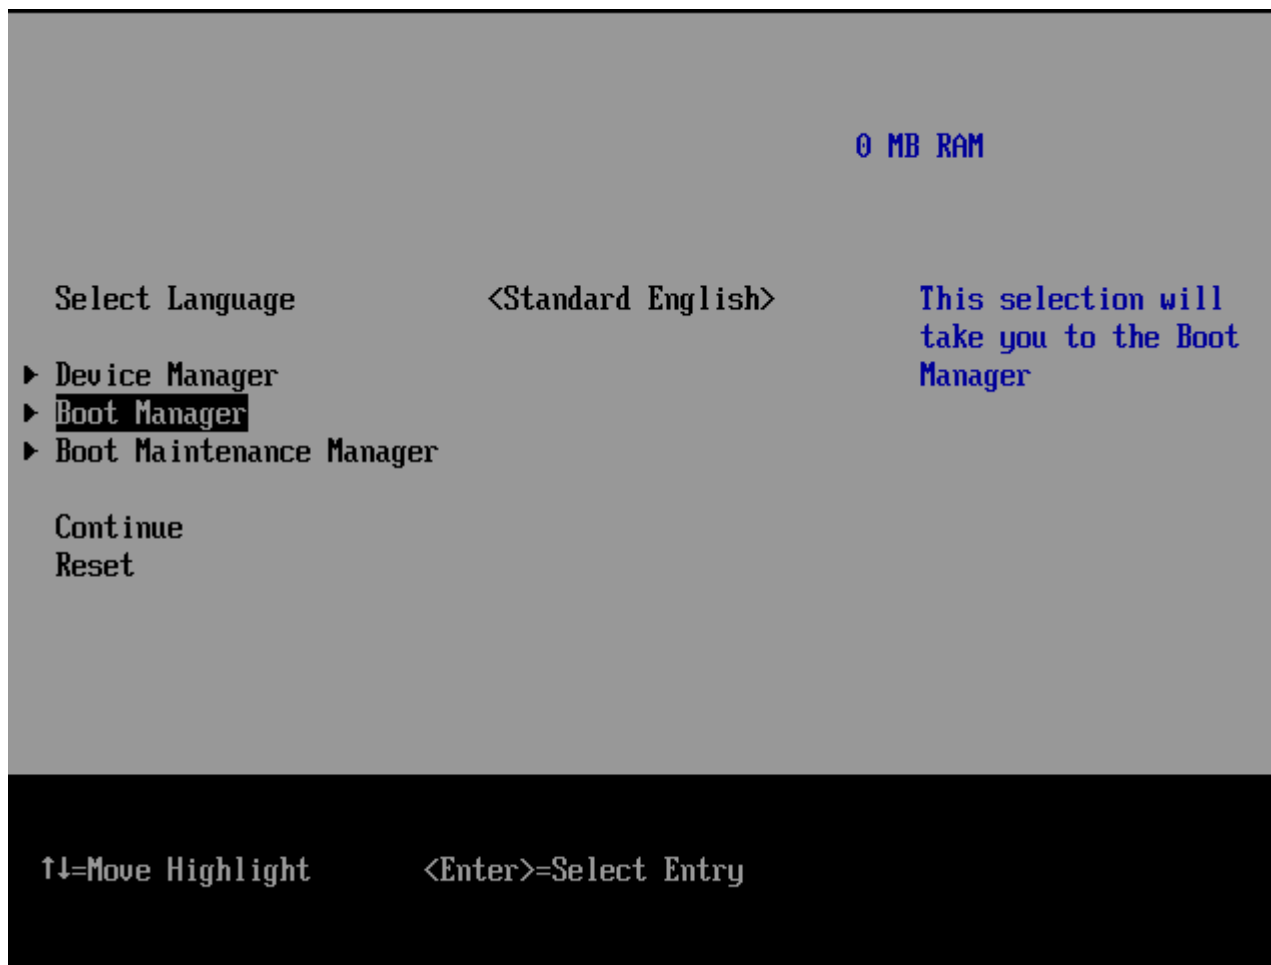
qemu: Execute

Run the following command:

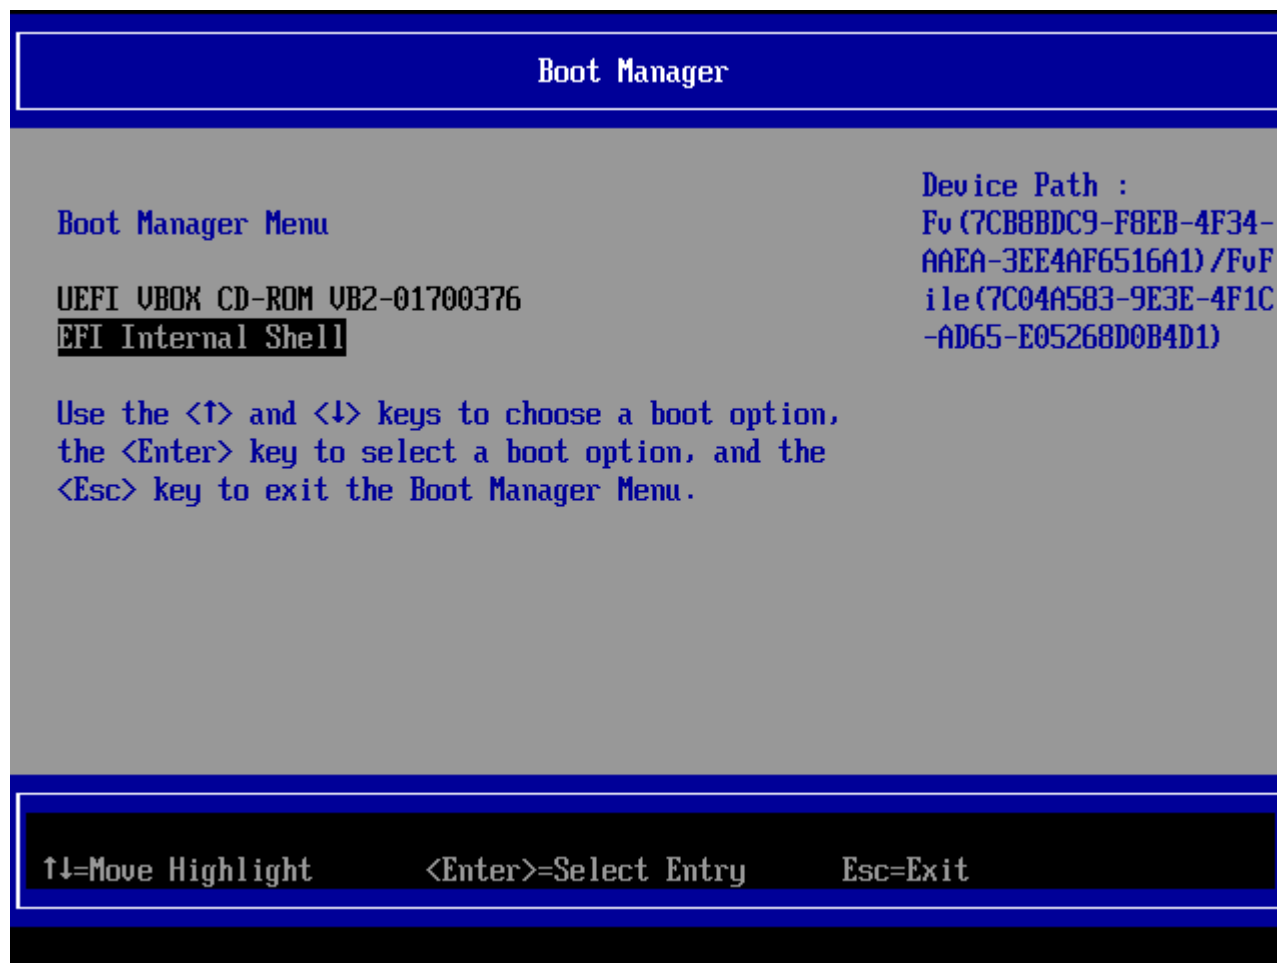
```
qemu-system-x86_64 -bios /path/to/bios.bin -m 1024 -drive  
format=raw,file=/path/to/boot.img
```



UEFI Example Program: Complete



UEFI Example Program: Shell



UEFI Example Program: Shell

```
UEFI Interactive Shell v2.2
EDK II
UEFI v2.70 (EDK II, 0x00010000)
Mapping table
  FS0: Alias(s):F0c::BLK0:
    PciRoot(0x0)/Pci(0x1,0x1)/Ata(0x0)
Press ESC in 5 seconds to skip startup.nsh or any other key to continue.
Shell>
Shell>
Shell>
Shell> FS0:
FS0:\> ls
Directory of: FS0:\
00/00/0000  00:00 <DIR>          2,048  EFI
          0 File(s)            0 bytes
          1 Dir(s)
FS0:\> ls EFI\BOOT
Directory of: FS0:\EFI\BOOT\
00/00/0000  00:00          2,560  BOOTX64.EFI
          1 File(s)        2,560 bytes
          0 Dir(s)
FS0:\> EFI\BOOT\BOOTX64.EFI
Hello World!
FS0:\> _
```

UEFI Documentation

- Go to <https://www.uefi.org/specifications/>
 - Download UEFI Specification Version 2.9 (released March 2021)
- There also used to be convenient Phoenix (BIOS/UEFI vendor) documentation
 - Still accessible through https://web.archive.org/web/20181012151104/http://wiki.phoenix.com/wiki/index.php/Category:UEFI_2.0

EFI_SYSTEM_TABLE and EFI_BOOT_SERVICES

- The EFI_SYSTEM_TABLE parameter to “efi_main” is the topmost data structure passed by the UEFI firmware which contains pointers to many other structures
- EFI_BOOT_SERVICES deserves to be saved separately
 - Declare a global variable EFI_BOOT_SERVICES *BootServices;
 - BootServices = SystemTable->BootServices;

Assignment 1: UEFI Memory Allocation

```
static VOID *AllocatePool(UINTN size)
{
    VOID *ptr;
    EFI_STATUS ret = BootServices->AllocatePool(EfiBootServicesData, size, &ptr);
    if (EFI_ERROR(ret))
        return NULL;
    return ptr;
}

static VOID FreePool(VOID *buf)
{
    BootServices->FreePool(buf);
}
```

Assignment 1: UEFI Memory Allocation

```
static VOID *AllocatePool(UINTN size)
{
    VOID *ptr;
    EFI_STATUS ret = BootServices->AllocatePool(EfiBootServicesData, size, &ptr);
    if (EFI_ERROR(ret))
        return NULL;
    return ptr;
}
```

```
static VOID FreePool(VOID *buf)
{
    BootServices->FreePool(buf);
}
```

- You need to select the allocation type, for this example we assume memory for Boot Services (EfiBootServicesData)

Assignment 1:

EFI_LOADED_IMAGE_PROTOCOL_GUID

```
#include <Protocol/LoadedImage.h>
```

```
EFI_GUID gEfiLoadedImageProtocolGuid = EFI_LOADED_IMAGE_PROTOCOL_GUID;
```

Get LoadedImage Protocol:

```
EFI_LOADED_IMAGE *li = NULL;
```

```
EFI_STATUS efi_status;
```

```
efi_status = BootServices->HandleProtocol(ImageHandle, &gEfiLoadedImageProtocolGuid,  
                                           (void **)&li);
```

```
if (EFI_ERROR(efi_status)) {  
    SystemTable->ConOut->OutputString(SystemTable->ConOut,  
                                       L"Cannot get LoadedImage for BOOTx64.EFI\r\n");  
    return efi_status;  
}
```


Assignment 1:

EFI_SIMPLE_FILE_SYSTEM_PROTOCOL

Get SimpleFileSystem Protocol:

```
#include <Protocol/SimpleFileSystem.h>
```

```
EFI_GUID gEfiSimpleFileSystemProtocolGuid =  
EFI_SIMPLE_FILE_SYSTEM_PROTOCOL_GUID;
```

```
EFI_FILE_IO_INTERFACE *fio = NULL;
```

```
efi_status = BootServices->HandleProtocol(li->DeviceHandle,  
    &gEfiSimpleFileSystemProtocolGuid,  
    (void **) &fio);
```

```
if (EFI_ERROR(efi_status)) {  
    SystemTable->ConOut->OutputString(SystemTable->ConOut, L"Cannot get fio\r\n");  
    return efi_status;  
}
```

Assignment 1:

EFI_SIMPLE_FILE_SYSTEM_PROTOCOL

Get Volume and File Handles:

```
EFI_FILE_PROTOCOL *vh = NULL;  
EFI_FILE_PROTOCOL *fh = NULL;
```

```
efi_status = fio->OpenVolume(fio, &vh);  
if (EFI_ERROR(efi_status)) {  
    SystemTable->ConOut->OutputString(SystemTable->ConOut, L"Cannot get vh!\r\n");  
    return efi_status;  
}
```

```
efi_status = vh->Open(vh, &fh, L"\\EFI\\BOOT\\KERNEL",  
    EFI_FILE_MODE_READ, 0);  
if (EFI_ERROR(efi_status)) {  
    SystemTable->ConOut->OutputString(SystemTable->ConOut, L"Cannot get fh!\r\n");  
    return efi_status;  
}
```

Assignment 1: EFI_SYSTEM_TABLE and EFI_BOOT_SERVICES

- Take a look at EFI_FILE_PROTOCOL
 - Specifically, Read, Write, Close methods, etc
 - Read the kernel file (what memory type do you need to use for the buffer?)
- Take a look at EFI_BOOT_SERVICES (BootServices)
 - ExitBootServices()
 - GetMemoryMap()
 - Jump to the kernel
 - Use an ELF position-independent executable (PIE) and call the entry point function

Assignment 1: Video Framebuffer

- Text mode vs. Graphics Mode
 - Video adapters and/or firmware can work in both modes
 - For VGA adapters, 8-bit ASCII fonts can be loaded for 80x25 (standard) or other text modes
 - Colors are supported in both modes
- Text mode manipulates entire **characters**, e.g., 80x25
- Graphics mode manipulates **pixels** (tiny dots on the screen), e.g., 640x480, 800x600, 1024x768, ...

Example: Text mode (TUI)

```
mc [ruslan@ruslan-ThinkPad-T470p]:/

Left      File      Command  Options  Right
<-- /home -----.[^]>  <-- / -----.[^]>
.n      Name      Size      Modify   time
/..      UP--DIR  Oct 20 03:28
/ruslan  28672    Feb  3 13:06

/home
Size      Modify   time
4096      Oct 20 04:06

/home
63G/366G (17%)

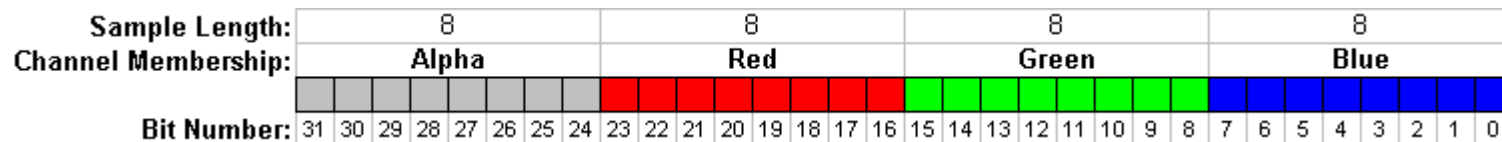
Hint: Want your plain shell? Press C-o, and get back to MC with C-o again.
ruslan@ruslan-ThinkPad-T470p:/$ [^]
1Help  2Menu  3View  4Edit  5Copy  6RenMov  7Mkdir  8Delete  9PullDn 10Quit
```

Assignment 1: Video Framebuffer

- Text mode can be emulated in graphics mode
 - Load some font and render its pixels manually
 - Used in Linux for the "framebuffer console"
 - "Terminal" in GNOME or KDE (Linux)
- Examples and flexibility:
 - `OutputString(L"...")` in UEFI is text mode (UTF-16)
 - If we render pixels ourselves, we are not restricted by any specific encoding or font

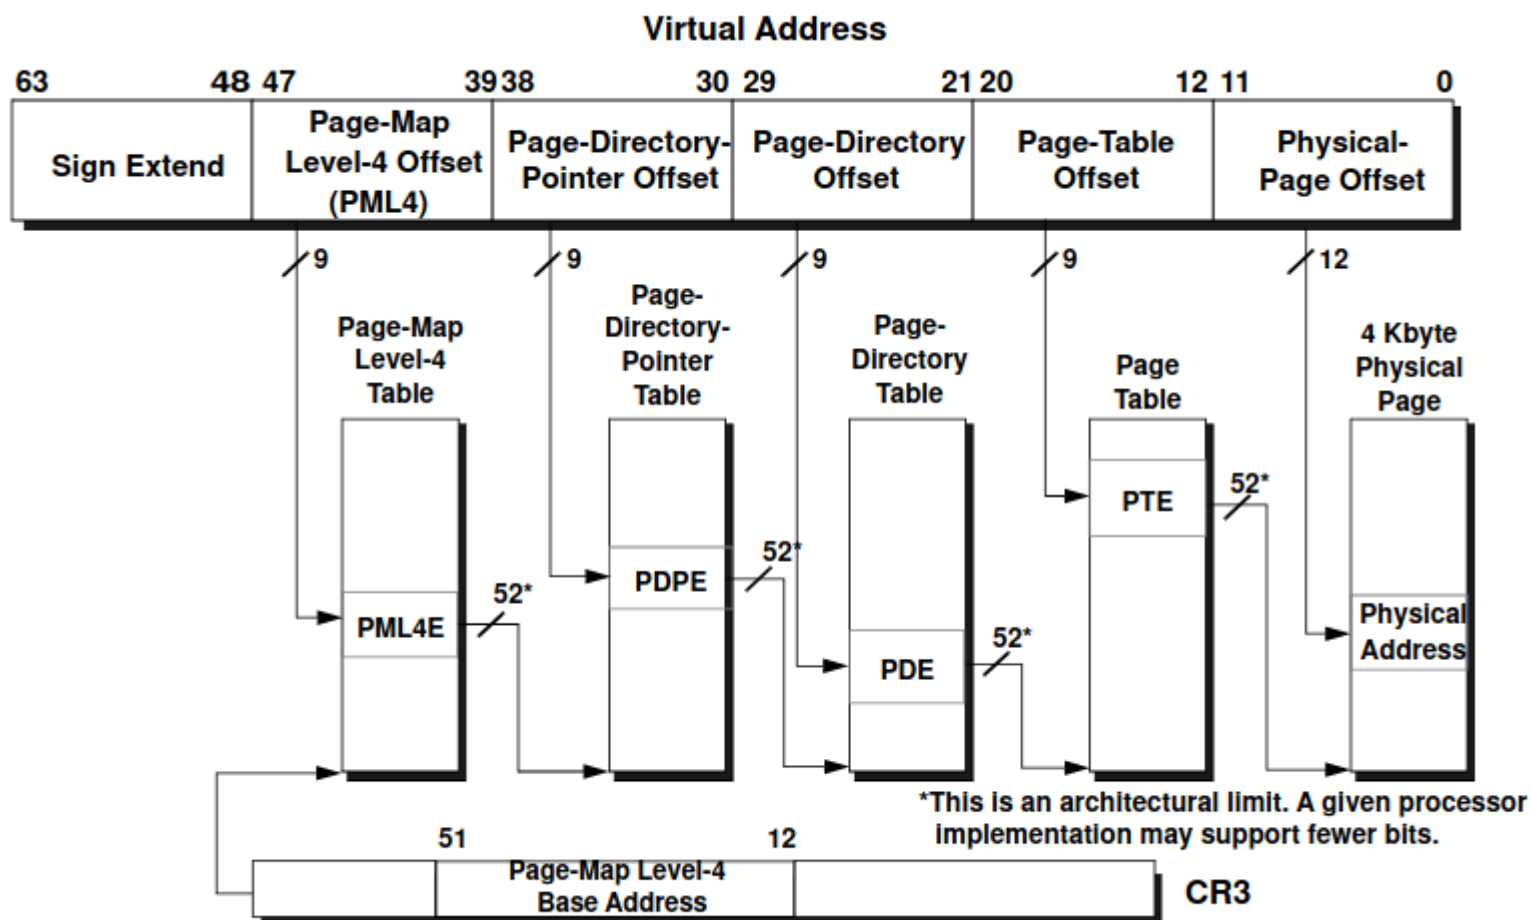
Assignment 1: Video Framebuffer

- Once you switch to graphics mode, you should not be using `OutputString()` anymore
- Draw some figures using pixels
 - Each pixel is a 32-bit (unsigned int) integer
 - `0xZZZZZZZZ`, 4 bytes (BGRA), 2 hexadecimal digits for each color component

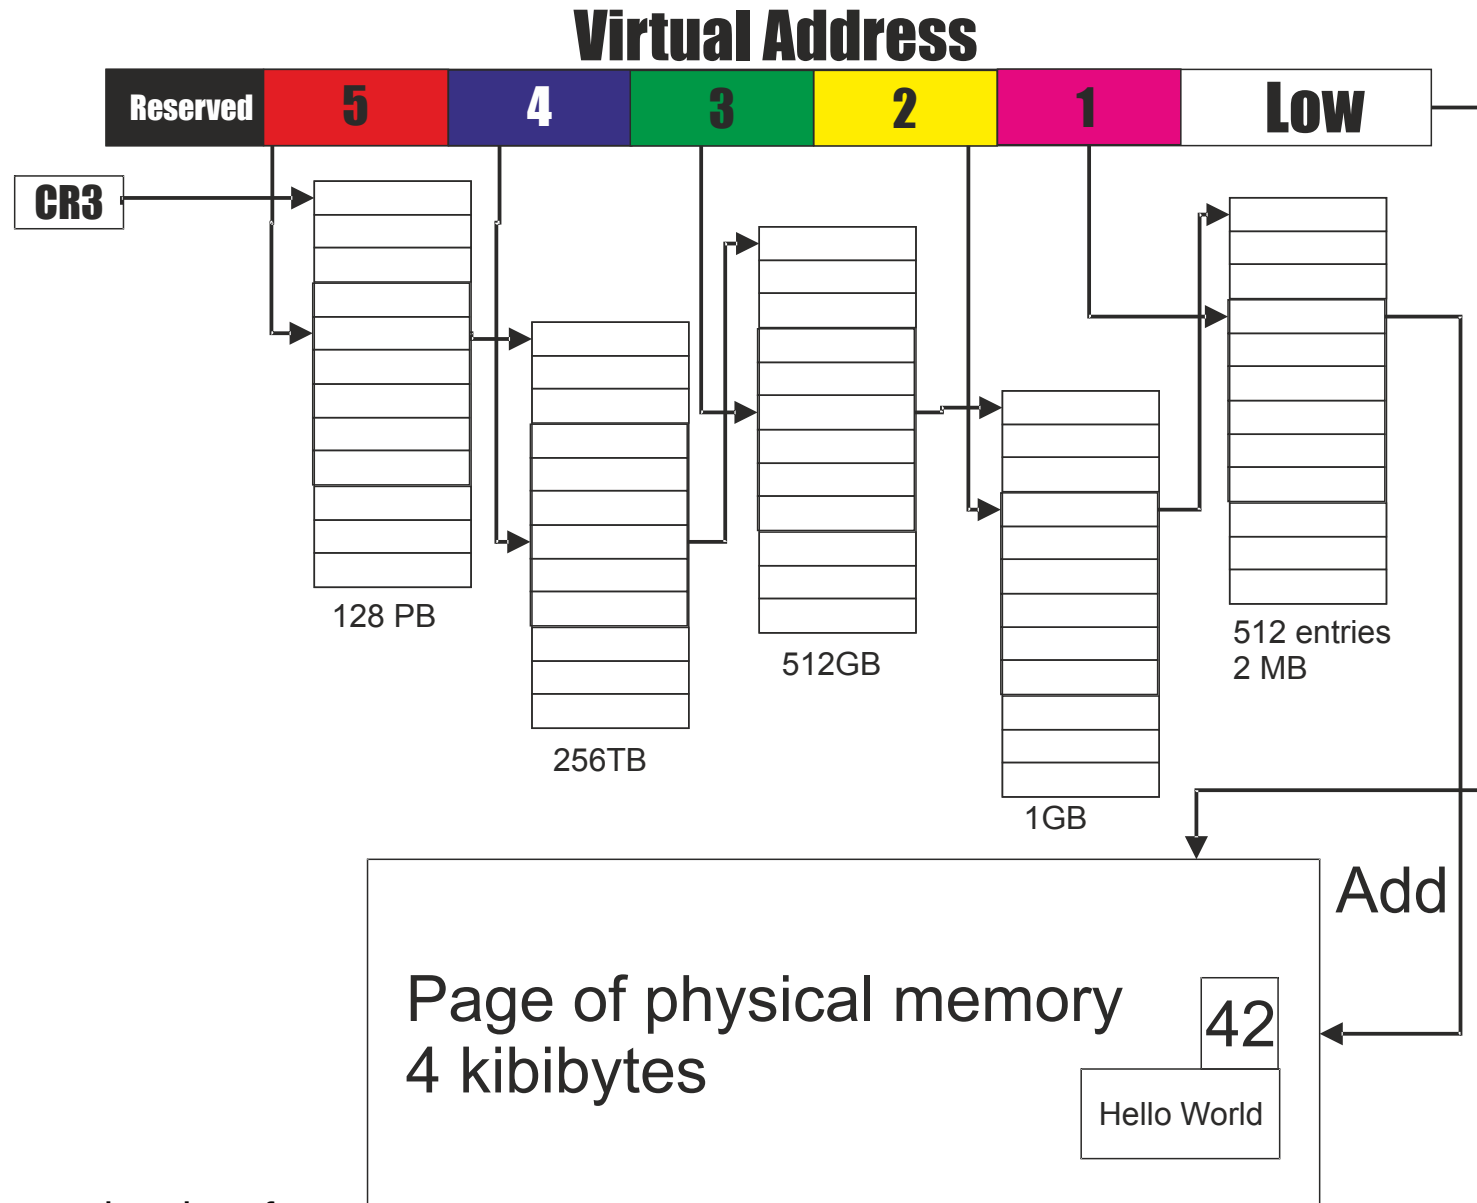


* The picture is taken from
https://en.wikipedia.org/wiki/RGBA_color_model

Traditional 48-bit paging (4 levels)



Extension: 57-bit paging (5 levels)



* The picture is taken from https://en.wikipedia.org/wiki/Intel_5-level_paging

Page Table

```
cat /proc/cpuinfo
```

```
ruslan@ruslan-ThinkPad-T470p: ~  
c cpuid aperfmperf pni pclmulqdq dtes64 monitor ds_cpl vmx smx est tm2 ssse3 sdb  
g fma cx16 xtpr pdcm pcid sse4_1 sse4_2 x2apic movbe popcnt tsc_deadline_timer a  
es xsave avx f16c rdrand lahf_lm abm 3dnowprefetch cpuid_fault epb invpcid_singl  
e pti ssbd ibrs ibpb stibp tpr_shadow vnmi flexpriority ept vpid ept_ad fsgsbase  
tsc_adjust bmi1 hle avx2 smep bmi2 erms invpcid rtm mpx rdseed adx smap clflush  
opt intel_pt xsaveopt xsavec xgetbv1 xsaves dtherm ida arat pln pts hwp hwp_noti  
fy hwp_act_window hwp_epp md_clear flush_l1d  
vmx flags      : vnmi preemption_timer invvpid ept_x_only ept_ad ept_1gb flexpr  
iority tsc_offset vtptr mtf vapic ept vpid unrestricted_guest ple shadow_vmcs pml  
ept_mode_based_exec  
bugs           : cpu_meltdown spectre_v1 spectre_v2 spec_store_bypass l1tf mds  
swapgs taa itlb_multihit srbds  
bogomips       : 5799.77  
clflush size   : 64  
cache_alignment : 64  
address sizes  : 39 bits physical, 48 bits virtual  
power management:  
ruslan@ruslan-ThinkPad-T470p:~$
```

100

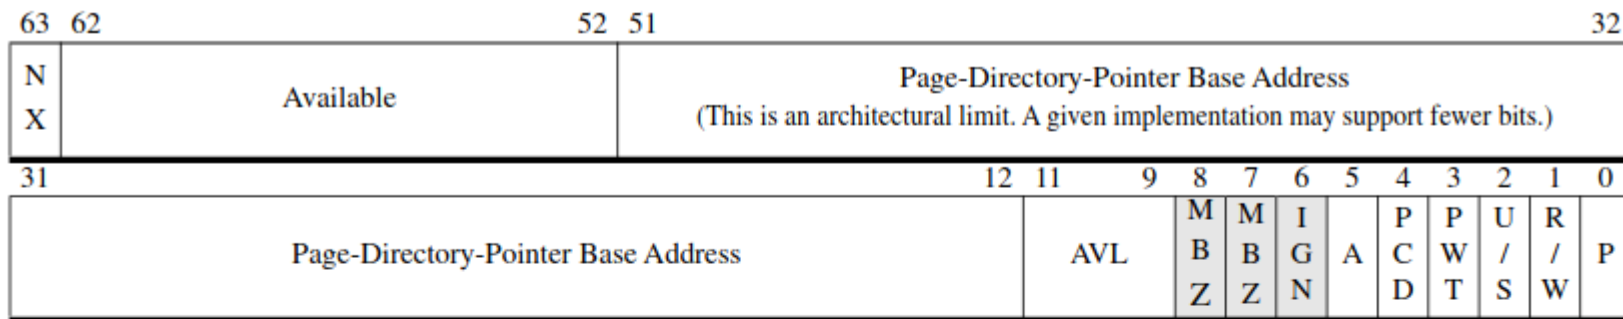


Figure 5-18. 4-Kbyte PML4E—Long Mode

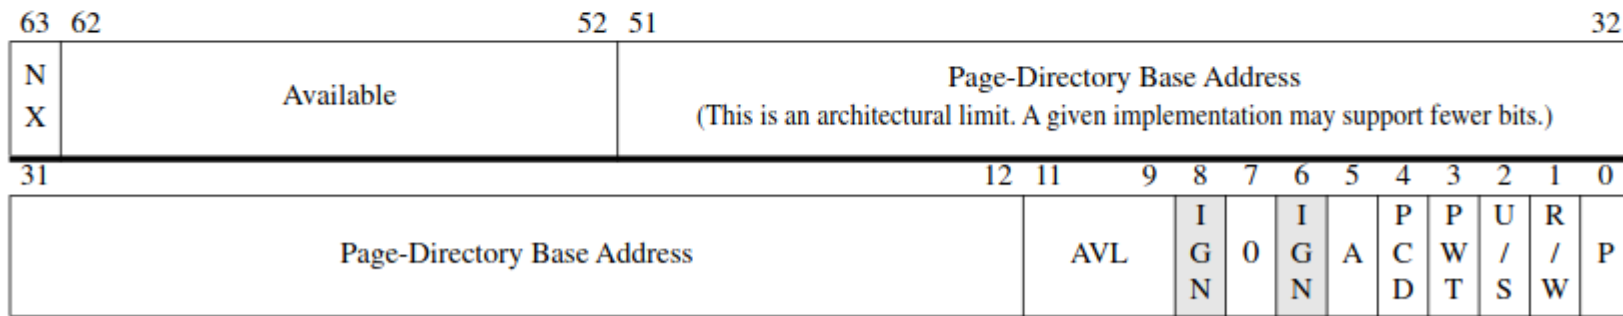


Figure 5-19. 4-Kbyte PDPE—Long Mode

100

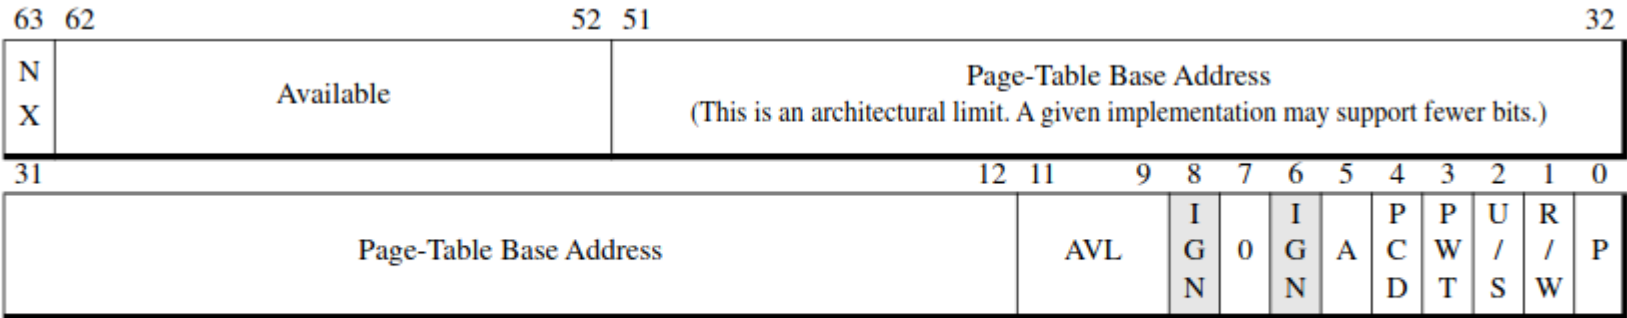


Figure 5-20. 4-Kbyte PDE—Long Mode

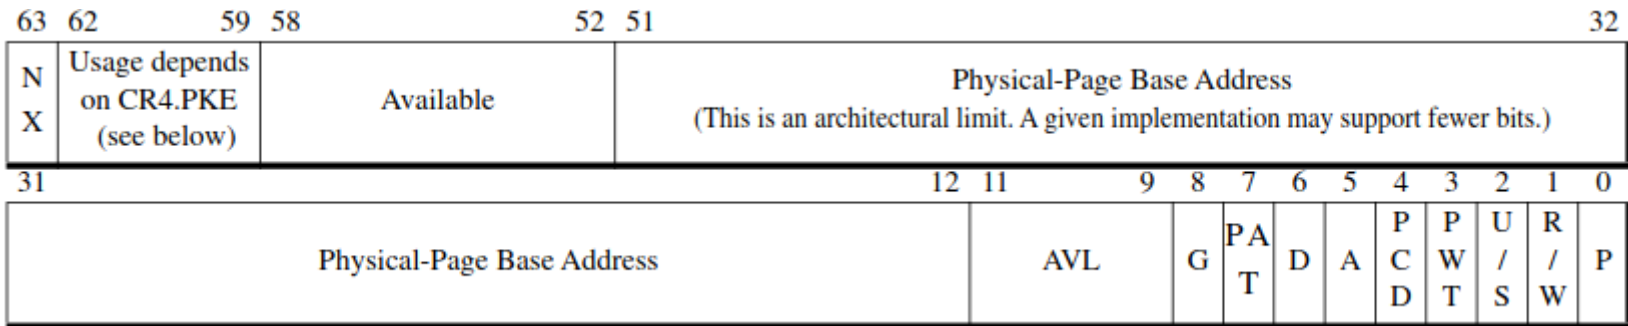


Figure 5-21. 4-Kbyte PTE—Long Mode

Example: PTE

```
typedef unsigned long long u64;

struct page_pte {
    u64 present:1;           // Bit P
    u64 writable:1;          // Bit R/W
    u64 user_mode:1;         // Bit U/S
    ...
    u64 page_address:40;     // 40+12 = 52-bit physical address (max)
    U64 avail:7;             // reserved, should be 0
    u64 pke:4;               // no MPK/PKE, should be 0
    u64 nonexecute:1;
};
```

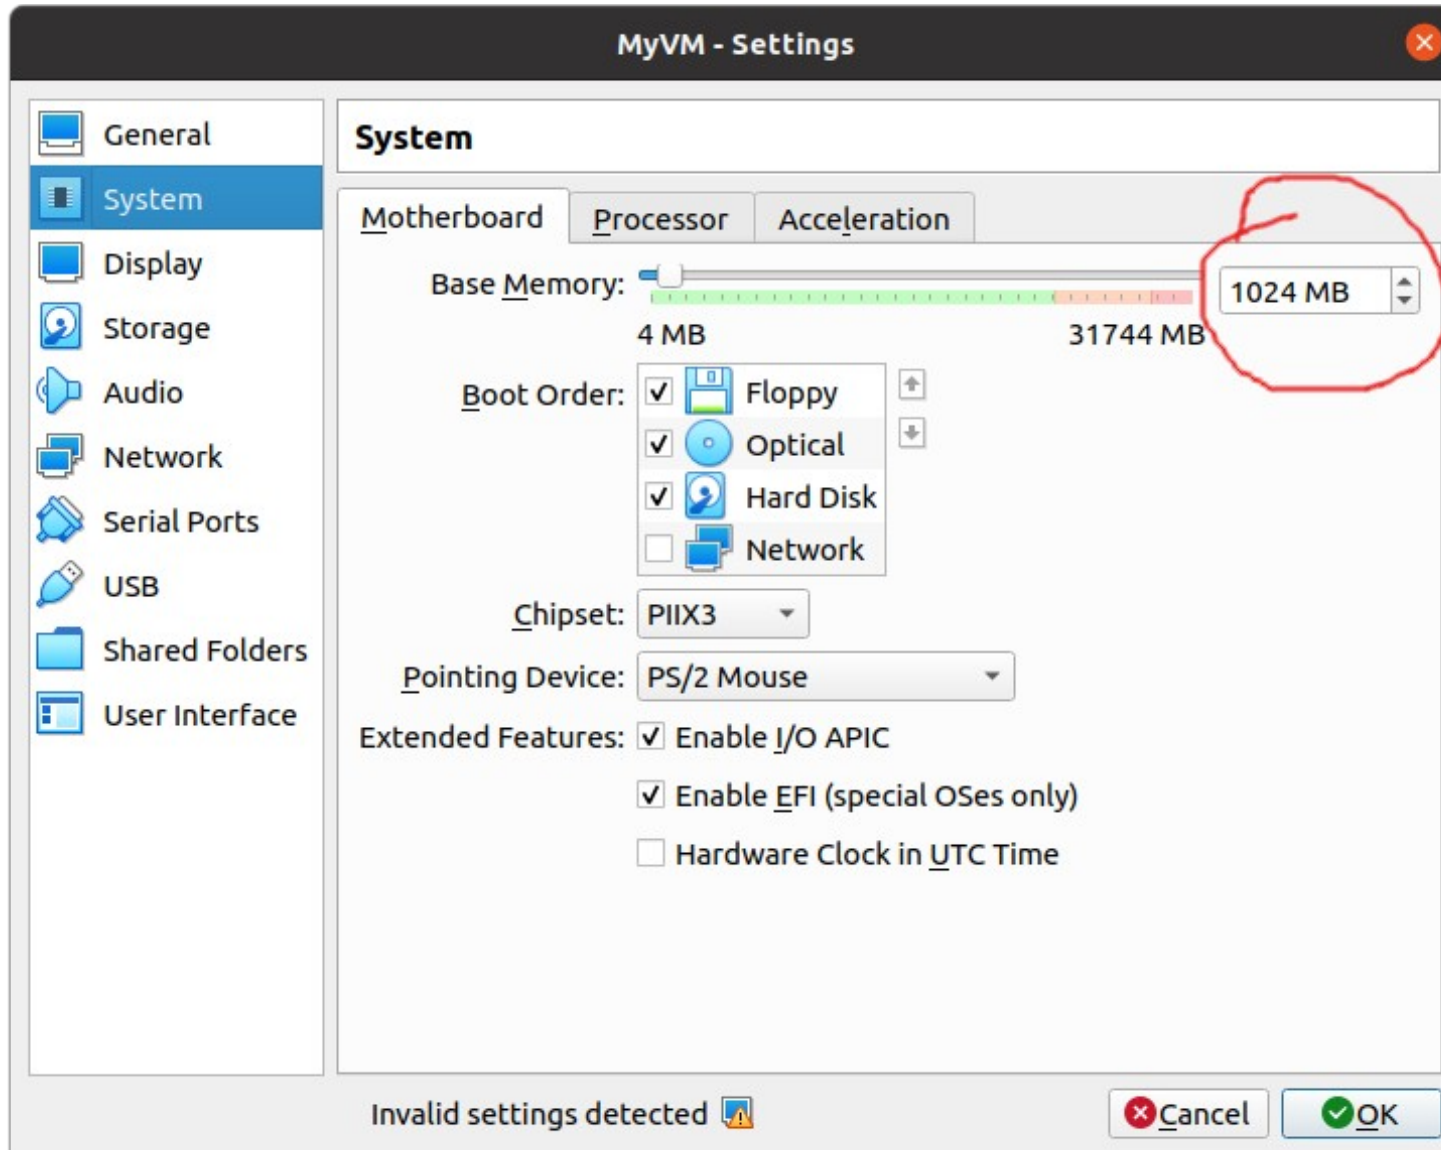
Example: PTE

```
typedef unsigned long long u64;

struct page_pte {
    u64 present:1;          // Bit P
    u64 writable:1;         // Bit R/W
    u64 user_mode:1;        // Bit U/S
    ...
    u64 page_address:40;    // 40+12 = 52-bit physical address (max)
    u64 avail:7;           // reserved, should be 0
    u64 pke:4;              // no MPK/PKE, should be 0
    u64 nonexecute:1;
};

// _page_ address, i.e. memory_address / 4096 (or 2^12)
p->page_address = 0xZZZZZZZZZZULL; // ULL is unsigned long long
p->writable = 1;
p->present = 1;
p->user_mode = 0;
...
p->avail = 0;
p->pke = 0;
p->nonexecute = 0;
```

Example: Initializing 4 GB



1GB should
be sufficient

Video RAM can
be after that

Example: Initializing 4 GB

For 4GB, we reference 1048576 physical pages

Using 2048 pages for PTEs, 4 pages for PDEs, 1 for PDPE, 1 for PMLE4E

Total: 2054 pages = 8413184 bytes, align at the 4096 boundary!

PTE:

```
struct page_pte *p; ... // Each entry is 8 bytes
```

```
for (int i = 0; i < 1048576; i++) { // 1048576/512 = 2048 pages
    p[i] = ... // physical pages 0,1,2,3,...,1048575 (absolute address)
}
```

PDE:

```
struct page_pde *pd = (struct page_pde *) (p + 1048576);
for (int j = 0; j < 2048; j++) { // 2048/512 = 4 pages
    struct page_pte *start_pte = p + 512 * j;
    page_addr = (u64) start_pte >> 12; // we record the page address
    ...
}
```

PDPE:

Reference 4 PDEs (1 page), everything else is empty

PMLE4E:

Just one reference to PDPE; everything else is empty

Aligning Pages

Allocate more space: e.g., $8413184 + 4095$

Align the allocated 'base':

```
(void *) (((unsigned long long) base + 4095) & (~4095ULL))
```

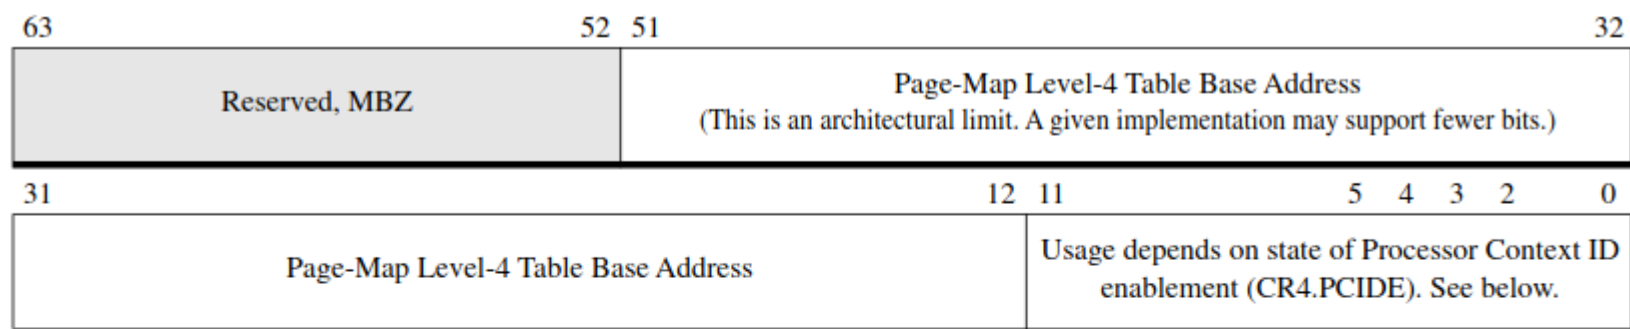
- Why long long?
 - The kernel uses 64-bit 'long' due to System V's ABI (aka the LP64 model)
 - The boot loader uses 32-bit 'long' due to EFI/Microsoft's ABI (aka the LLP64 model)
 - 'int' is 32 bit and 'long long' is 64 bit in either case

Loading Page Table

```
void write_cr3(unsigned long long cr3_value)
{
    asm volatile ("mov %0, %%cr3"
                  :
                  : "r" (cr3_value)
                  : "memory");
}
```

Loading Page Table

```
void write_cr3(unsigned long long cr3_value)
{
    asm volatile ("mov %0, %%cr3"
                  :
                  : "r" (cr3_value)
                  : "memory");
}
```



No PCID
(PCIDE=0)!

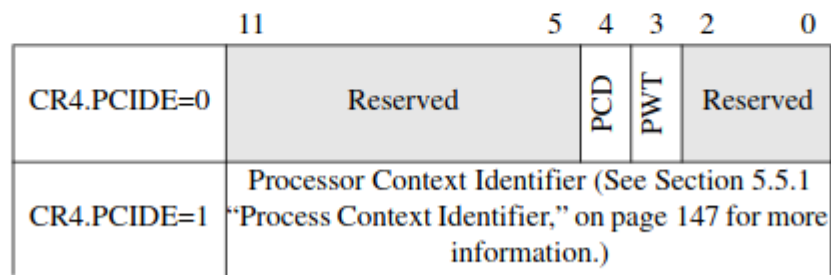


Figure 3-6. Control Register 3 (CR3)—Long Mode

* The picture is taken from <https://www.amd.com/system/files/TechDocs/24593.pdf>

Loading Page Table

Page-Level Writethrough (PWT) Bit. Bit 3. Page-level writethrough indicates whether the highest-level page-translation table has a writeback or writethrough caching policy. When PWT=0, the table has a writeback caching policy. When PWT=1, the table has a writethrough caching policy.

Page-Level Cache Disable (PCD) Bit. Bit 4. Page-level cache disable indicates whether the highest-level page-translation table is cacheable. When PCD=0, the table is cacheable. When PCD=1, the table is not cacheable.

PCD=0, PWT=0